



Use Analytics Zoo to build an intelligent recommendation system on Office Depot

Kai Huang

Mar 26th, 2020

OUTLINE

- **Background and use case overview**
- **Introduction to Analytic Zoo**
- **Recommenders on Analytics Zoo**
- **Performance and deployment by Office Depot**
- **Conclusion**



Why Recommendation Systems?

- Help customers choose from a variety of products.
- Maintain user satisfaction and royalty.
- Turn ordinary users into potential customers.
- Increase revenue per user visit.
-



Big Data Journey for Recommendation



Stage I :

Office Depot tried to build intelligent models for product recommendation using Python/SAS/R.

Challenges:

They can not process this amount of data on a single machine:

- Over 100,000,000 distinct sessions monthly.
- More than 300,000 active products selling online.
- Training data often exceed 10G.

Big Data Journey for Recommendation

Stage II :

Office Depot incorporated Spark and AWS cloud into their workflow.




Challenge:

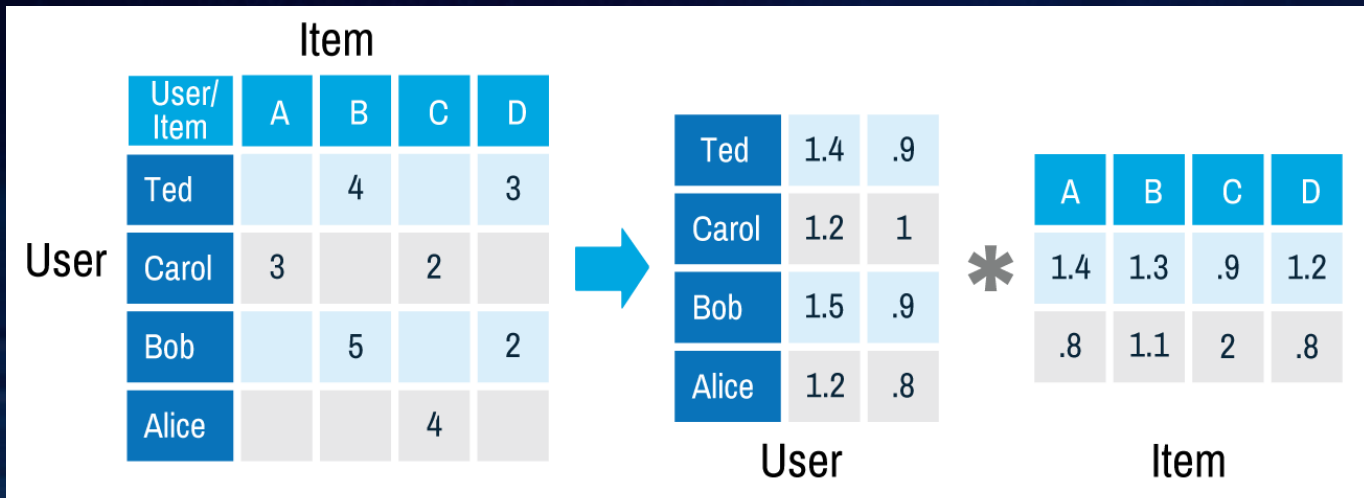
Deep learning libraries such as TensorFlow/Keras/PyTorch cannot run directly on Spark clusters.

Why deep learning?

- Better performance on larger data.
- Less manual feature engineering needed.
- Easier to involve complex functions and combine different architectures.



Collaborative Filtering (ALS)




- The Collaborative filtering approach works by decomposing the user-item interaction matrix into the product of two lower dimensionality rectangular matrices.
- Spark ALS (Alternating Least Squares) implementation runs matrix factorization in a parallel fashion and therefore has a pretty good scalability and performance.

Collaborative Filtering (ALS)

	i_1	i_2	i_3	i_4	i_5
u_1	1	1	1	0	1
u_2	0	1	1	0	0
u_3	0	1	1	1	0
u_4	1	0	1	1	1

(a) user-item matrix



(b) user latent space

Figure 1: An example illustrates MF's limitation. From data matrix (a), u_4 is most similar to u_1 , followed by u_3 , and lastly u_2 . However in the latent space (b), placing p_4 closest to p_1 makes p_4 closer to p_2 than p_3 , incurring a large ranking loss.

Limitations of matrix factorization:

- Simple choice of the interaction function will hinder the performance.
- Data sparse problem.
- Not able to do incremental training.
- Cold start problem.
- Not able to capture the latest purchase intent.

...



Distributed, High-Performance
Deep Learning Framework
for Apache Spark

<https://github.com/intel-analytics/BigDL>



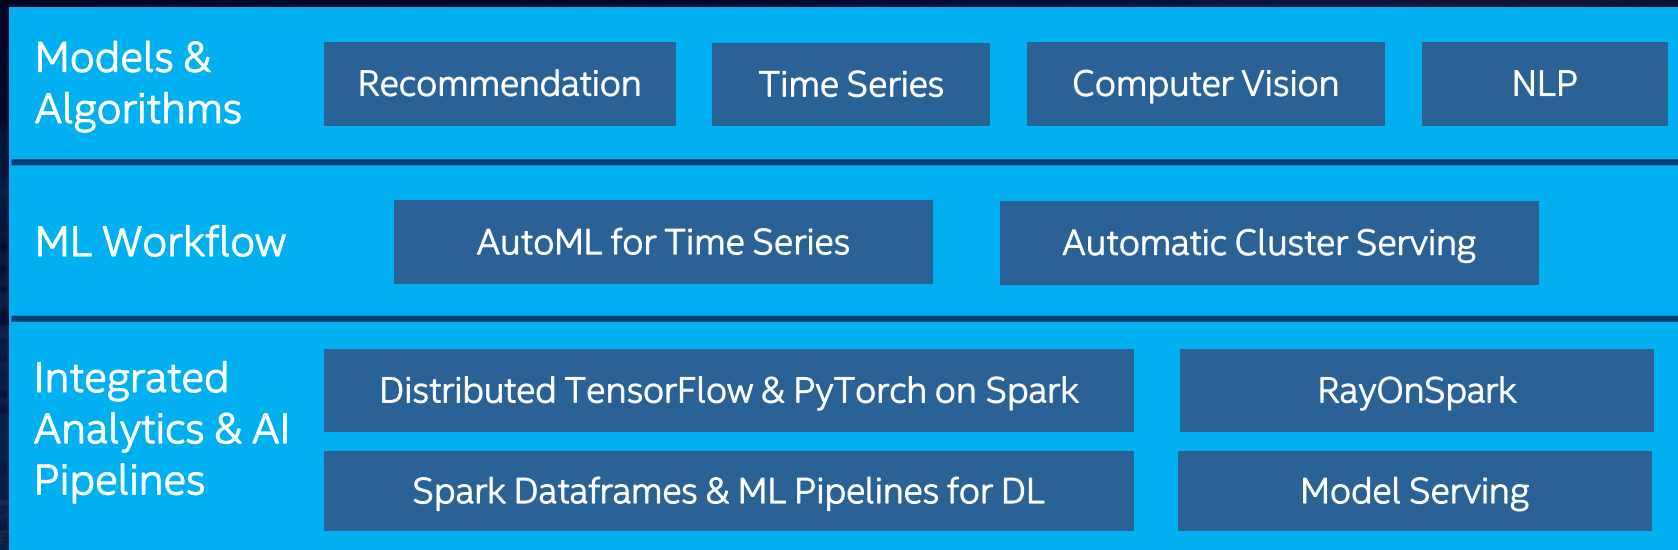
A unified analytics and AI platform
for distributed Tensorflow, Keras, PyTorch and Ray
on Apache Spark

<https://github.com/intel-analytics/analytics-zoo>

Accelerating Data Analytics + AI Solutions At Scale

Analytics Zoo

Unified Big Data Analytics and AI Platform



Library & Framework

Distributions
(Cloudera/Databricks/....)

Distributed Analytics
(Spark/Flink/Ray/...)

DL Frameworks
(TF/PyTorch/...)

Python Libraries
(Numpy/Pandas/...)

<https://github.com/intel-analytics/analytics-zoo>

Unified Big Data Analytics and AI Platform

Seamless Scaling from Laptop to Production

Prototype on **laptop**
using sample data



Experiment on **clusters**
with history data



Production deployment w/
distributed data pipeline



Production
Data pipeline



- Easily prototype the **integrated data analytics & AI solution**
- **“Zero” code change** from laptop to distributed cluster
- **Directly access production data** (Hadoop/Hive/HBase) without data copy
- Seamlessly deployed on **production big data clusters**

Real-World Applications

NLP Based Customer Service Chatbot for Microsoft Azure*

<https://software.intel.com/en-us/articles/use-analytics-zoo-to-inject-ai-into-customer-service-platforms-on-microsoft-azure-part-1>

<https://www.infoq.com/articles/analytics-zoo-qa-module/>

Industrial Product Defect Detection in Midea*

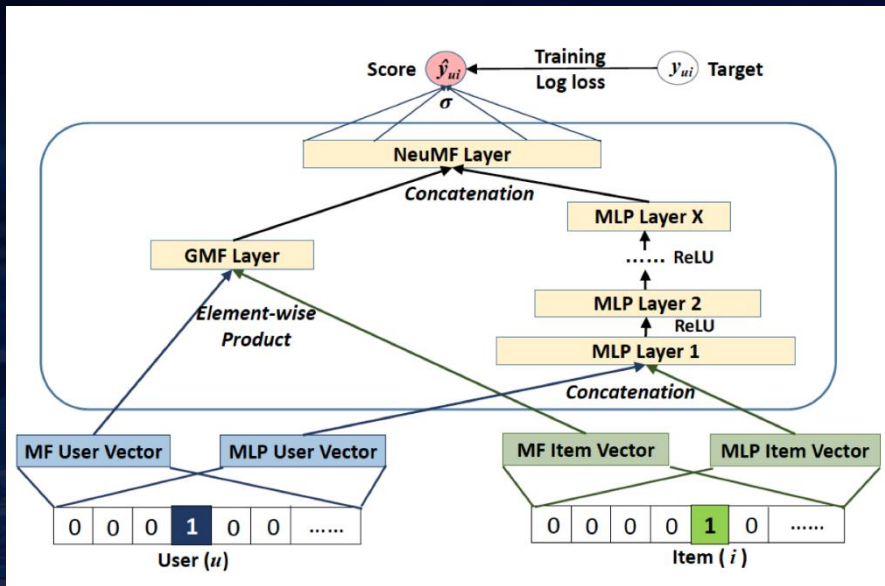
<https://software.intel.com/en-us/articles/industrial-inspection-platform-in-midea-and-kuka-using-distributed-tensorflow-on-analytics>

Unsupervised Time Series Anomaly Detection for Baosight*

<https://software.intel.com/en-us/articles/lstm-based-time-series-anomaly-detection-using-analytics-zoo-for-apache-spark-and-bigdl>

Any many more...

Neural Collaborative Filtering (NCF)

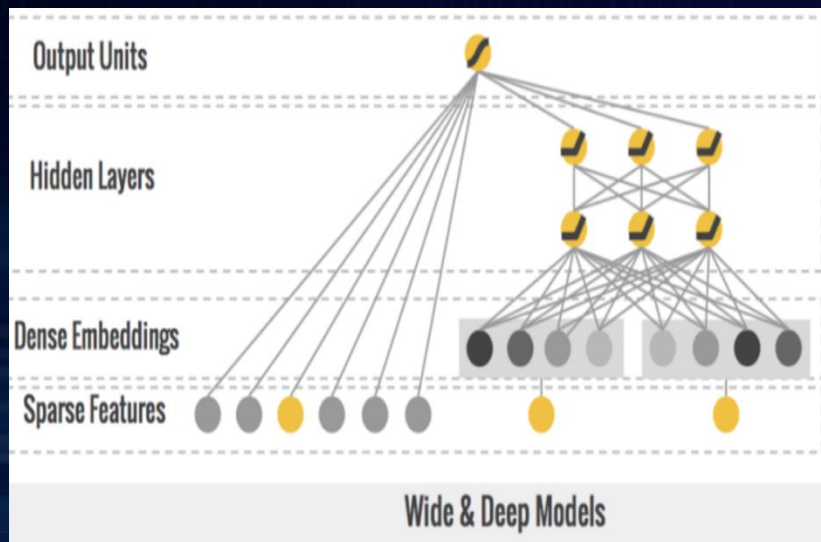


- NCF stimulates matrix factorization using DNN approach and is served as a guideline for deep learning methods for recommendation services.
- It combines GMF with MLP to model user-item interactions.

```
01. from zoo.models.recommendation import NeuralCF
02.
03. ncf = NeuralCF(user_count, item_count, class_num, user_embed=20,
04.                item_embed=20, hidden_layers=[40, 20, 10],
05.                include_mf=True, mf_embed=20)
06. ncf.compile(optimizer="adam",
07.             loss="sparse_categorical_crossentropy",
08.             metrics=['accuracy'])
09. ncf.fit(train_rdd,
10.         nb_epoch,
11.         batch_size,
12.         validation_data=val_rdd)
```

<https://github.com/intel-analytics/analytics-zoo/tree/master/apps/recommendation-ncf>

Wide & Deep Learning



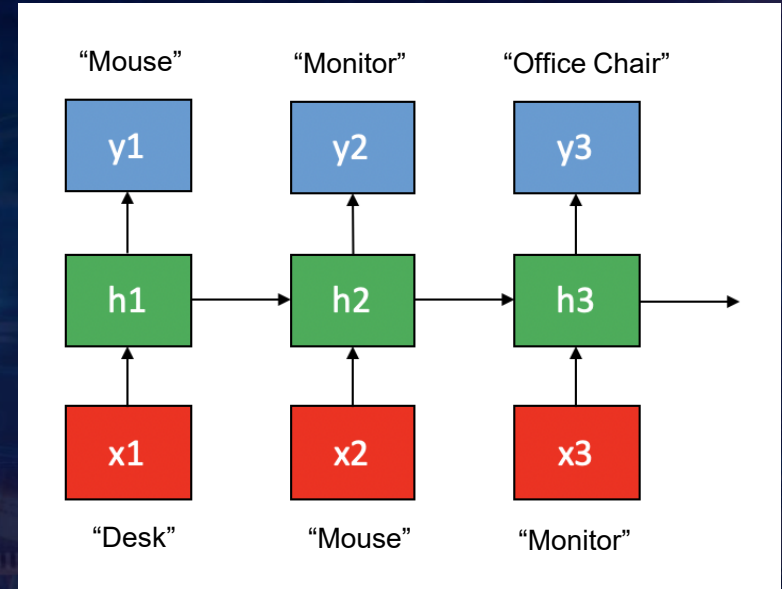
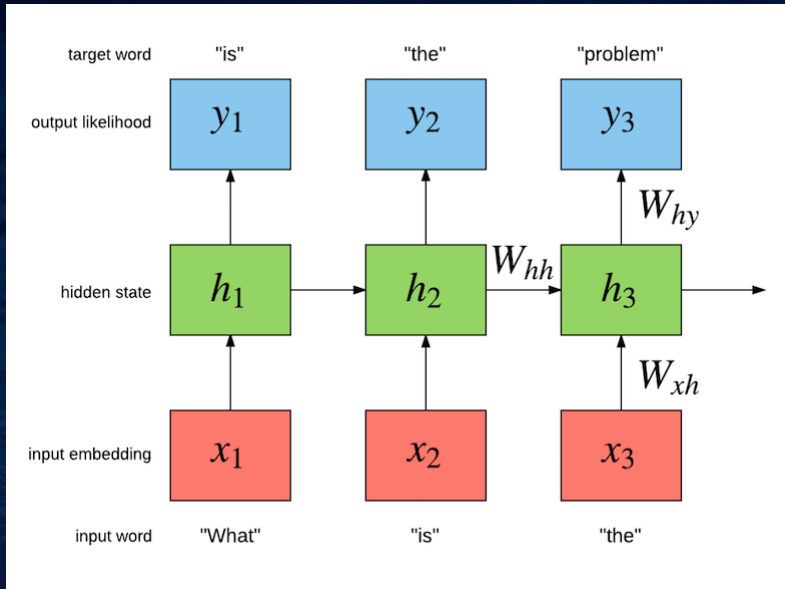
- Wide and Deep learning model can take rich data as input.
- The wide part can effectively memorize sparse feature interactions using cross-product feature transformations.
- The deep part can generalize to previously unseen feature interactions through low dimensional user and item embeddings similar to NCF.

```
01. from zoo.models.recommendation import WideAndDeep
02.
03. # column_info can be shared by feature generation and model, where you can specify
04. # columns and their dimensions for each part of the model.
05. wnd = WideAndDeep(user_count, class_num, column_info,
06.                    model_type="wide_n_deep", hidden_layers=[40, 20, 10])
```

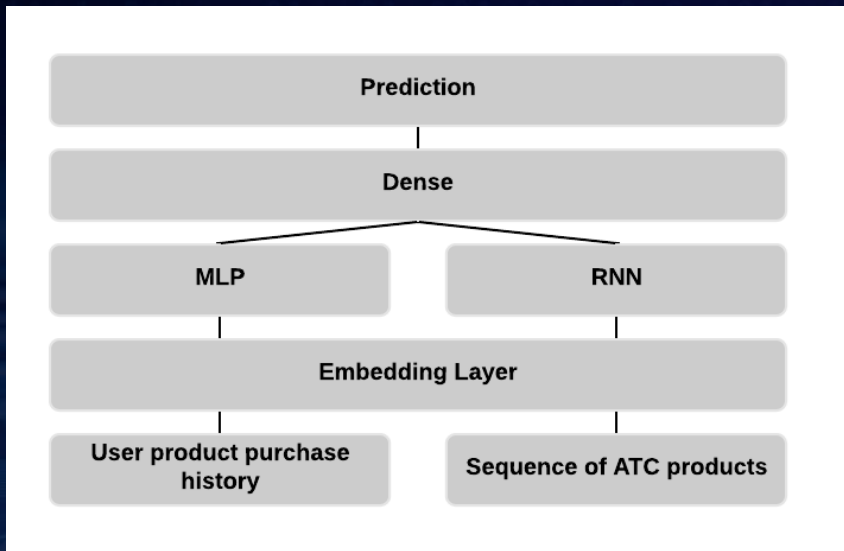
<https://github.com/intel-analytics/analytics-zoo/tree/master/apps/recommendation-wide-n-deep>

Session Recommender

- Each user session in an e-commerce system could be modeled as a sequence of web pages.
- A deep RNN could track how users browse the website using multiple hidden layers.



Session Recommender



The Good:

- Can catch the latest purchase intent from current session behavior and adjust its product recommendation in real time.
- Can work with both anonymous / identified customers.
- No pre-filtering mechanism required, simpler serving architect.

The Bad:

- Sequence window size is hard to set.
- Online inference requires lots of resources.

```
01. from zoo.models.recommendation import SessionRecommender
02.
03. model = SessionRecommender(item_count, item_embed=100, rnn_hidden_layers=[40, 20],
04.                             session_length=5, include_history=True,
05.                             mlp_hidden_layers=[40, 20], history_length=10)
```

Performance Comparison

Offline measurement:

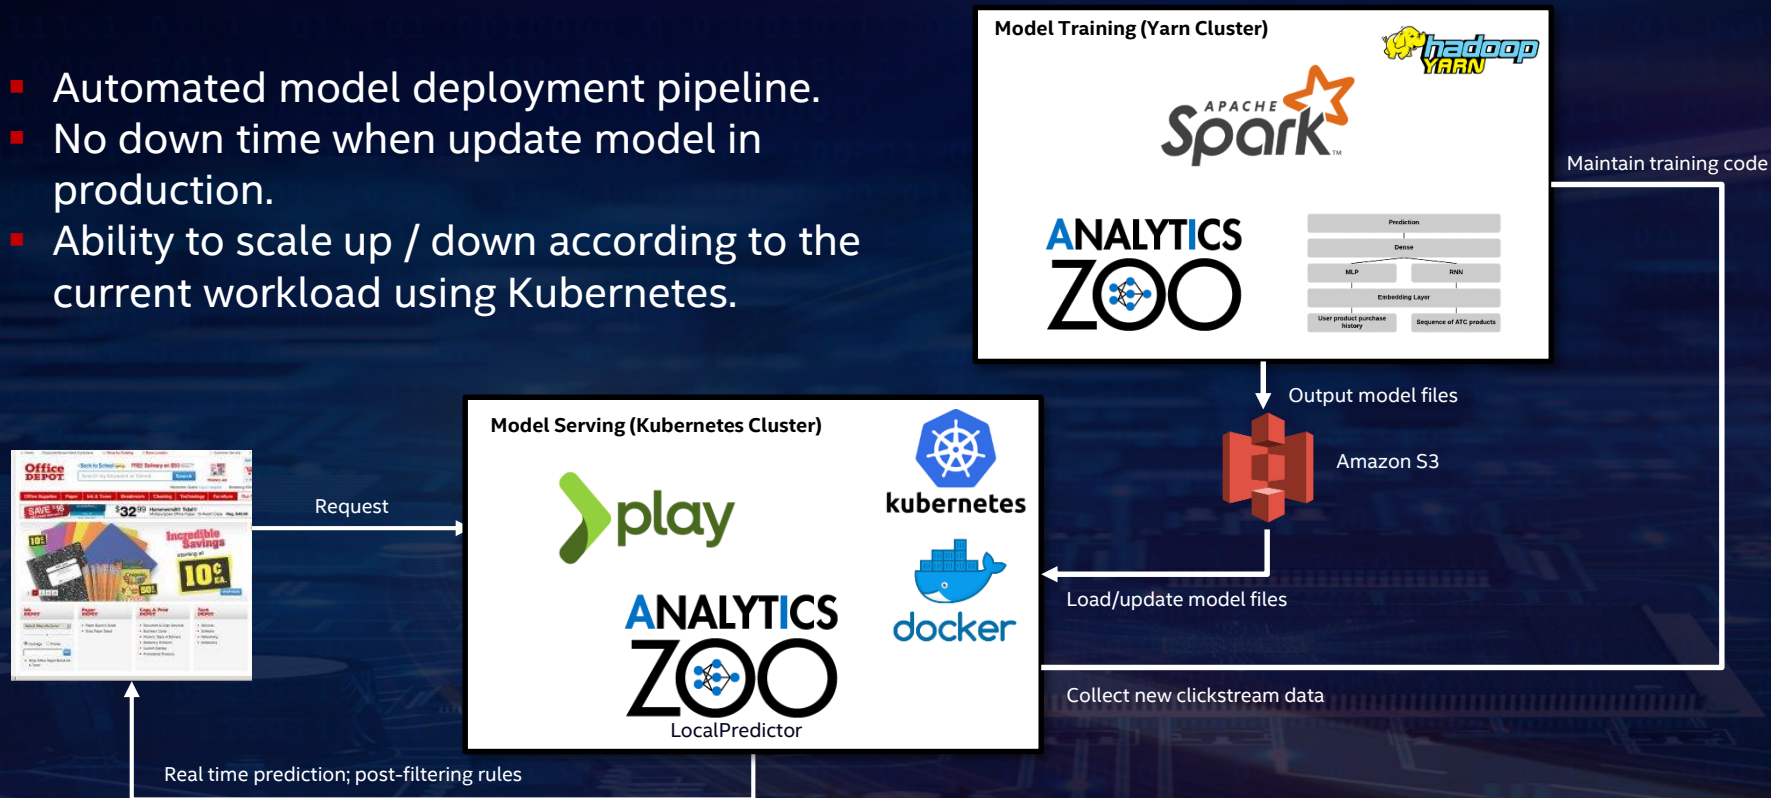
Method	Top 5 Accuracy
Session Recommender	52.3%
Wide & Deep	45.2%
NCF	46.7%
ALS	16.2%

Online measurement:

Online A/B testing shows the test group using Session Recommender lifted sales by 1% and average order value by 1.6% compared to control group.

Recommendation System In Production

- Automated model deployment pipeline.
- No down time when update model in production.
- Ability to scale up / down according to the current workload using Kubernetes.



Conclusion and Takeaways

- Analytics Zoo integrates well into existing big data pipelines.
- Analytics Zoo provides model serving API for high performance real-time inference.
- Deep learning based recommendation provides more flexibility to combine different model architectures for different use cases.
- Lots of NLP algorithms (for example, transformers) can be utilized for recommendation.
- Check out the joint blog for more information:

<https://software.intel.com/en-us/articles/real-time-product-recommendations-for-office-depot-using-apache-spark-and-analytics-zoo-on>

Analytics Zoo on Ali E-MR



+



Alibaba Cloud
aliyun.com

Analytics Zoo is already out-of-box on Ali EMR:



* Version upgrade for Analytics Zoo is on-going.

For more information and support,
contact Wesley:

Email: wesley.du@intel.com
DingTalk:



More Information on Analytics Zoo

- Project websites
 - <https://analytics-zoo.github.io/master/>
 - <https://github.com/intel-analytics/analytics-zoo>
 - <https://github.com/intel-analytics/bigdl>
- Tutorials
 - CVPR 2018: <https://jason-dai.github.io/cvpr2018/>
 - AAAI 2019: <https://jason-dai.github.io/aaai2019/>
- “BigDL: A Distributed Deep Learning Framework for Big Data”
 - *In proceedings of ACM Symposium on Cloud Computing 2019 (SOCC'19)*
 - <https://dl.acm.org/doi/10.1145/3357223.3362707>
- Use cases
 - *Microsoft Azure, CERN, MasterCard, Baosight, Tencent, Midea, etc.*
 - <https://analytics-zoo.github.io/master/#powered-by/>




ANALYTICS ZOO



Unified Analytics + AI Platform
Distributed TensorFlow, Keras and BigDL on Apache Spark
<https://github.com/intel-analytics/analytics-zoo>





LEGAL DISCLAIMERS

- Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Learn more at intel.com, or from the OEM or retailer.
- No computer system can be absolutely secure.
- Tests document performance of components on a particular test, in specific systems. Differences in hardware, software, or configuration will affect actual performance. Consult other sources of information to evaluate performance as you consider your purchase. For more complete information about performance and benchmark results, visit <http://www.intel.com/performance>.

Intel, the Intel logo, Xeon, Xeon phi, Lake Crest, etc. are trademarks of Intel Corporation in the U.S. and/or other countries.

*Other names and brands may be claimed as the property of others.

© 2019 Intel Corporation