

What is Analytics Zoo



Distributed, High-Performance
Deep Learning Framework
for Apache Spark



<https://github.com/intel-analytics/bigdl>




Unified Analytics + AI Platform
Distributed TensorFlow, Keras, PyTorch and BigDL
on Apache Spark



<https://github.com/intel-analytics/analytics-zoo>

Accelerating Data Analytics + AI Solutions At Scale



BigDL: A Distributed Deep Learning Framework for Big Data

Jason (Jinquan) Dai, Yiheng Wang, Xin Qiu, Ding Ding, Yao Zhang, Yanzhang Wang, Xianyan Jia, Cherry (Li) Zhang, Yan Wan, Zhichao Li, Jiao Wang, Shengsheng Huang, Zhongyuan Wu, Yang Wang, Yuhao Yang, Bowen She, Dongjie Shi, Qi Lu, Kai Huang, Guoqiong Song

AI on



Distributed, High-Performance
Deep Learning Framework
for Apache Spark*

<https://github.com/intel-analytics/bigdl>



Analytics + AI Platform

Distributed TensorFlow*, Keras*,
PyTorch* and BigDL on Apache Spark*

<https://github.com/intel-analytics/analytics-zoo>

Accelerating Data Analytics + AI Solutions At Scale

Agenda

- **Motivation**
- **BigDL Execution Model**
- **Experimental Evaluation**
- **Real-World Applications**
- **Future Work**

Real-World ML/DL Systems Are Complex Big Data Analytics Pipelines

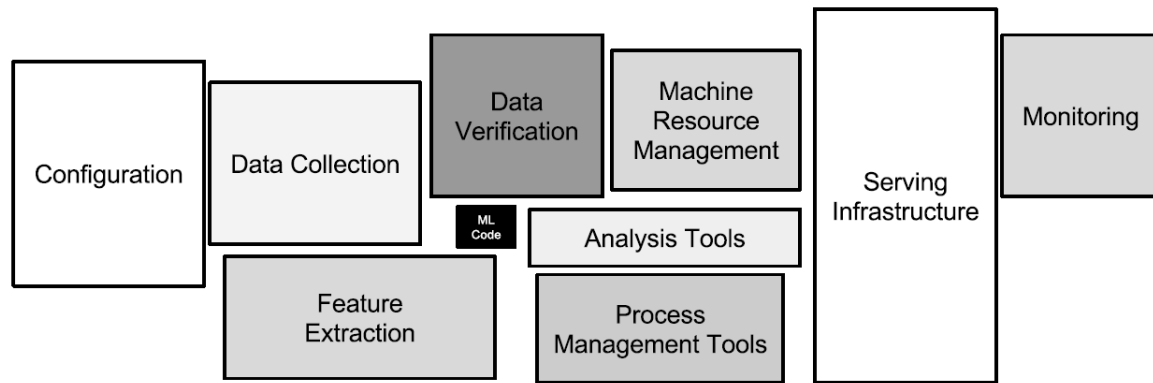



Figure 1: Only a small fraction of real-world ML systems is composed of the ML code, as shown by the small black box in the middle. The required surrounding infrastructure is vast and complex.

“Hidden Technical Debt in Machine Learning Systems”,
Sculley et al., Google, NIPS 2015 Paper

Chasm b/w Deep Learning and Big Data Communities



Big Data Analysis Challenges

Real-World data analytics and deep learning pipelines are challenging

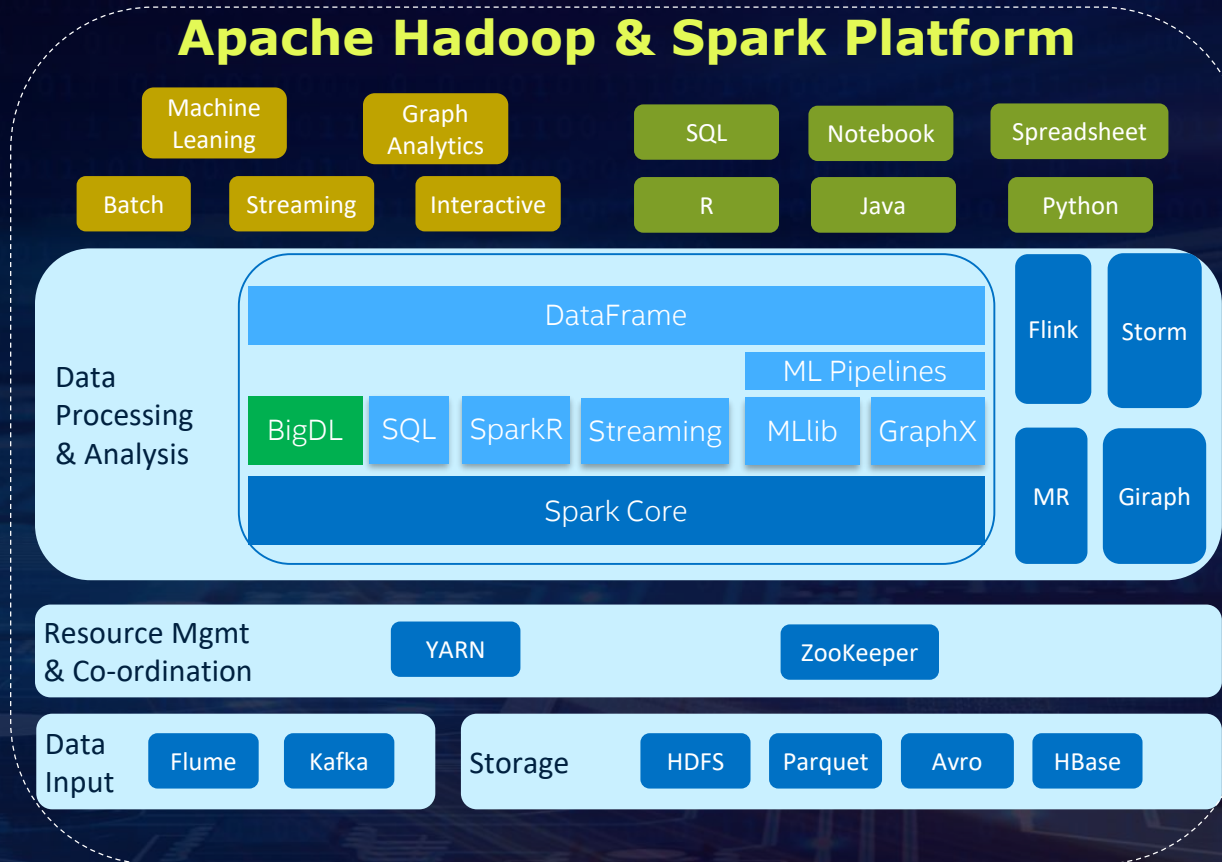
- Deep learning benchmarks (ImageNet, SQuAD , etc.)
 - Curated and explicitly labelled Dataset
 - Suitable for dedicated DL systems
- Real-world production data pipeline
 - Dynamic, messy (and possibly implicitly labeled) dataset
 - Suitable for integrated data analytics and DL pipelines using BigDL
- Problems with “connector approaches”
 - TFX, TensorFlowOnSpark, Project Hydrogen, etc.
 - Adaptation overheads, impedance mismatch

The background of the slide is a dark blue, high-tech image. It features a close-up, slightly angled view of a circuit board or microchip. Overlaid on this image is a pattern of white binary code (0s and 1s) that appears to be floating or falling from the top, creating a digital rain effect. The overall aesthetic is futuristic and technological.

BigDL Introduction

Unified Big Data Analytics Platform

Apache Hadoop & Spark Platform

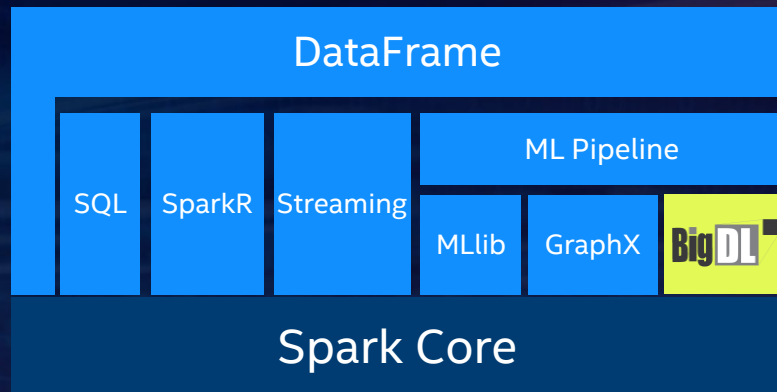


BigDL

Bringing Deep Learning To Big Data Platform



- **Distributed** deep learning framework for Apache Spark
- Make deep learning more accessible to **big data users** and **data scientists**
 - Write deep learning applications as **standard Spark programs**
 - Run on existing Spark/Hadoop clusters (**no changes needed**)
- Feature parity with popular deep learning frameworks
 - E.g., Caffe, Torch, Tensorflow, etc.
- High performance (on CPU)
 - Powered by Intel MKL and multi-threaded programming
- Efficient scale-out
 - Leveraging Spark for distributed training & inference



<https://github.com/intel-analytics/BigDL>

<https://bigdl-project.github.io/>

BigDL Run as Standard Spark Programs

Standard Spark jobs

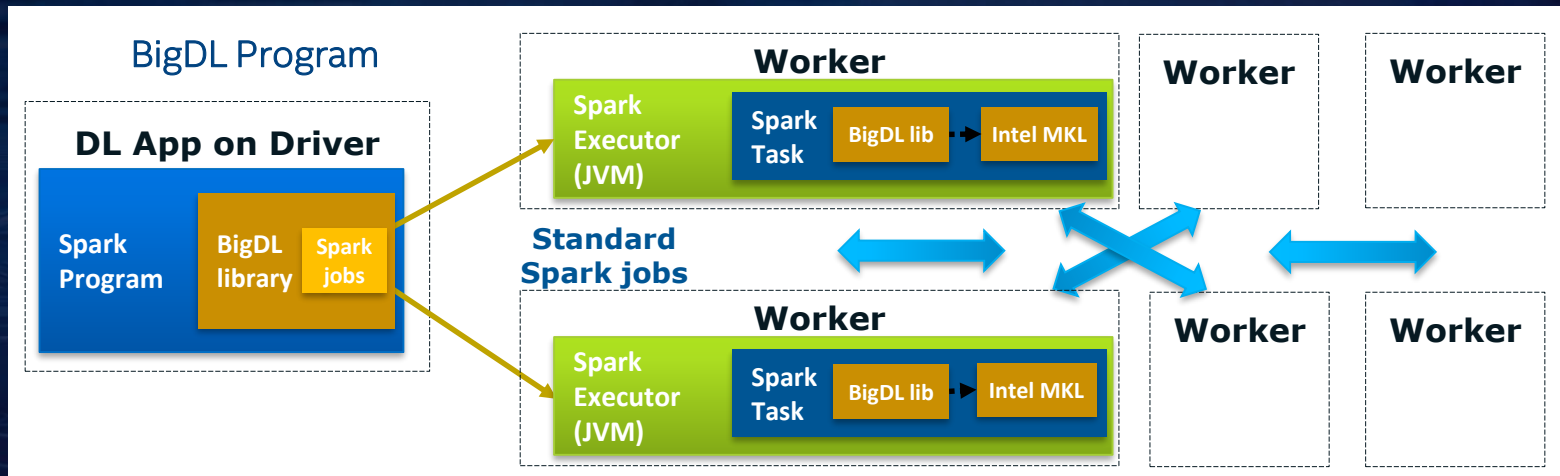
- No changes to the Spark or Hadoop clusters needed

Iterative

- Each iteration of the training runs as a Spark job

Data parallel

- Each Spark task runs the same model on a subset of the data (batch)



The background of the slide is a dark blue, high-contrast image of a computer circuit board. The board's intricate patterns of lines and components are visible, though slightly blurred. Overlaid on this background is a semi-transparent layer of white binary code (0s and 1s) arranged in horizontal lines, reminiscent of a digital rain or data stream effect. The overall aesthetic is technological and futuristic.

BigDL Execution Model

Distributed Training in BigDL


Data Parallel, Synchronous Mini-Batch SGD

```
Prepare training data as an RDD of Samples
Construct an RDD of models (each being a replica of the original model)

for (i <- 1 to N) {
  // "model forward-backward" job
  for each task in the Spark job:
    read the latest weights
    get a random batch of data from local Sample partition
    compute errors (forward on local model replica)
    compute gradients (backward on local model replica)

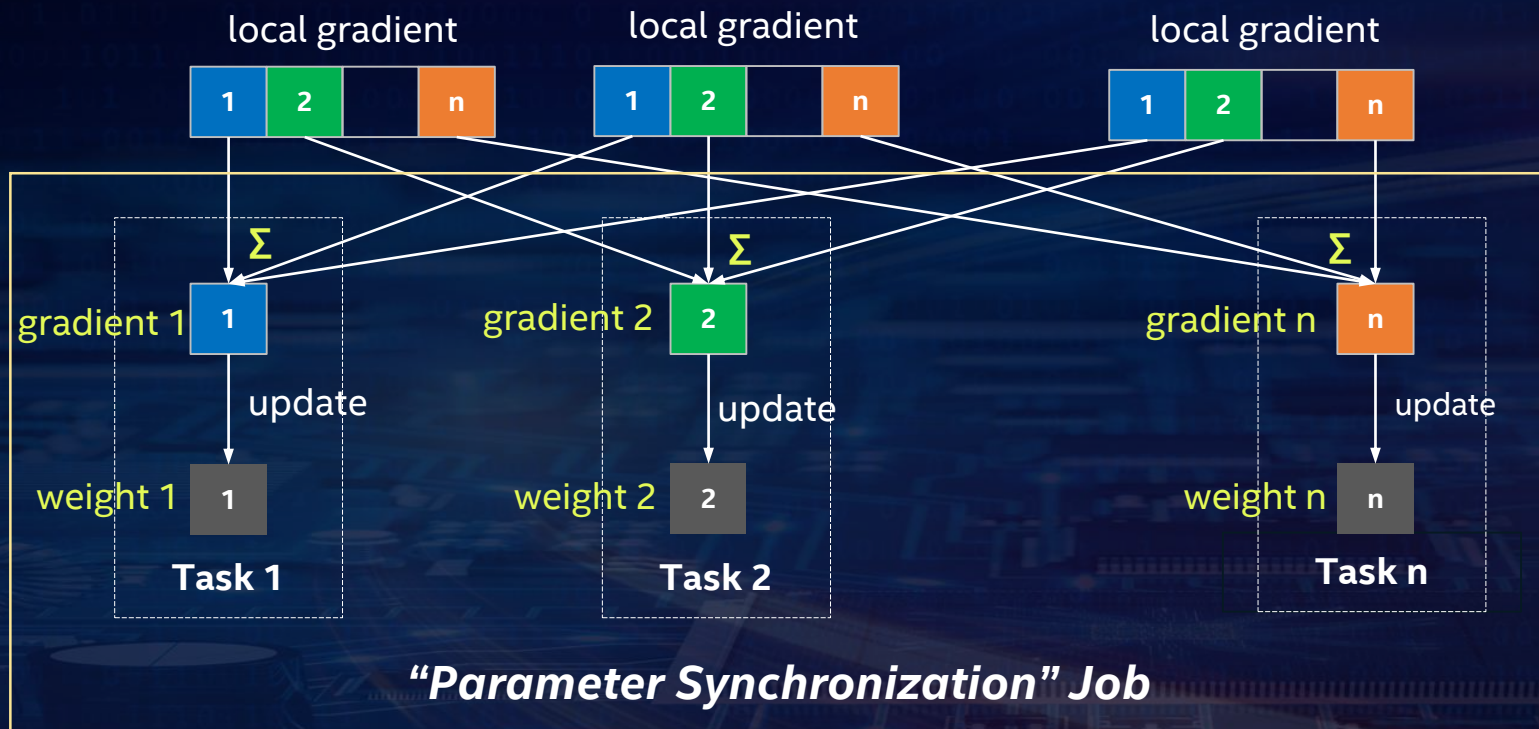
  // "parameter synchronization" job
  aggregate (sum) all the gradients
  update the weights per specified optimization method
}
```

Data Parallel Training



“Model Forward-Backward” Job

Parameter Synchronization



Parameter Synchronization

```
For each task  $n$  in the "parameter synchronization" job {  
  shuffle the  $n^{th}$  partition of all gradients to this task  
  aggregate (sum) the gradients  
  updates the  $n^{th}$  partition of the weights  
  broadcast the  $n^{th}$  partition of the updated weights  
}
```

"Parameter Synchronization" Job

(managing n^{th} partition of the parameters - similar to a parameter server)

AllReduce Operation (directly on top of primitives in Spark)

- Gradient aggregation: *shuffle*
- Weight sync: *task-side broadcast*
- *In-memory persistence*

Difference vs. Classical AllReduce

Classical AllReduce architecture


- Multiple long-running, potentially stateful tasks
- Interact with each other (in a blocking fashion for synchronization)
- Require fine-grained data access and in-place data mutation
- Not directly supported by existing big data systems

BigDL implementation

- Run a series of short-lived Spark jobs (e.g., two jobs per mini-batch)
- Each task in the job is stateless and non-blocking
- Automatically adapt to the dynamic resource changes (e.g., *preemption*, *failures*, *resource sharing*, etc.)
- Built on top of existing primitives in Spark (e.g., *shuffle*, *broadcast*, and *in-memory data persistence*)

Experimental Evaluation

Computing Performance




NCF training on single node:

- PyTorch 0.4 on Nvidia P100 GPU
- BigDL 0.7.0 and Spark 2.1.0 on a dual-socket Intel Skylake 8180 server (56 cores and 384GB)

The training performance of NCF using the BigDL implementation is 1.6x faster than the reference PyTorch implementation, as reported by MLPerf

MLPerf 0.5 training results URL: <https://mlperf.org/training-results-0-5>

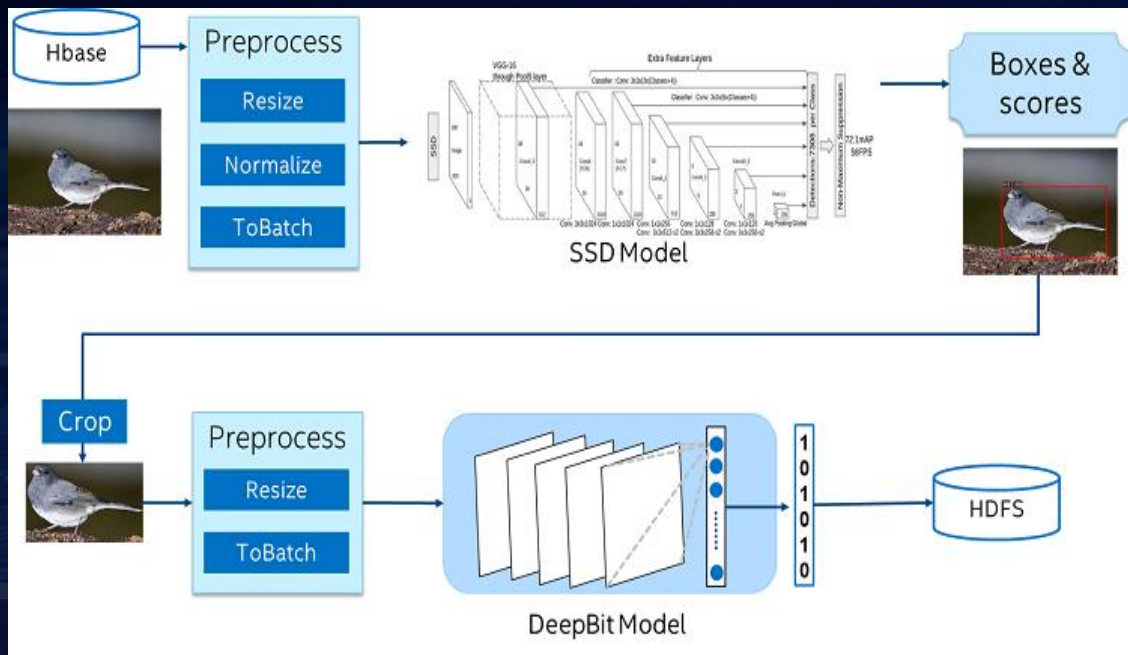
Training Scalability



Throughput of ImageNet Inception v1 training (w/ BigDL 0.3.0 and dual-socket Intel Broadwell 2.1 GHz); the throughput scales almost linear up to 128 nodes (and continue to scale reasonably up to 256 nodes).

Real-World Applications


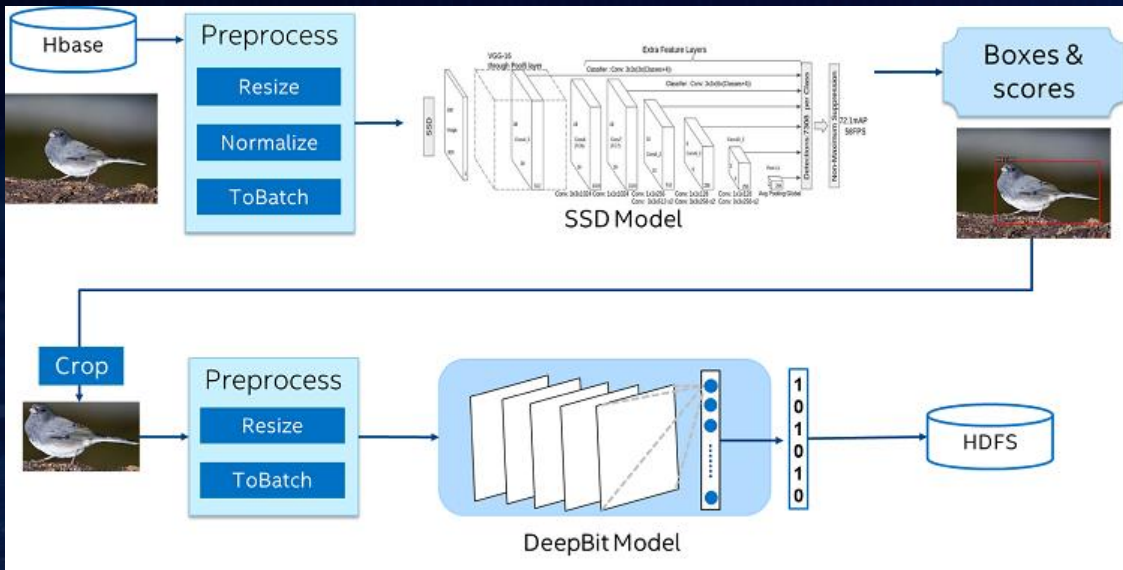
Object Detection and Image Feature Extraction at JD.com



Problem with previous "connector approach" (similar to CaffeOnSpark)

- Very complex and error-prone in managing large-scale distributed systems
- Impedance mismatch
 - Mismatch in the parallelism for data processing and for model compute

Object Detection and Image Feature Extraction at JD.com



- Implement the entire data analysis and deep learning pipeline under a **unified** programming paradigm on Spark
- Greatly improves the efficiency of **development** and **deployment**
- Efficiently scale out on Spark with superior performance (**3.83x** speed-up vs. GPU servers) as benchmarked by JD

<https://software.intel.com/en-us/articles/building-large-scale-image-feature-extraction-with-bigdl-at-jdcom>

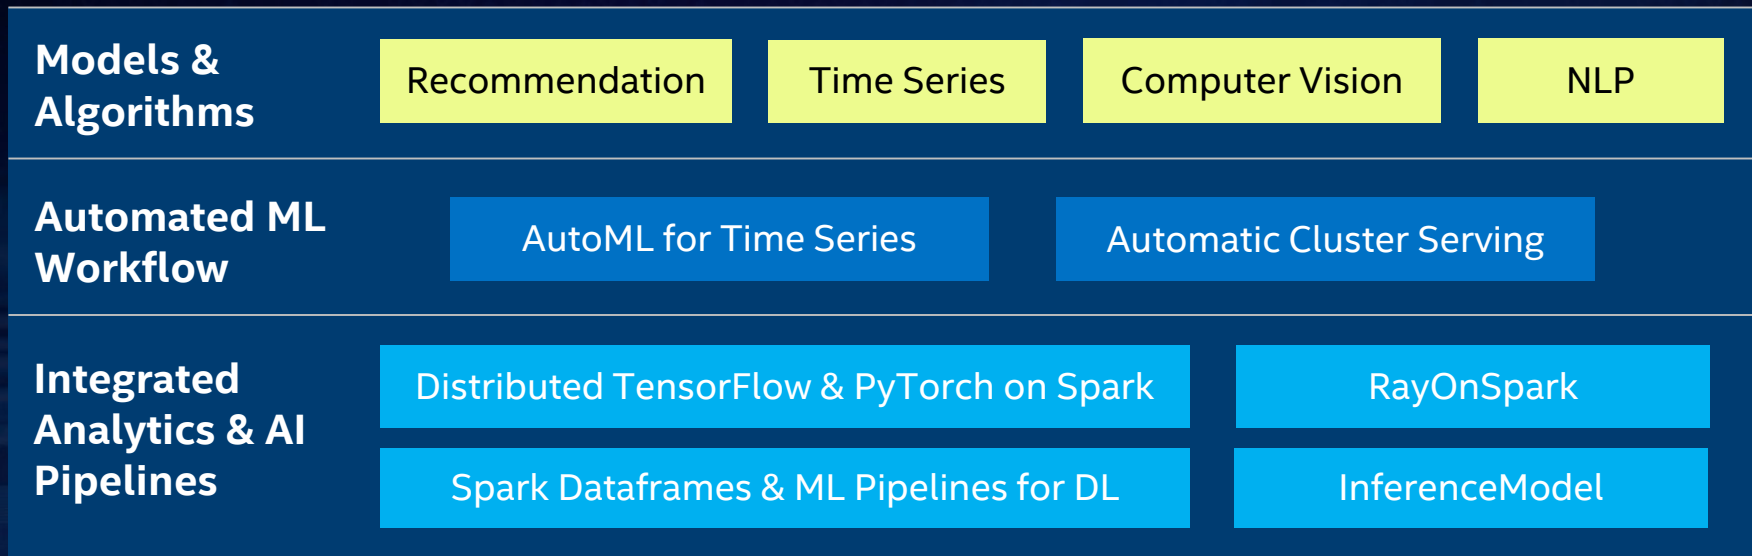
SOCC 2019



Future Work

Analytics Zoo

Unified Analytics + AI Platform for Big Data



Compute Environment

Laptop

K8s Cluster

Hadoop Cluster

Spark Cluster

DL Frameworks
(TF/PyTorch/OpenVINO/...)

Distributed Analytics
(Spark/Flink/Ray/...)

Python Libraries
(Numpy/Pandas/sklearn/...)

Powered by oneAPI

<https://github.com/intel-analytics/analytics-zoo>

SOCC 2019

The background is a dark blue, stylized image of a circuit board. It features various electronic components like capacitors and integrated circuits, with glowing lines representing circuit traces. Overlaid on this is a pattern of white binary code (0s and 1s) arranged in horizontal rows, creating a digital or data-themed aesthetic.

Q & A

What is Analytics Zoo



Distributed, High-Performance
Deep Learning Framework
for Apache Spark



<https://github.com/intel-analytics/bigdl>



Unified Analytics + AI Platform
Distributed TensorFlow, Keras, PyTorch and BigDL
on Apache Spark



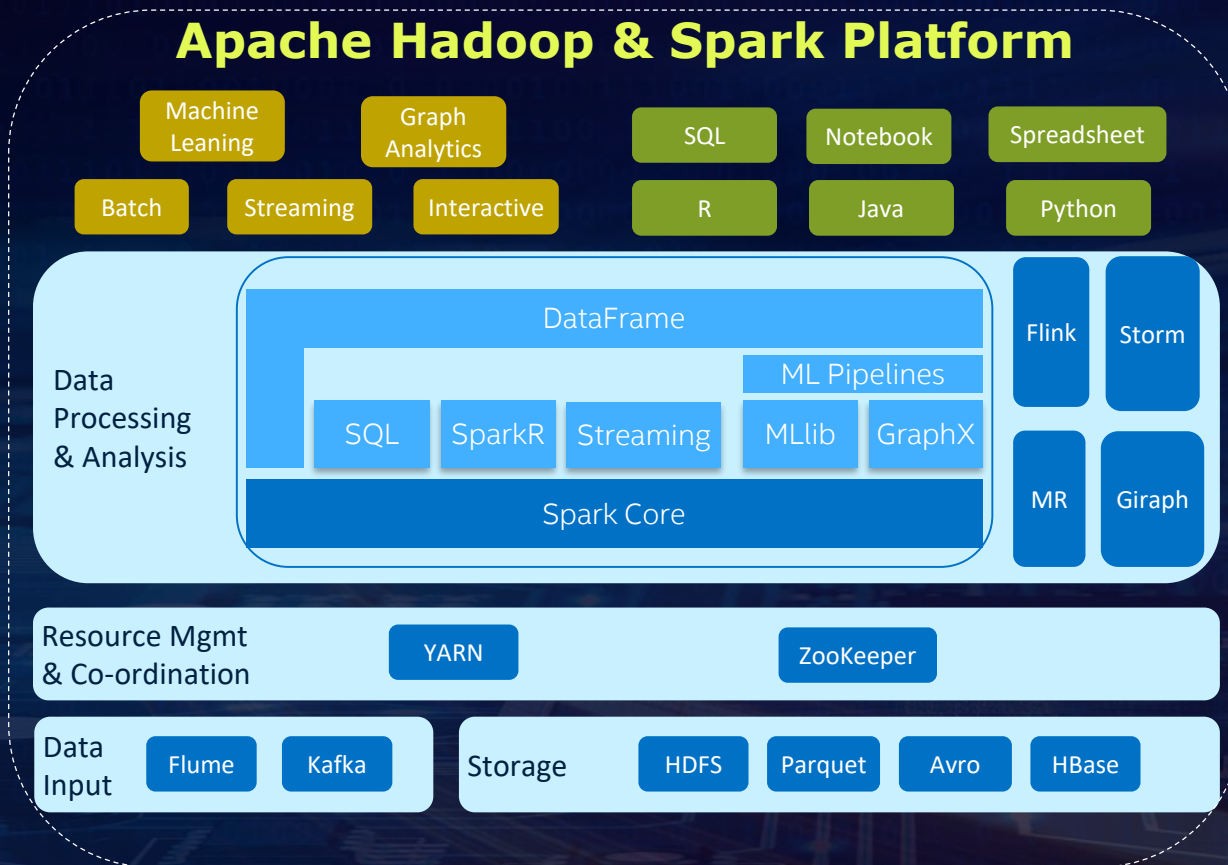
<https://github.com/intel-analytics/analytics-zoo>

Accelerating Data Analytics + AI Solutions At Scale

Appendix

Unified Big Data Analytics Platform

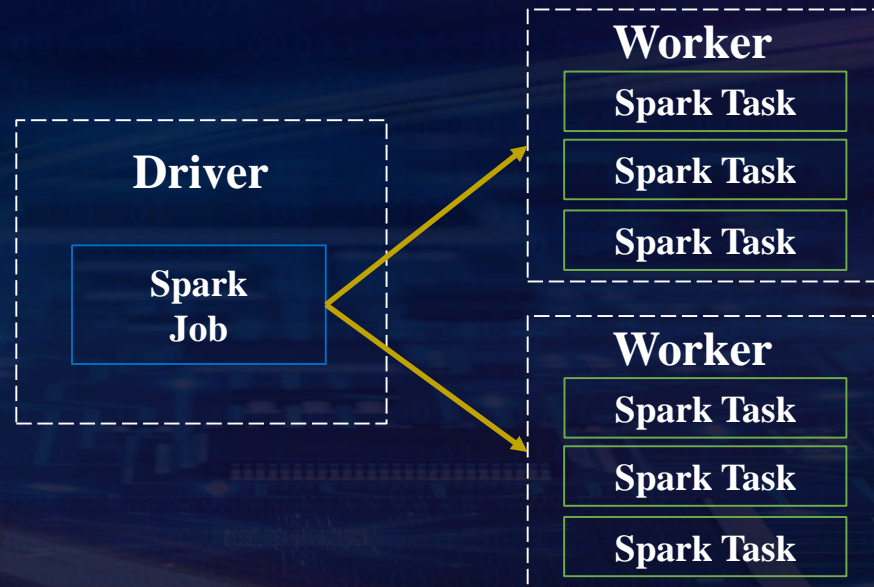
Apache Hadoop & Spark Platform




Apache Spark

Low Latency, Distributed Data Processing Framework

- A Spark cluster consists of a single *driver* node and multiple *worker* nodes
- A Spark *job* contains many Spark *tasks*, each working on a data *partition*
- Driver is responsible for scheduling and dispatching the tasks to workers, which runs the actual Spark tasks



Training Scalability




Overheads of parameter synchronization (as a fraction of average model computation time) of ImageNet Inception-v1 training in BigDL

Source: Scalable Deep Learning with BigDL on the Urika-XC Software Suite
(<https://www.cray.com/blog/scalable-deep-learning-bigdl-urika-xc-software-suite/>)

Reducing Scheduling Overheads Using Drizzle

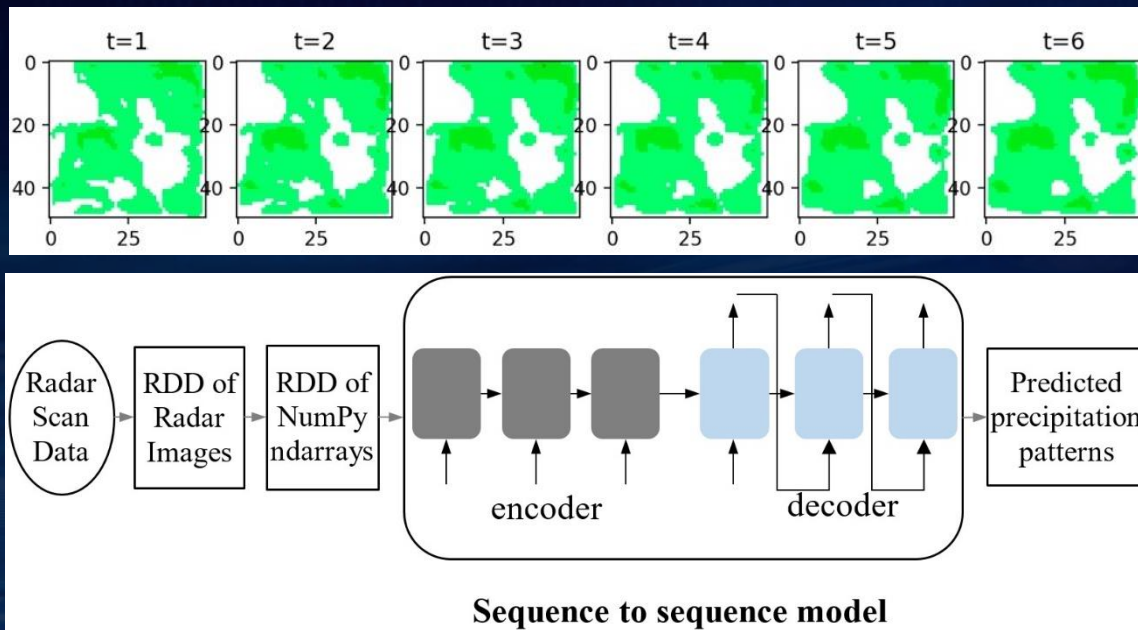
Scaling to even larger (>500) workers

- Iterative model training
 - Same operations run repeatedly
- Drizzle
 - A low latency execution engine for Spark
 - *Group scheduling* for multiple iterations of computations at once



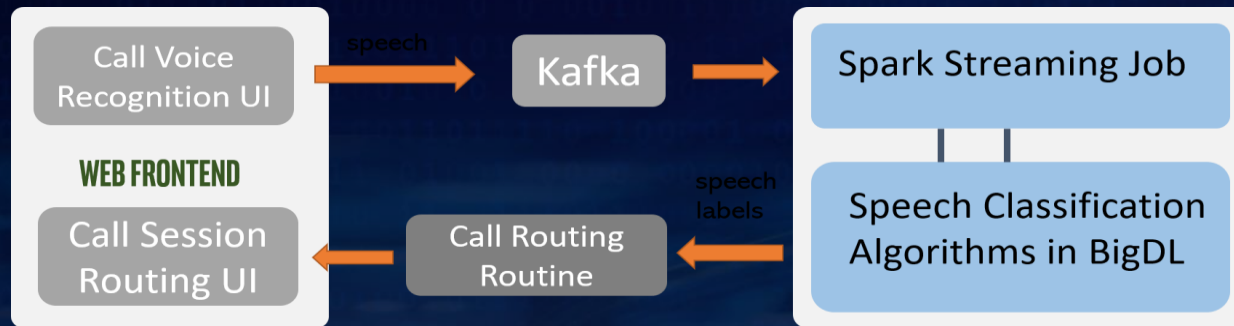
Source: Accelerating Deep Learning Training with BigDL and Drizzle on Apache Spark, Shivaram Venkataraman, Ding Ding, and Sergey Ermolin.
(<https://rise.cs.berkeley.edu/blog/accelerating-deep-learning-training-with-bigdl-and-drizzle-on-apache-spark/>)

Precipitation nowcasting using sequence-to-sequence models in **Cray**



- Running data processing on a Spark cluster, and deep learning training on GPU cluster not only brings high data movement overheads, but hurts the development productivity due to the fragmented workflow
- Using a single unified data analysis and deep learning pipeline on Spark and BigDL improves the efficiency of development and deployment


Real-time streaming speech classification in GigaSpaces



The end-to-end workflow of real-time streaming speech classification on Kafka, Spark Streaming and BigDL in GigaSpaces.

- BigDL allows neural network models to be directly applied in standard distributed streaming architecture for Big Data (using Apache Kafka and Spark Streaming), and efficiently scales out to a large number of nodes in a transparent fashion.

<https://www.gigaspace.com/blog/gigaspace-to-demo-with-intel-at-strata-data-conference-and-microsoft-ignite/>



LEGAL DISCLAIMERS

- Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Learn more at intel.com, or from the OEM or retailer.
- No computer system can be absolutely secure.
- Tests document performance of components on a particular test, in specific systems. Differences in hardware, software, or configuration will affect actual performance. Consult other sources of information to evaluate performance as you consider your purchase. For more complete information about performance and benchmark results, visit <http://www.intel.com/performance>.

Intel, the Intel logo, Xeon, Xeon phi, Lake Crest, etc. are trademarks of Intel Corporation in the U.S. and/or other countries.

*Other names and brands may be claimed as the property of others.

© 2019 Intel Corporation