What is Analytics Zoo

BigDL

Distributed, High-Performance

Deep Learning Framework

for Apache Spark

https://github.com/intel-analytics/bigdl

ANALYTICS ZOO

Unified Analytics + AI Platform

Distributed TensorFlow, Keras, PyTorch and BigDL on Apache Spark

https://github.com/intel-analytics/analytics-zoo

Accelerating Data Analytics + AI Solutions At Scale


intel
ANALYTICS ZOO
Cluster Serving
2020/06/12

Overview

What's Analytics Zoo

Analytics + AI Platform

Distributed TensorFlow*, Keras*,
PyTorch* and BigDL on Apache Spark*

https://github.com/intel-analytics/analytics-zoo

Accelerating Data Analytics + AI Solutions At Scale

Machine Learning VS Deep Learning

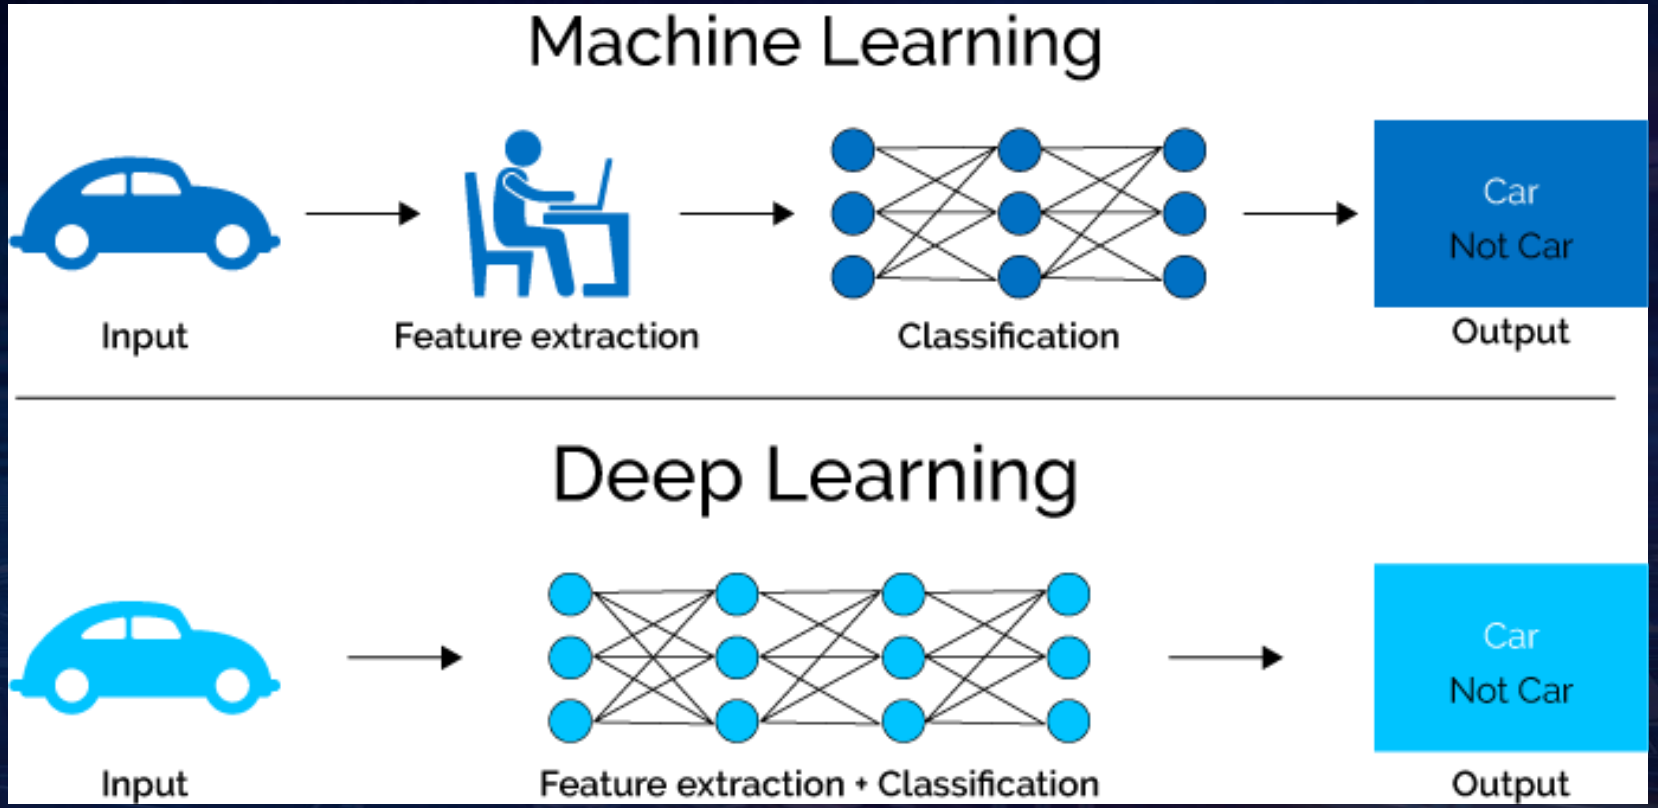

Feature Visualization

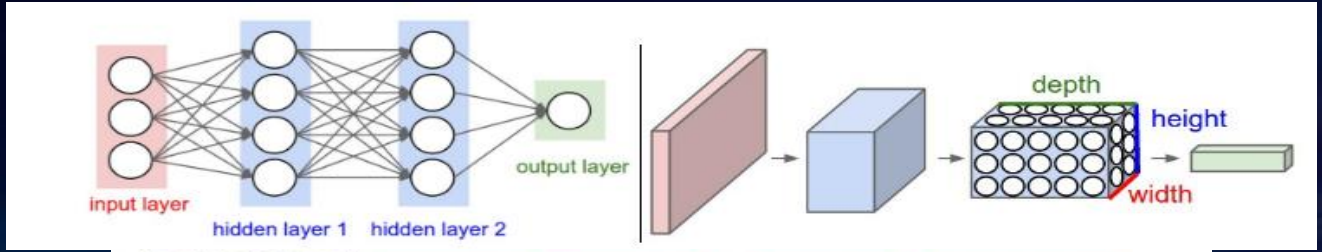

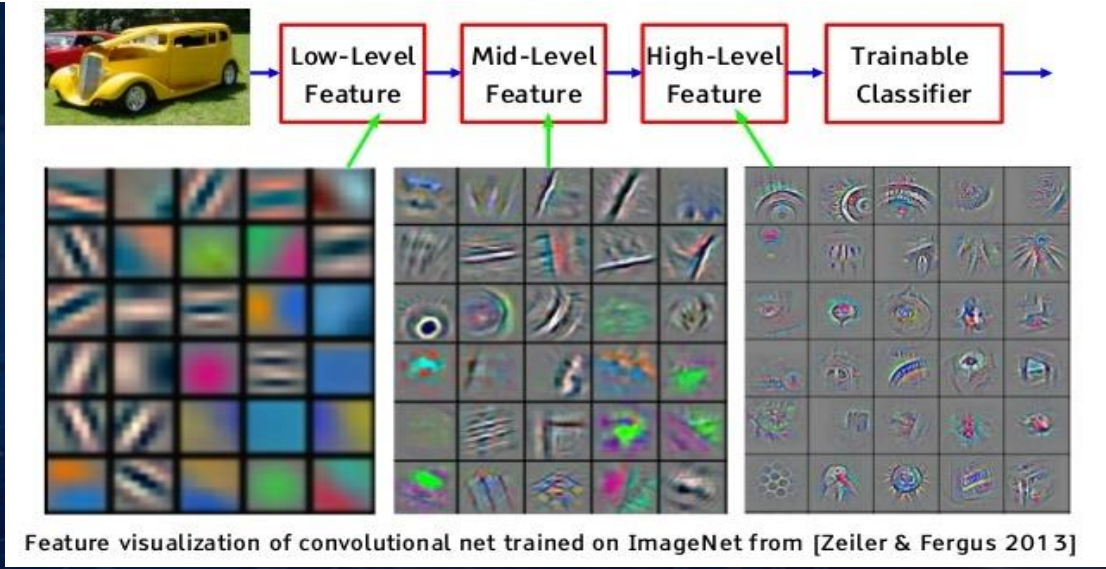

Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Data & Performance Relationship

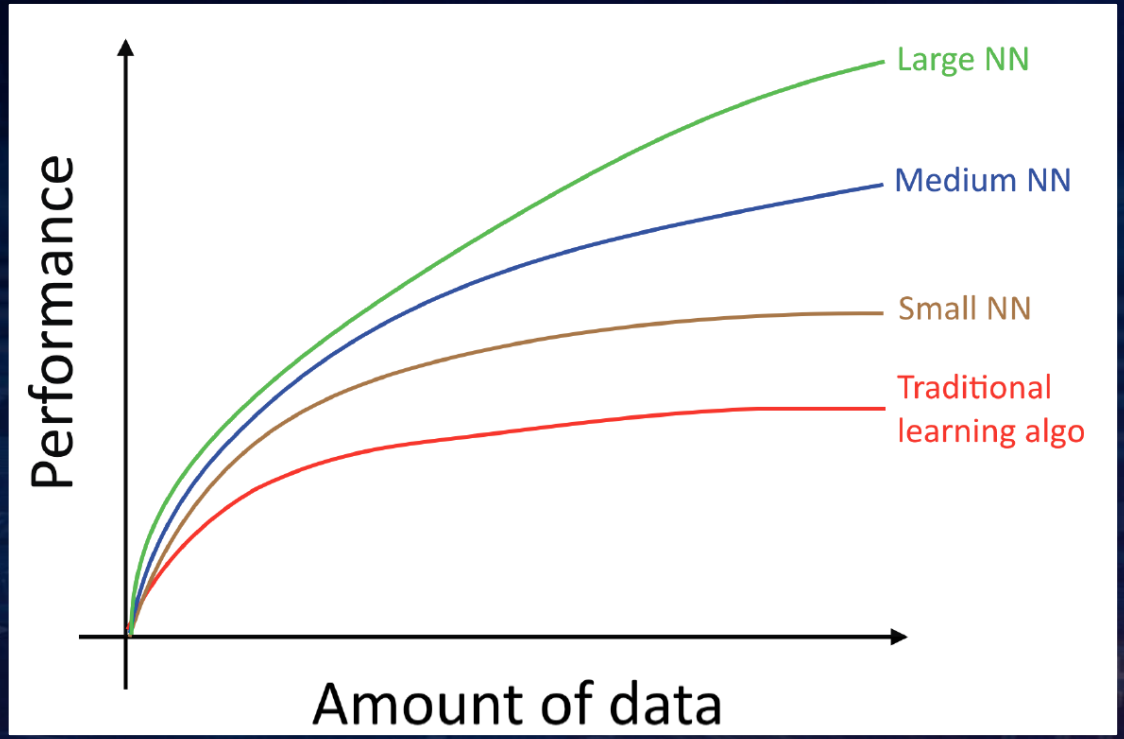

“Machine Learning Yearning”, Andrew Ng, 2016

Real-World ML/DL Applications Are Complex Data Analytics Pipelines

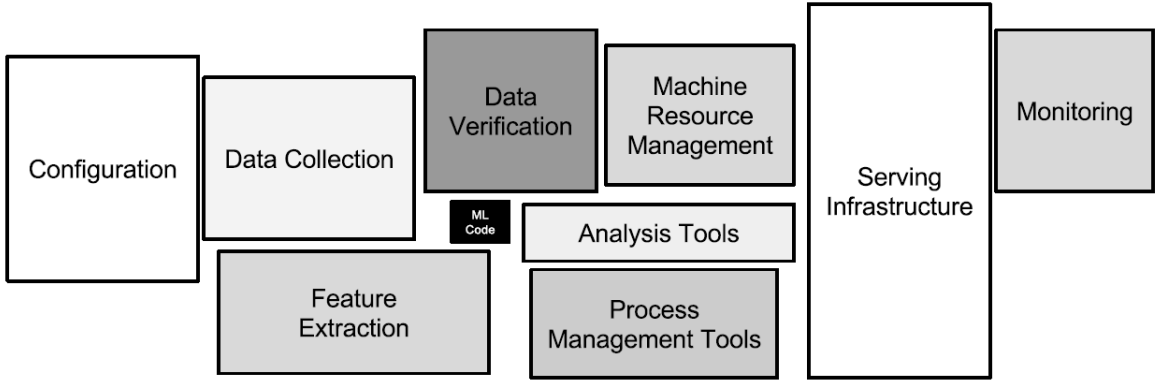

Figure 1: Only a small fraction of real-world ML systems is composed of the ML code, as shown by the small black box in the middle. The required surrounding infrastructure is vast and complex.

"Hidden Technical Debt in Machine Learning Systems", Sculley et al., Google, NIPS 2015 Paper

Data Analytics Pipeline from production perspective

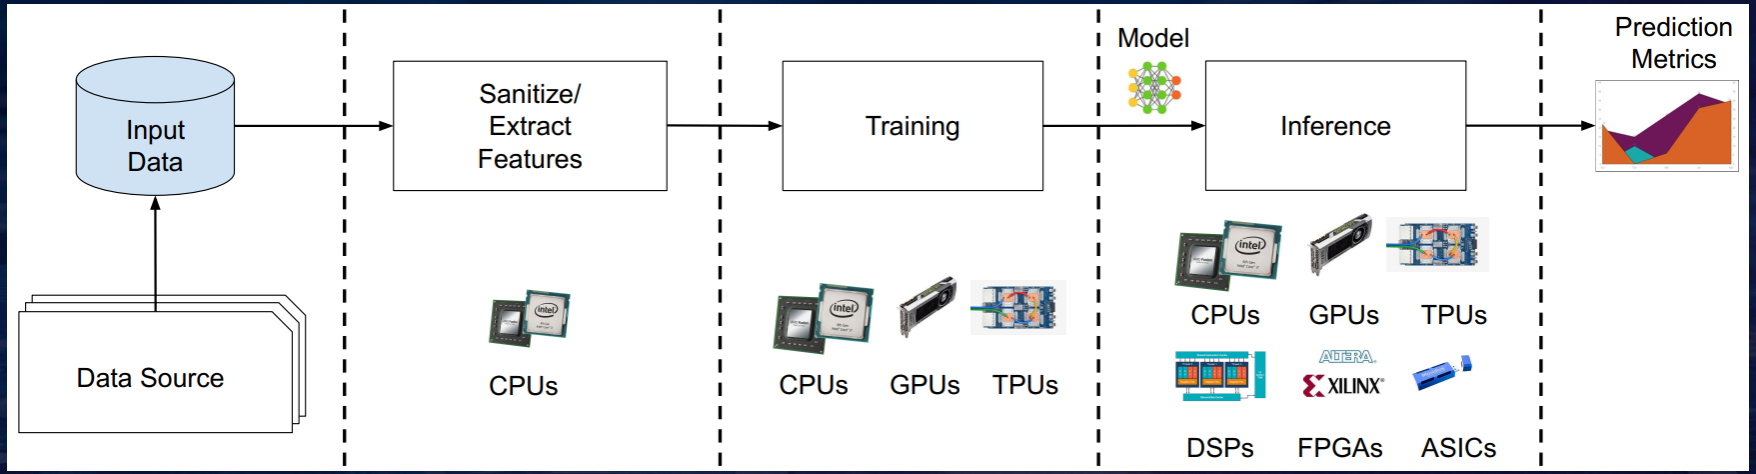

End-to-End Big Data Analytics and AI Pipeline

Seamless Scaling from Laptop to Production with ANALYTICS ZOO

- **"Zero" code change from laptop to distributed cluster**
- **Directly access production data (Hadoop/Hive/HBase) without data copy**
- **Easily prototype the end-to-end pipeline**
- **Seamlessly deployed on production big data clusters**

Analytics Zoo

Unified Analytics + AI Platform for Big Data

Use case

- Recommendation
- Anomaly Detection
- Text Classification
- Text Matching

Model

- Image Classification
- Object Detection
- Seq2Seq
- Transformer
- BERT

Feature Engineering

- image
- 3D image
- text
- Time series

High Level Pipelines

- tfpark: Distributed TF on Spark
- Distributed Keras w/ autograd on Spark
- nnframes: Spark Dataframes & ML Pipelines for Deep Learning
- Distributed Model Serving (batch, streaming & online)

Backend/ Library

- TensorFlow
- Keras
- BigDL
- NLP Architect
- Apache Spark
- Apache Flink
- MKLDNN
- OpenVINO
- Intel® Optane™ DCPMM
- DL Boost (VNNI)

https://github.com/intel-analytics/analytics-zoo

What's Analytics Zoo

Analytics + AI Platform

Distributed TensorFlow*, Keras*,
PyTorch* and BigDL on Apache Spark*

https://github.com/intel-analytics/analytics-zoo

Accelerating Data Analytics + AI Solutions At Scale

What's Serving

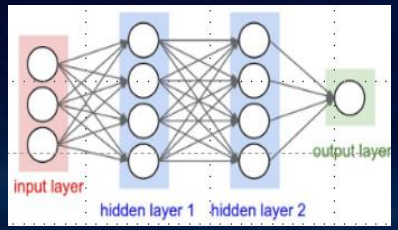

model

Input Data → Preprocessing → Predict → Postprocessing → Result

Example of TFServing

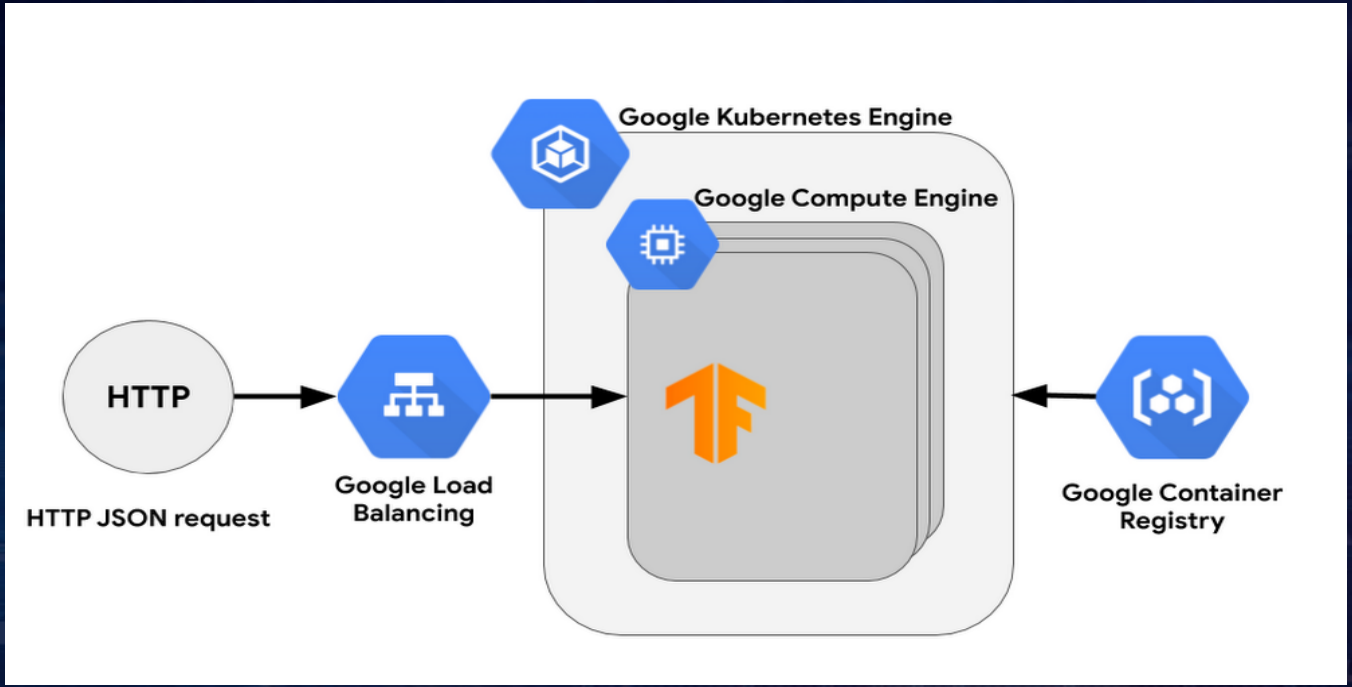

Distributed Model Serving

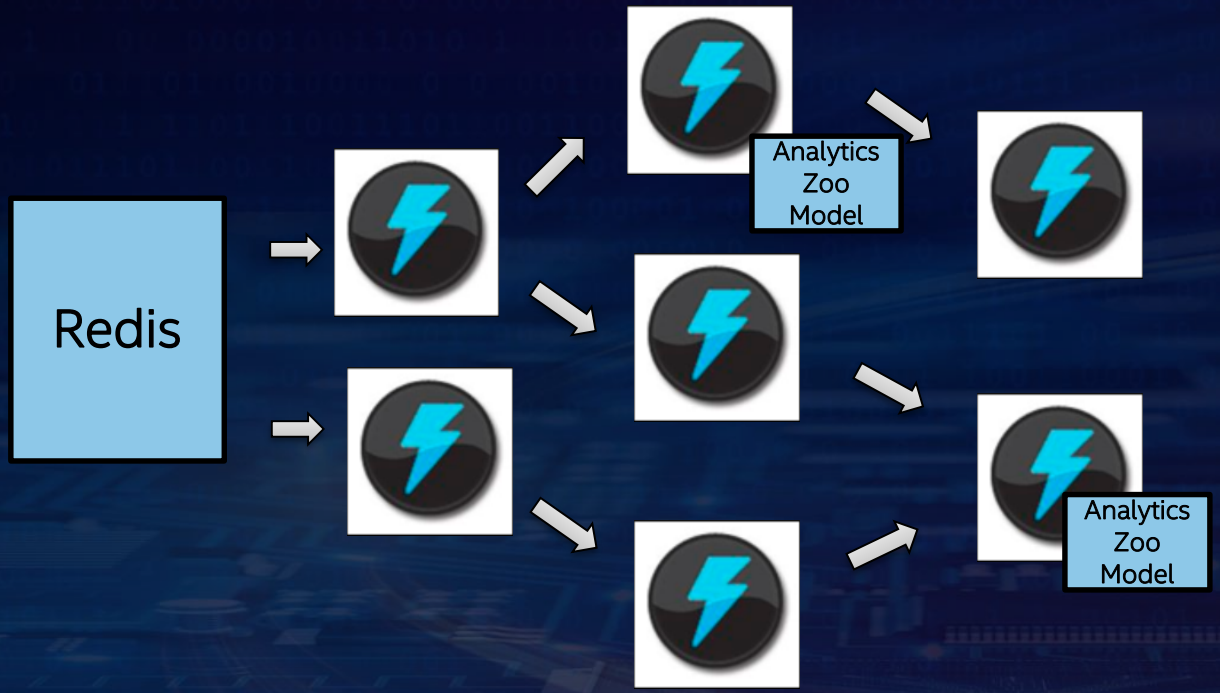

Distributed model serving in Web Service, Flink, Kafka, Storm, etc.

- **Plain Java or Python API, with OpenVINO and DL Boost (VNNI) support**

Main Version of Cluster Serving

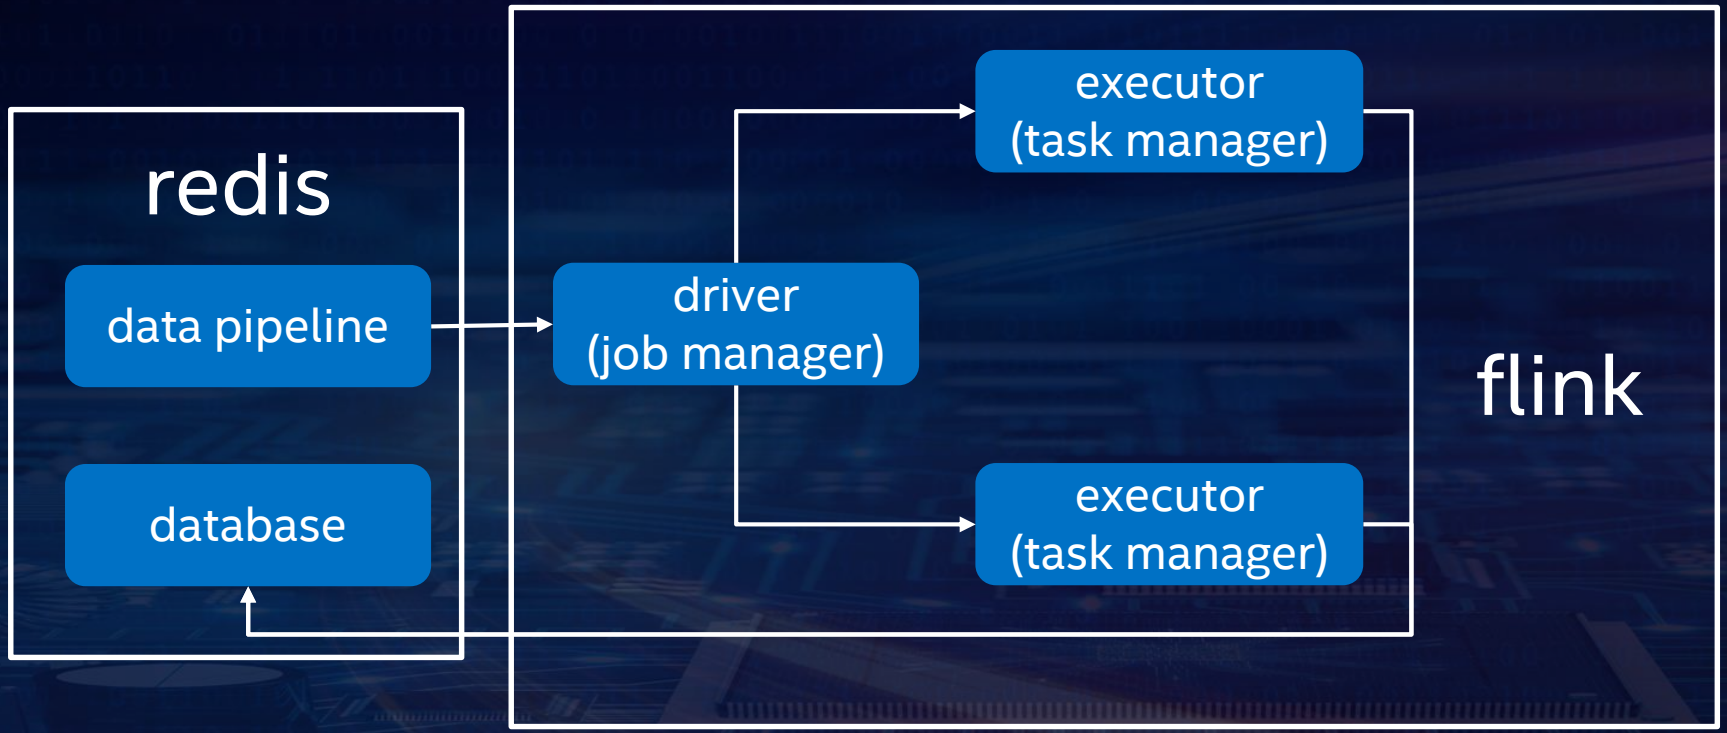

Version based on Spark Streaming is also supported.

Data pipeline User Perspective

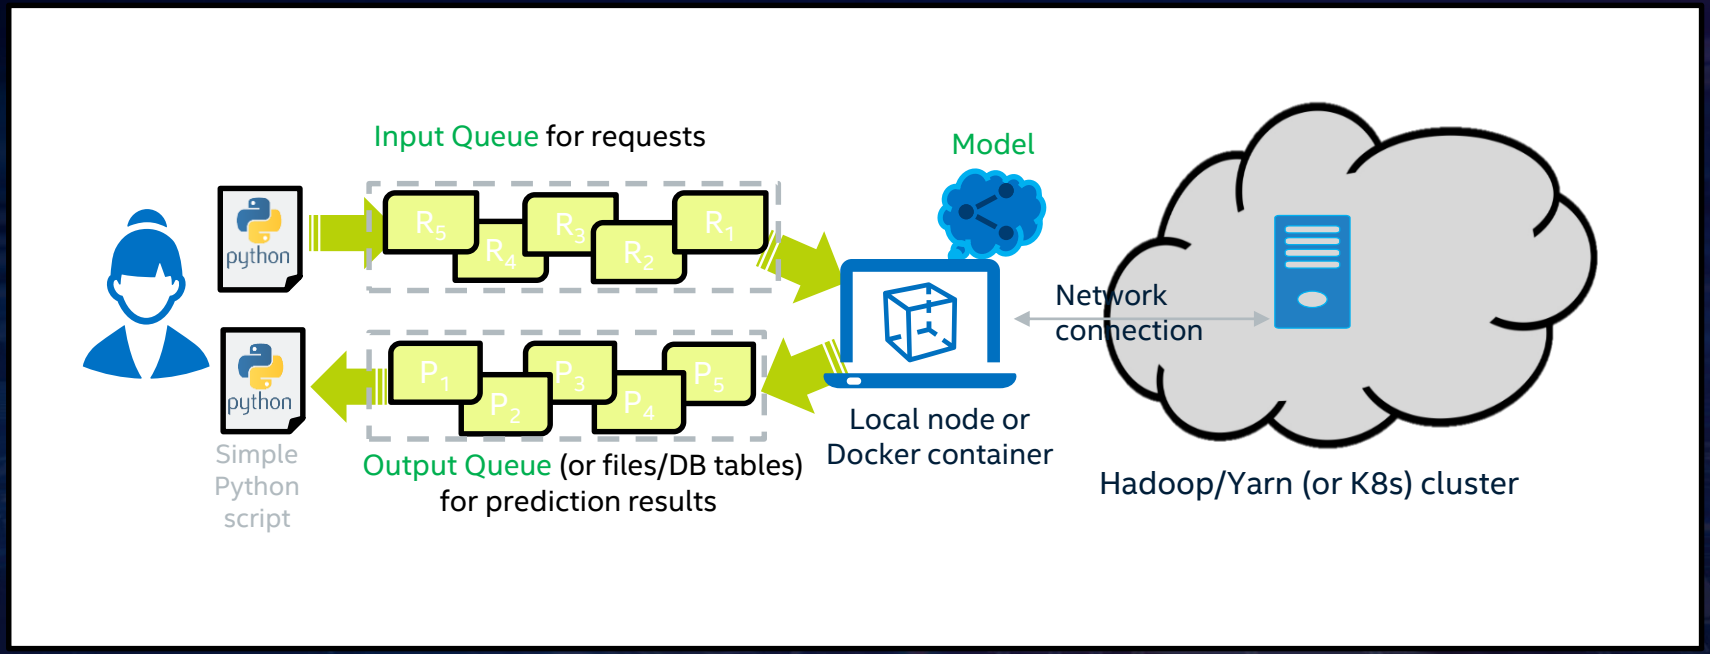

Deploy Your Own Cluster Serving

One command to pull docker image, and customize your config, then call cluster-serving-start to start your serving

example of config:

```
## Analytics-zoo Cluster Serving
model:
  # model path must be set
  path: resources
data:
  # redis address
  src: XXXXXX:6379

…
```

API Introductions

http API

data are represented by json format, and call http post method to enqueue your data into pipeline
(http API is compatible with TFServing)

python API

data are represented by ndarray, and call python method to enqueue your data into pipeline

Use Case – Medical Imaging Analysis

Consider a Very large medical image of patient, the mission is to determine if tumor exists

End-to-end pipeline would contain image preprocessing, predict, all reduce analysis

tumor → 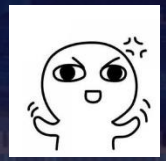 → if num > 1% the condition is bad →

Advantages

Wide Range Deep Learning model support

Tensorflow, Caffe, OpenVINO, Pytorch, BigDL

Low Latency

Continuous Streaming pipeline is supported by Apache Flink, also Spark version is supported for users who are more familiar.

High Throughput & Scalability

Optimization of multithread control, and could easily scale out to clusters.

Very Quick Start

```
docker run -itd --name cluster-serving --net=host intelanalytics/zoo-cluster-serving:0.7.0
```

Log into the container using `docker exec -it cluster-serving bash`.

We already prepared `analytics-zoo` and `opencv-python` with pip in this container. And prepared model in `model` directory with following structure.

```
cluster-serving |
                -- | model
                  -- frozen_graph.pb
                  -- graph_meta.json
```

Start Cluster Serving using `cluster-serving-start`.

Run python program `python quick_start.py` to push data into queue and get inference result.

Then you can see the inference output in console.

```
image: fish1.jpeg, classification-result:class: 5's prob: 0.18204997
image: dog1.jpeg, classification-result:class: 267's prob: 0.27166227
image: cat1.jpeg, classification-result:class: 292's prob: 0.32633427
```

End-to-End Big Data and AI Pipelines

Seamless Scaling from Laptop to Production

Unified Analytics + AI Platform

Distributed TensorFlow*, Keras*, PyTorch* & BigDL on Apache Spark*

https://github.com/intel-analytics/analytics-zoo

https://github.com/intel-analytics/analytics-zoo/blob/master/docs/docs/ClusterServingGuide/ProgrammingGuide.md

Analytics-Zoo

rm

Distributed Te DL on Apache Spark*

https://

*Other names and brands may be claimed as the property of others.

What is Analytics Zoo

BigDL

Distributed, High-Performance

Deep Learning Framework

for Apache Spark

https://github.com/intel-analytics/bigdl

ANALYTICS ZOO

Unified Analytics + AI Platform

Distributed TensorFlow, Keras, PyTorch and BigDL on Apache Spark

https://github.com/intel-analytics/analytics-zoo

Accelerating Data Analytics + AI Solutions At Scale

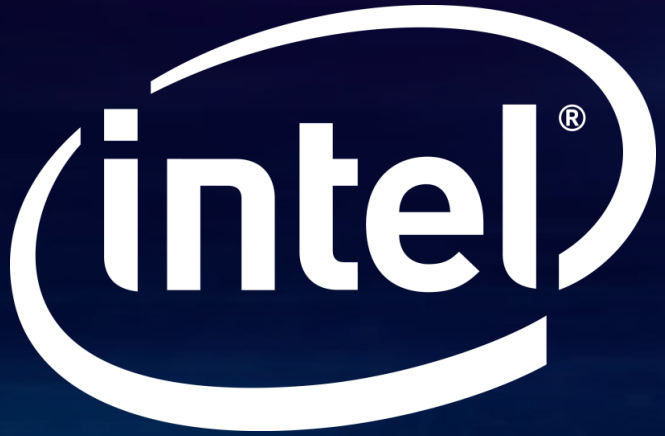
intel®

LEGAL DISCLAIMERS

- Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Learn more at intel.com, or from the OEM or retailer.
- No computer system can be absolutely secure.
- Tests document performance of components on a particular test, in specific systems. Differences in hardware, software, or configuration will affect actual performance. Consult other sources of information to evaluate performance as you consider your purchase. For more complete information about performance and benchmark results, visit **http://www.intel.com/performance**.

What is Analytics Zoo

BigDL

Distributed, High-Performance

Deep Learning Framework

for Apache Spark

https://github.com/intel-analytics/bigdl

ANALYTICS ZOO

Unified Analytics + AI Platform

Distributed TensorFlow, Keras, PyTorch and BigDL on Apache Spark

https://github.com/intel-analytics/analytics-zoo

Accelerating Data Analytics + AI Solutions At Scale