

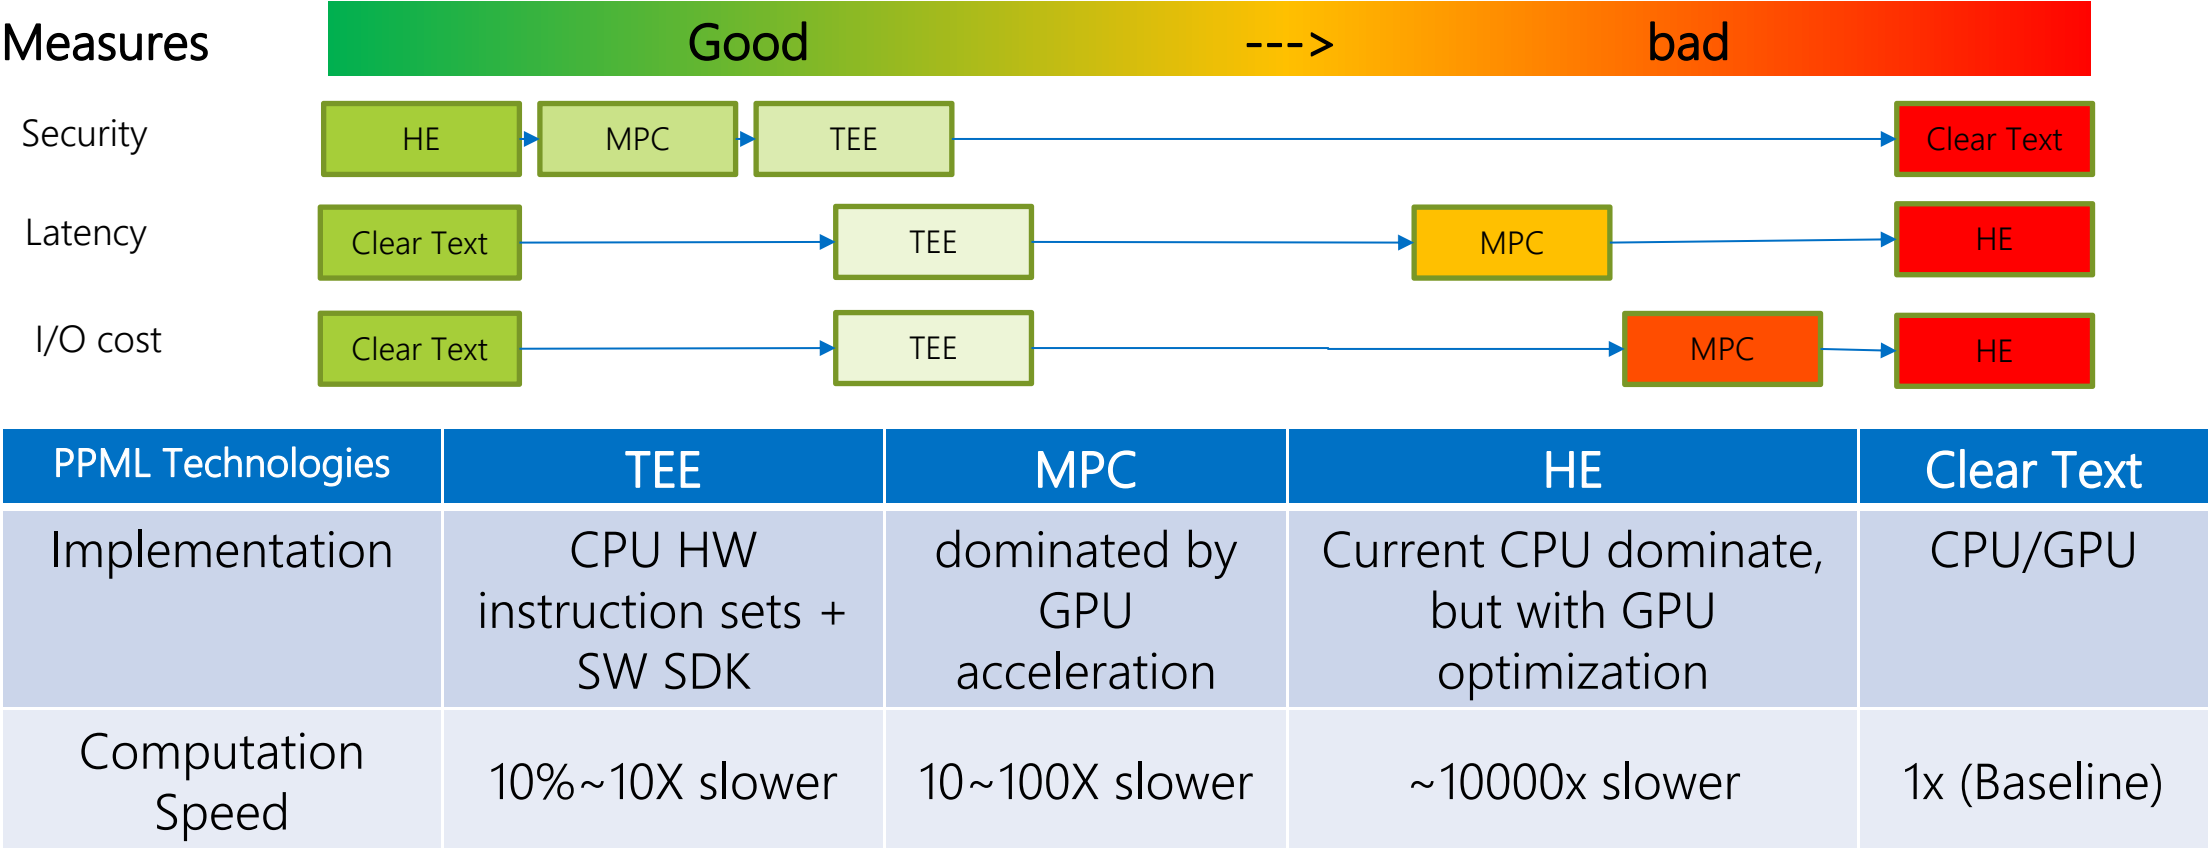
# Analytics Zoo and PPML

Dongjie Shi

# Agenda

- TEE
- Intel SGX
- Graphene-SGX
- Analytics Zoo
- Analytics Zoo Secured Cluster Serving on Graphene-SGX

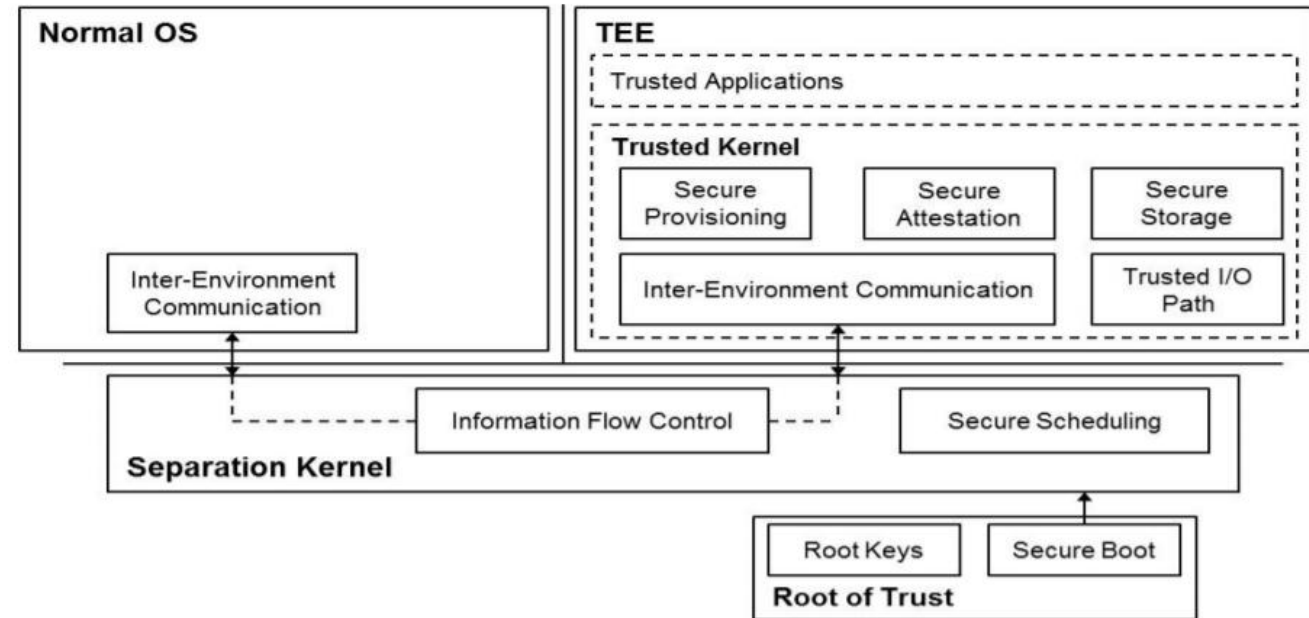
# Comparison of PPML Technologies



\*Other names and brands may be claimed as the property of others.

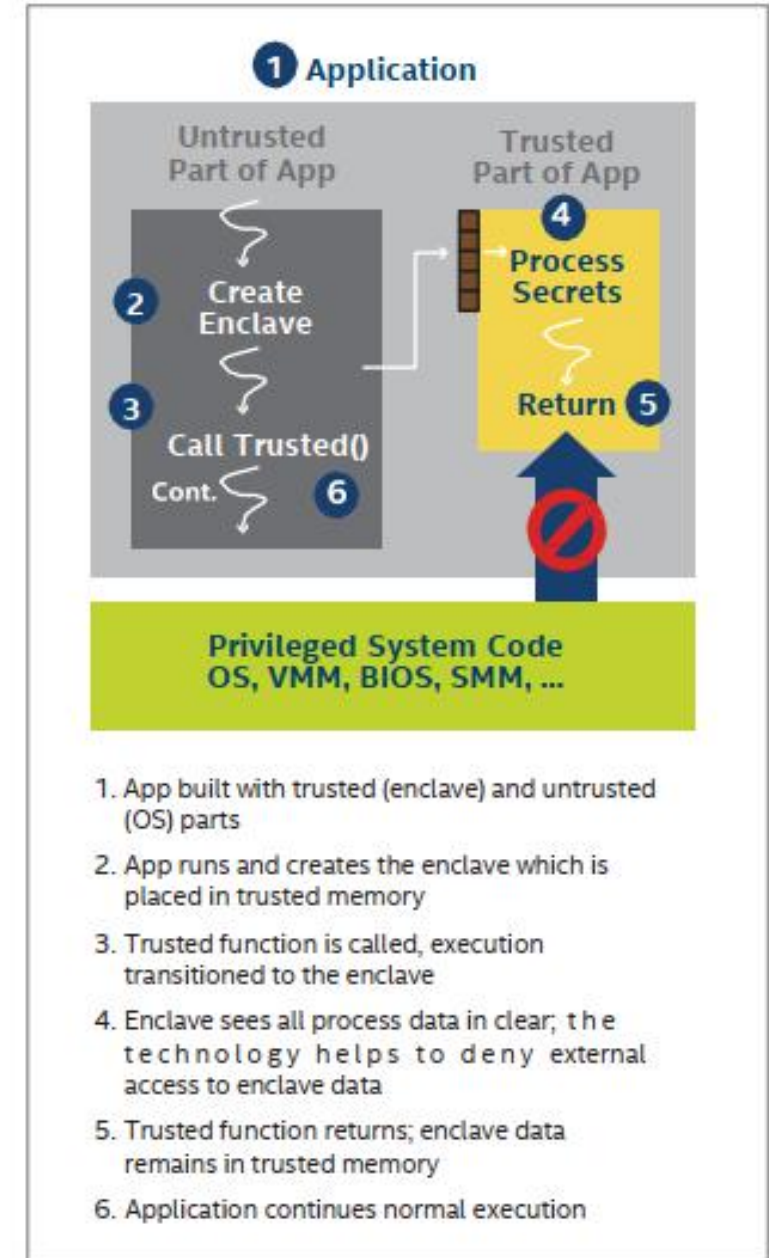
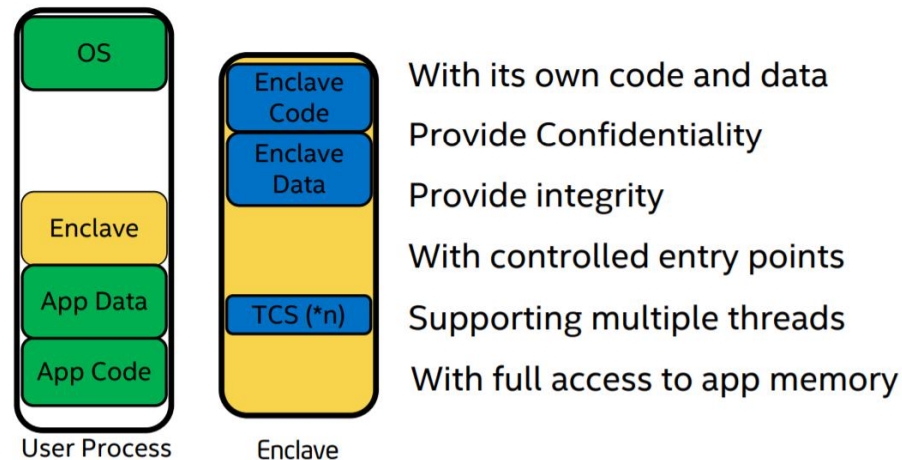
# TEE: Trusted Execution Environment

- TEE is a tamper resistant processing environment that runs on a separation kernel.
- Goals of TEE
  - Isolated Execution
    - TEE/Normal OS may be malicious
  - Secure Storage
    - Integrity, Confidentiality, Freshness
  - Remote Attestation
    - determine the level of trust in the integrity of attestator
  - Secure Provisioning
    - remotely manage and update its data in a secure way
  - Trusted I/O Path
    - protects authenticity, and optionally confidentiality, of communication between TEE and peripherals



# Intel® SGX (Intel® Software Guard Extensions)

- SGX protects selected code and data from disclosure or modification.
  - Enhances confidentiality and integrity
  - Low learning curve
  - Remotely attest and provision
  - Help Significantly reduce attack surface
- The primary SGX abstraction is an enclave: an isolated execution environment within the virtual address space of a process.



# Graphene-SGX:

## A Practical Library OS for **Unmodified** Applications on SGX

### Case Studies



Apache Flink

Big Data



TensorFlow

OpenVINO™

AI



Big Data + AI

### Properties

- Mature, widely deployed code
- Complex with millions of lines of code
- Written in C, Python, Java
- Use OS services

### Refactoring Challenges

- Takes time and effort
- May introduce bugs in mature code
- Need a separate SDK for every language
- Developer expertise to know how to refactor
- No access to source code of 3rd party apps

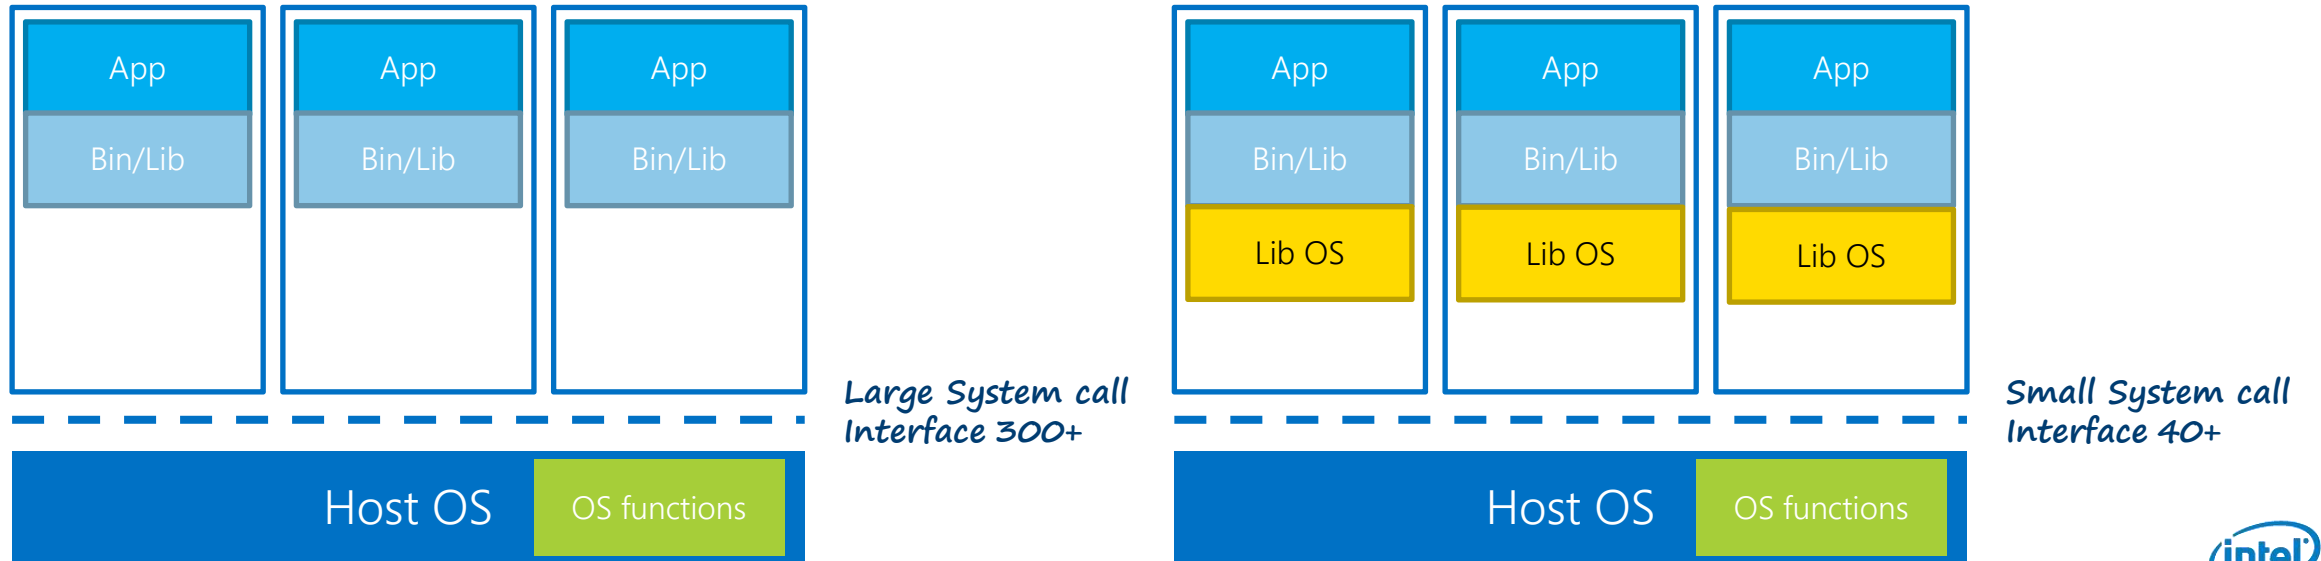


# Graphene-SGX: Objectives & Solution

- Objectives
  - Develop a tool to secure an application that:
    - Provides isolation without any code modification
    - Supports policy-based security enforcement (Manifest, trusted, allowed)
    - Does not compromise performance
    - Provides transparent support for attestation and secure OS services
- Solution
  - **Graphene:** A Library OS for running unmodified Linux applications inside SGX enclaves

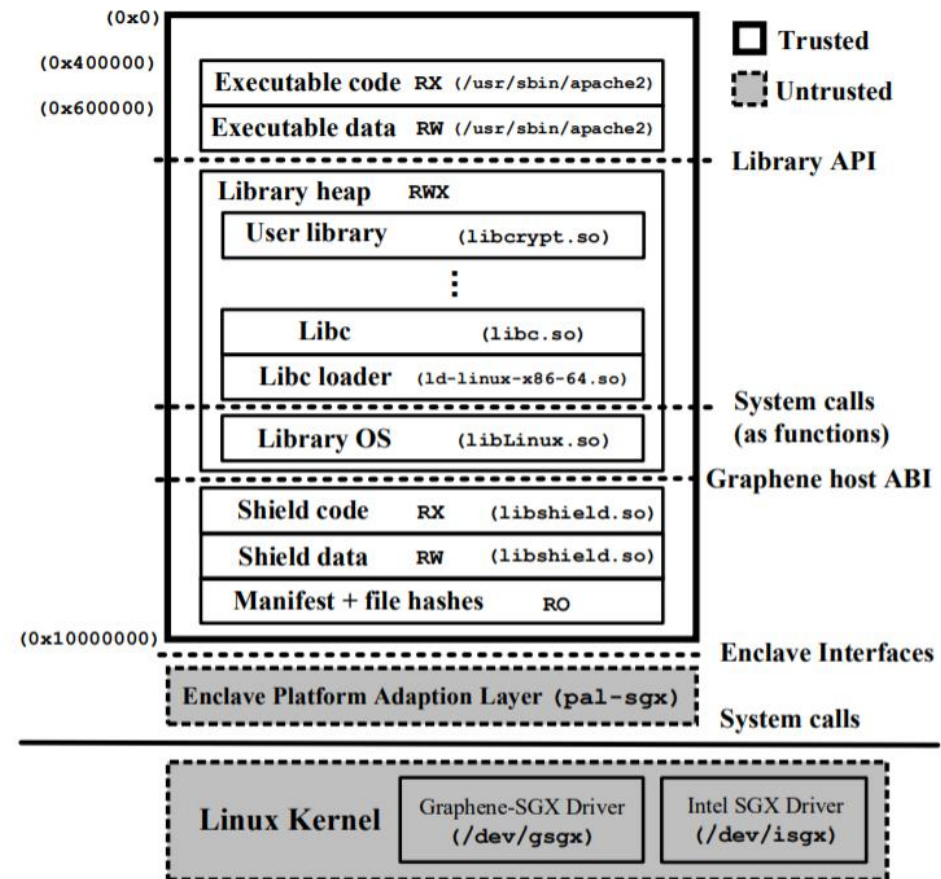
# Library OS Background

- Approach was championed by several OS designs in early 90's with a focus on performance
- Runs a part of OS functionality as a library in application
- Communication to the host OS with a small fixed set of abstractions
- Offers a secure alternative to large traditional OS system call interface





# Graphene-SGX Architecture



# AI on Big Data



Distributed, High-Performance  
**Deep Learning Framework**  
for Apache Spark\*

<https://github.com/intel-analytics/bigdl>



**Unified Analytics + AI Platform**  
for TensorFlow\*, PyTorch\*, Keras\*, BigDL,  
Ray\* and Apache Spark\*

<https://github.com/intel-analytics/analytics-zoo>

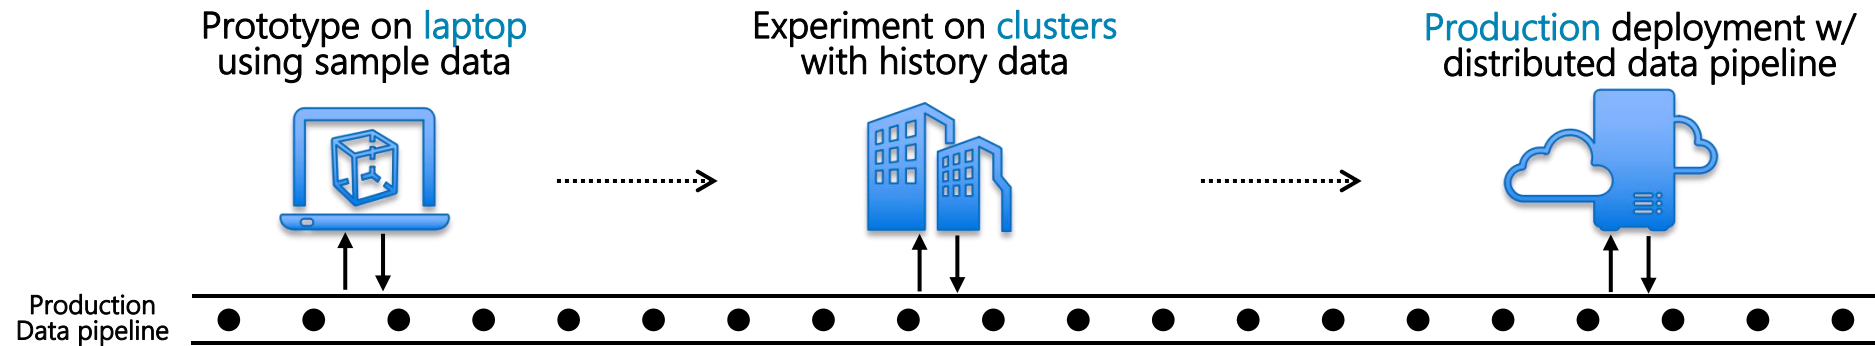
**Accelerating Data Analytics + AI Solutions At Scale**



Software

# Integrated Big Data Analytics and AI

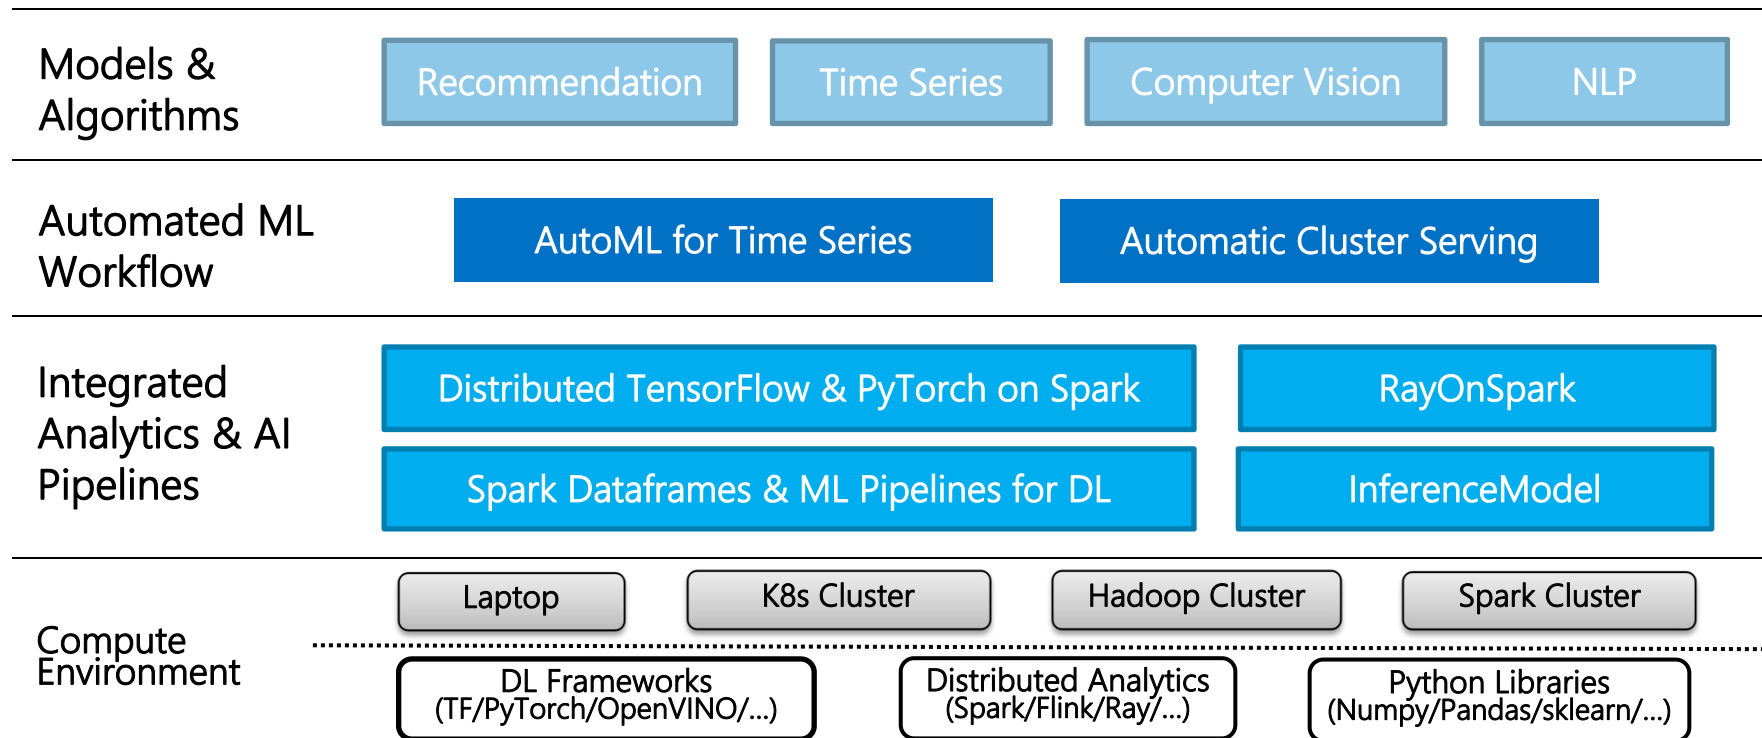
## Seamless Scaling from Laptop to Distributed Big Data



- Easily prototype **end-to-end pipelines** that apply AI models to big data
- **"Zero" code change** from laptop to distributed cluster
- Seamlessly deployed on **production Hadoop/K8s clusters**
- **Automate the process** of applying machine learning to big data

# Analytics Zoo

## Unified Data Analytics and AI Platform

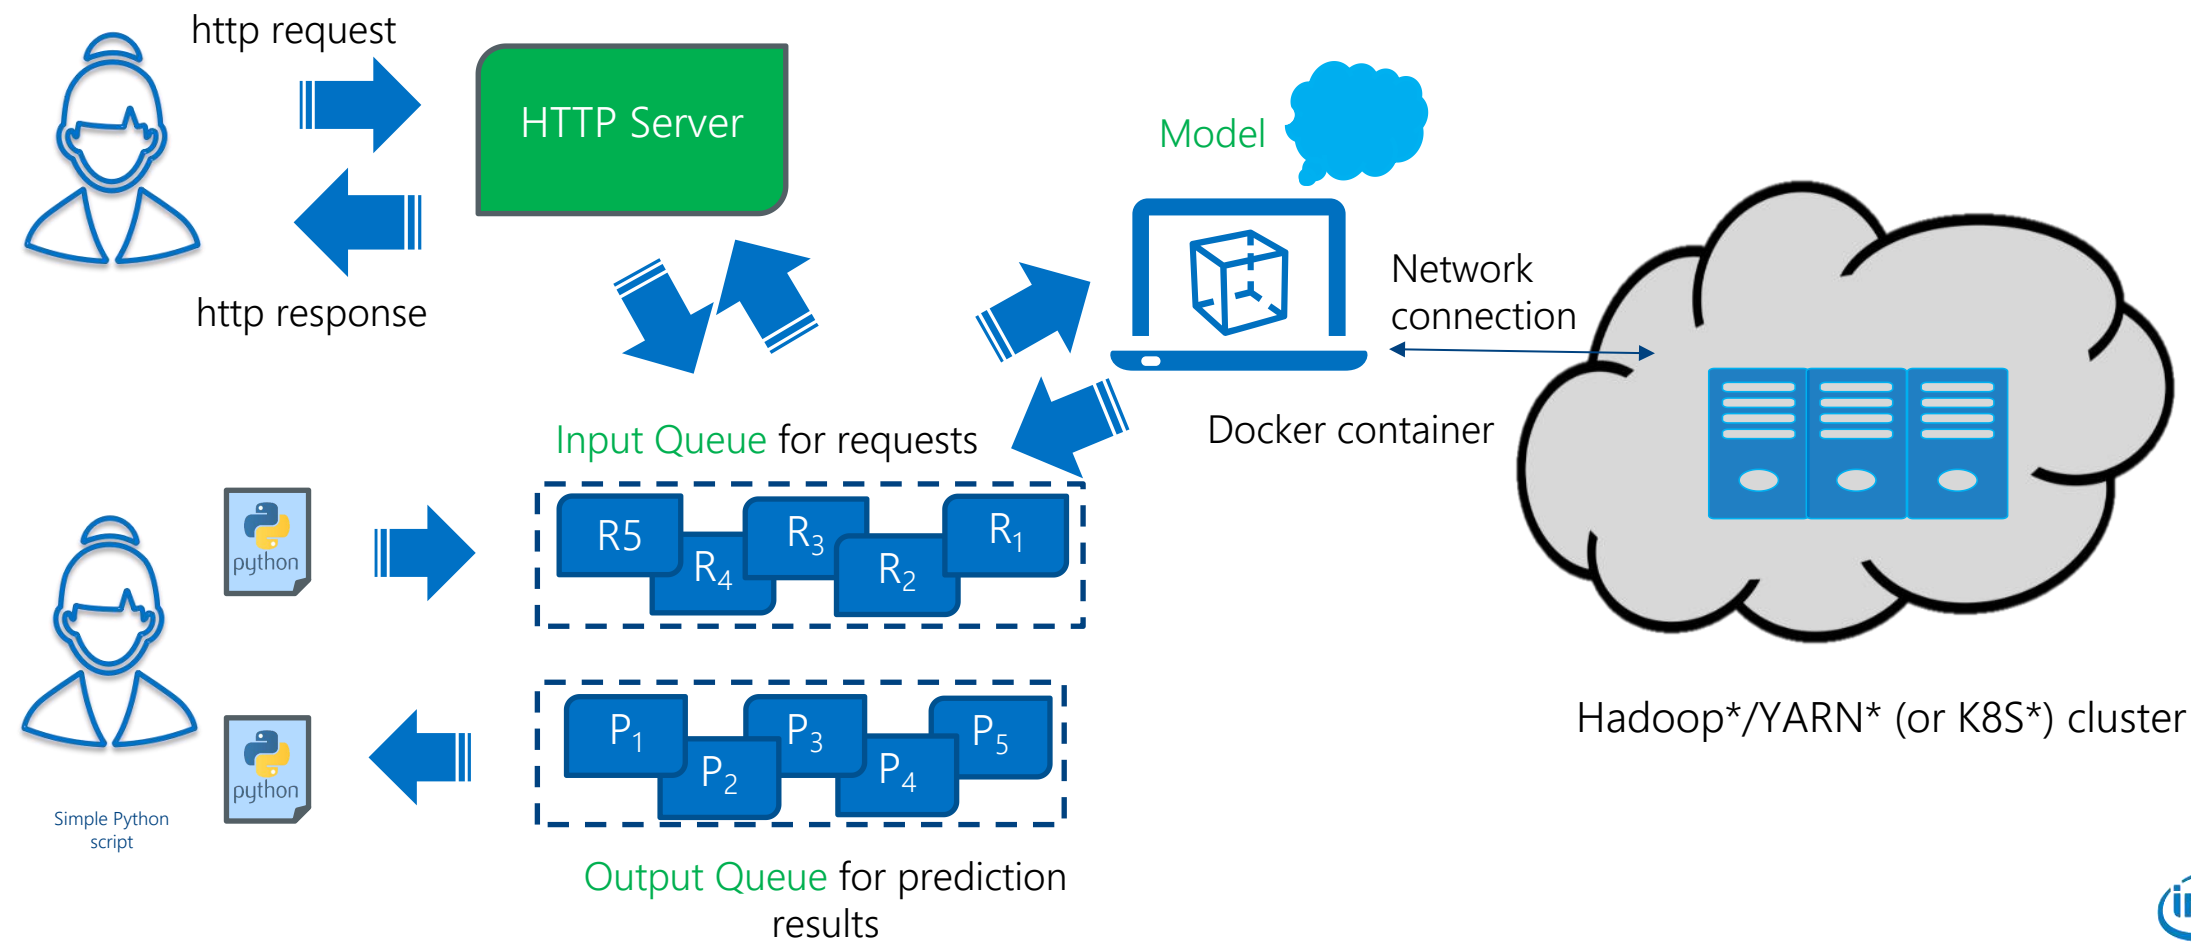


Powered by oneAPI

<https://github.com/intel-analytics/analytics-zoo>





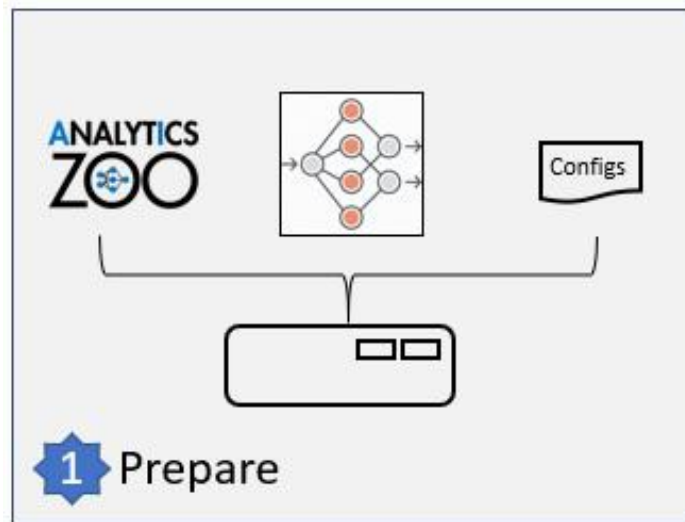
# Analytics Zoo Cluster Serving



\*Other names and brands may be claimed as the property of others.

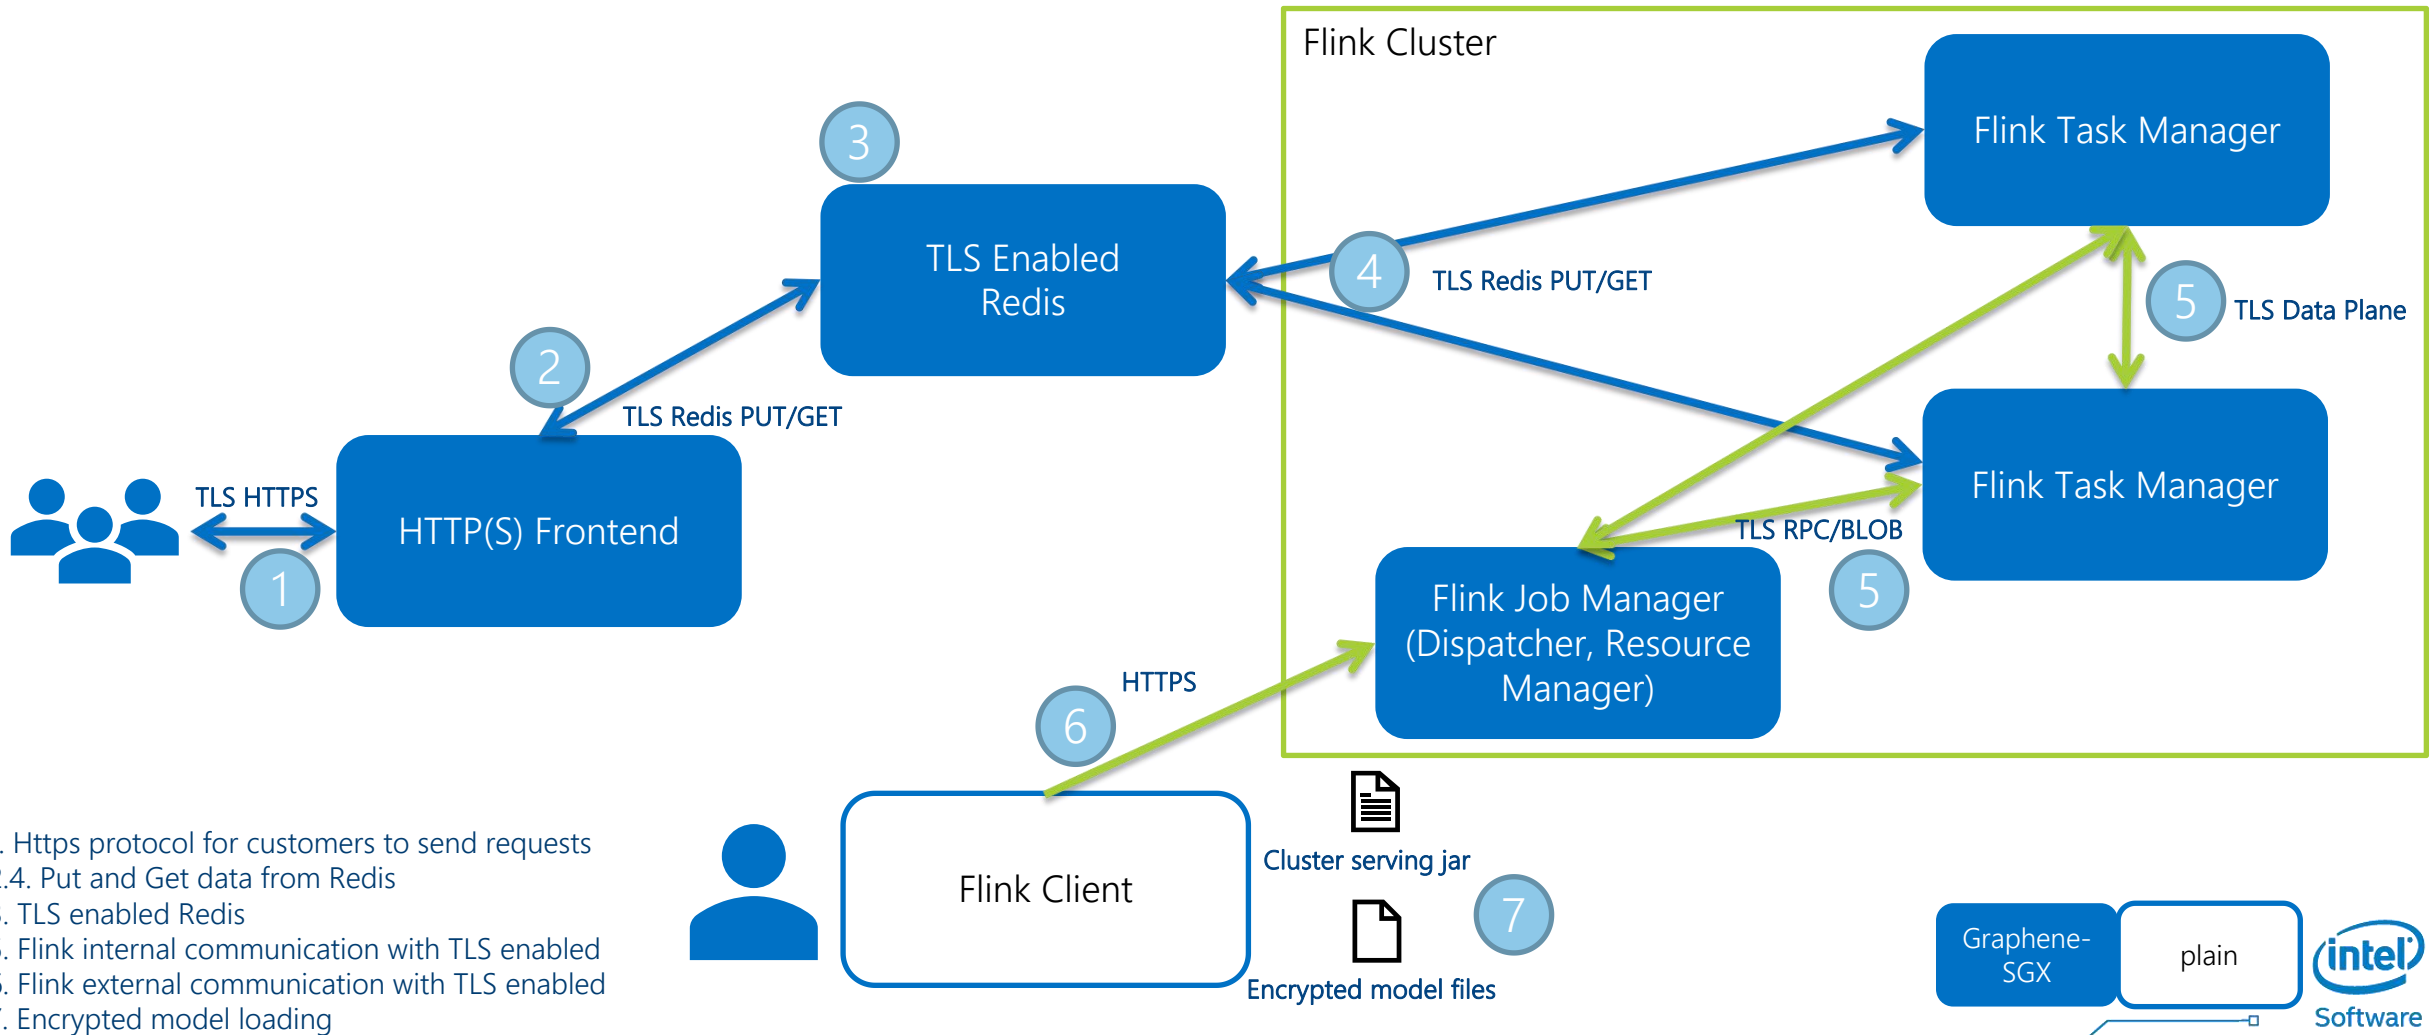
# Cluster Serving Workflow Overview

1. Install and prepare Cluster Serving environment on a local node
2. Launch the Cluster Serving service
3. Distributed, real-time (streaming) inference





# Secured Cluster Serving on Graphene-SGX



# Secured HTTP Frontend & Secured Redis

- Secured HTTP Frontend
  - TLS HTTPS enabled REST requests
  - TLS enabled Jedis GET/PUT to Redis
- TLS Enabled Redis
  - Rebuild Redis with BUILD\_TLS=yes
  - HTTP Frontend TLS enabled Jedis GET/PUT to Redis
  - Flink Source/Sink TLS enabled Jedis GET/PUT to Redis



# Secured Flink Cluster

- Flink SSL Setup and Internal and External Connectivity

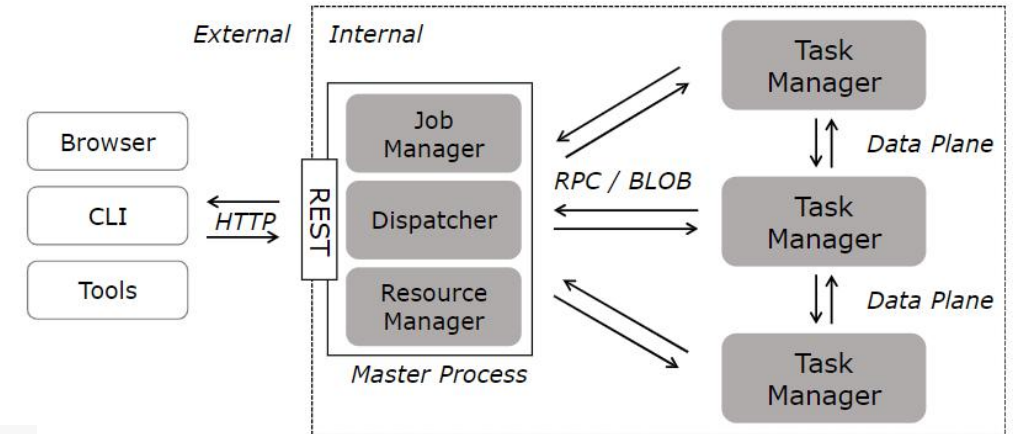
- Internal

- nano flink-conf.yaml

```
security.ssl.internal.enabled: true
security.ssl.internal.keystore: /home/sgx/glorysdj/keys/keystore.pkcs12
security.ssl.internal.truststore: /home/sgx/glorysdj/keys/keystore.pkcs12
security.ssl.internal.keystore-password: XXXXXXXXXr
security.ssl.internal.truststore-password: XXXXXXXXX
security.ssl.internal.key-password: XXXXXXXXXr
```

- External

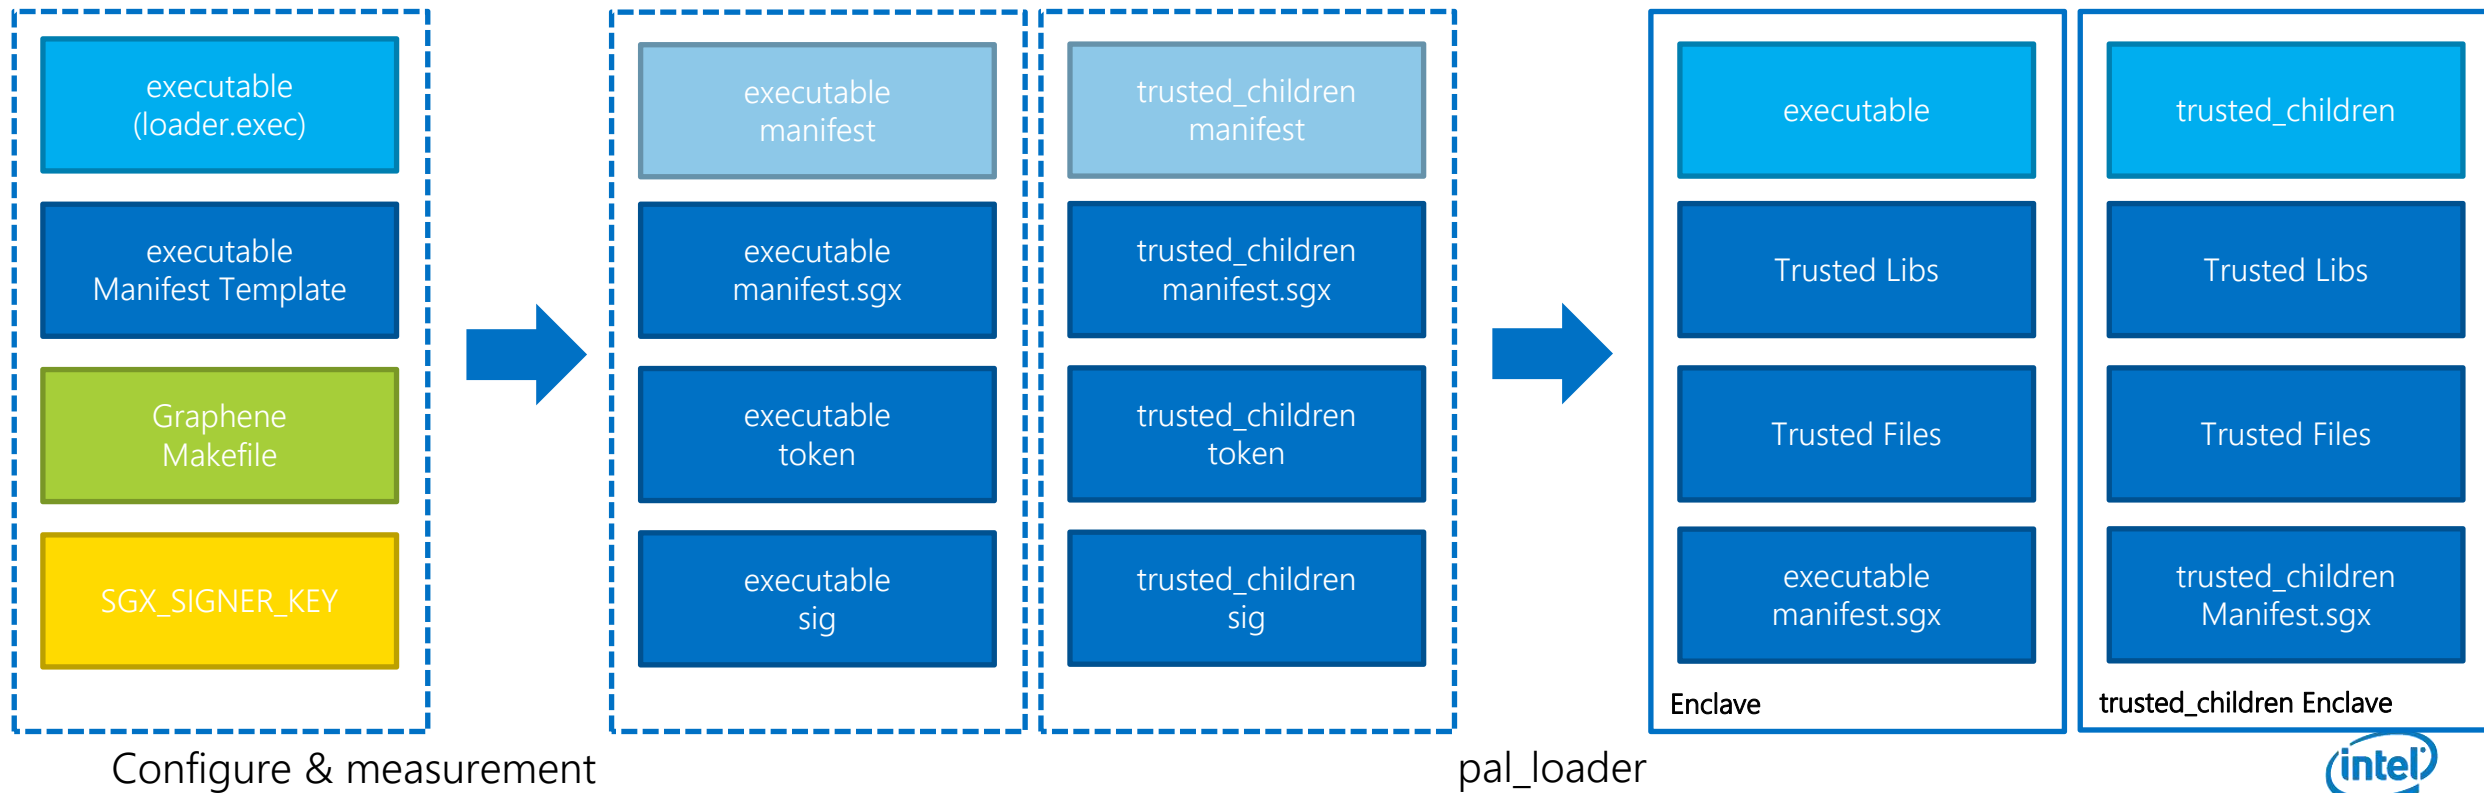
```
security.ssl.rest.enabled: true
security.ssl.rest.keystore: /path/to/flink/conf/rest.keystore
security.ssl.rest.truststore: /path/to/flink/conf/rest.truststore
security.ssl.rest.keystore-password: rest_keystore_password
security.ssl.rest.truststore-password: rest_truststore_password
security.ssl.rest.key-password: rest_keystore_password
```



# Encrypted Model loading

- Encrypted Model loading
  - trait **EncryptSupportive**
    - def encryptWith**AES256**(content: String, secret: String, salt: String): String
    - def decryptWithAES256(content: String, secret: String, salt: String): String
    - def encryptFileWithAES256(filePath: String, secret: String, salt: String, outputFile: String, encoding: String = "UTF-8")
    - def decryptFileWithAES256(filePath: String, secret: String, salt: String): String
    - def decryptFileWithAES256(filePath: String, secret: String, salt: String, outputFile: String)
  - **InferenceModel**
    - def doLoadEncryptedOpenVINO(modelPath: String, weightPath: String, secret: String, salt: String, batchSize: Int = 0)
  - **ClusterServing**
    - secret/salt = jedis.hget(Conventions.MODEL\_SECURED\_KEY, Conventions.MODEL\_SECURED\_SECRET/...)
    - model.doLoadEncryptedOpenVINO(defPath, weightPath, secret, salt, coreNum)

# Run Executable on Graphene-SGX



# Run Analytics Zoo Secured Cluster Serving on Graphene-SGX

- Two Executables(loader.exec)
  - loader.exec = file:redis-server
  - loader.exec = file:/bin/bash
    - child enclaves: `sgx.trusted_children.ls`, `sgx.trusted_children.cat`, `sgx.trusted_children.rm`, ..., **`sgx.trusted_children.java`**  
**= `file:/usr/bin/java`**
- Trusted Files: `sgx.trusted_files.keys_keystore_***` = `file:/home/sgx/keys/***`
- Allowed Files: `sgx.allowed_files.jvm` = `file:/usr/lib/jvm`
- Enlarge enclave size: `sgx.enclave_size` = 32G
- Enlarge thread num: `sgx.thread_num` = 1024
- Run the commands
  - `SGX=1 ./pal_loader redis-server ...`
  - `SGX=1 ./pal_loader bash.manifest -c "....."`

Thanks!