

# Context-aware Fast Food Recommendation with Ray on Apache Spark at Burger King

LUYANG WANG Burger King Corporation

KAI HUANG Intel Corporation

# Agenda

## LUYANG WANG

- Food recommendation use case
- TxT model in detail




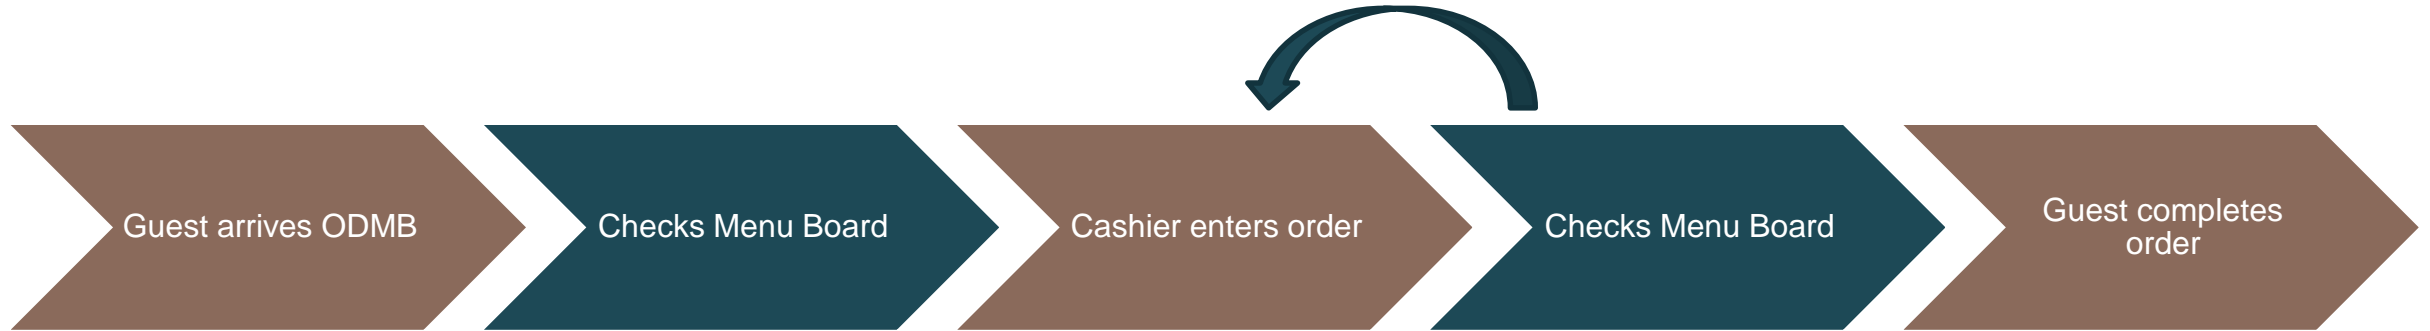
## KAI HUANG

- AI on big data
- Distributed training pipeline with Ray on Apache Spark

# Food Recommendation Use Case

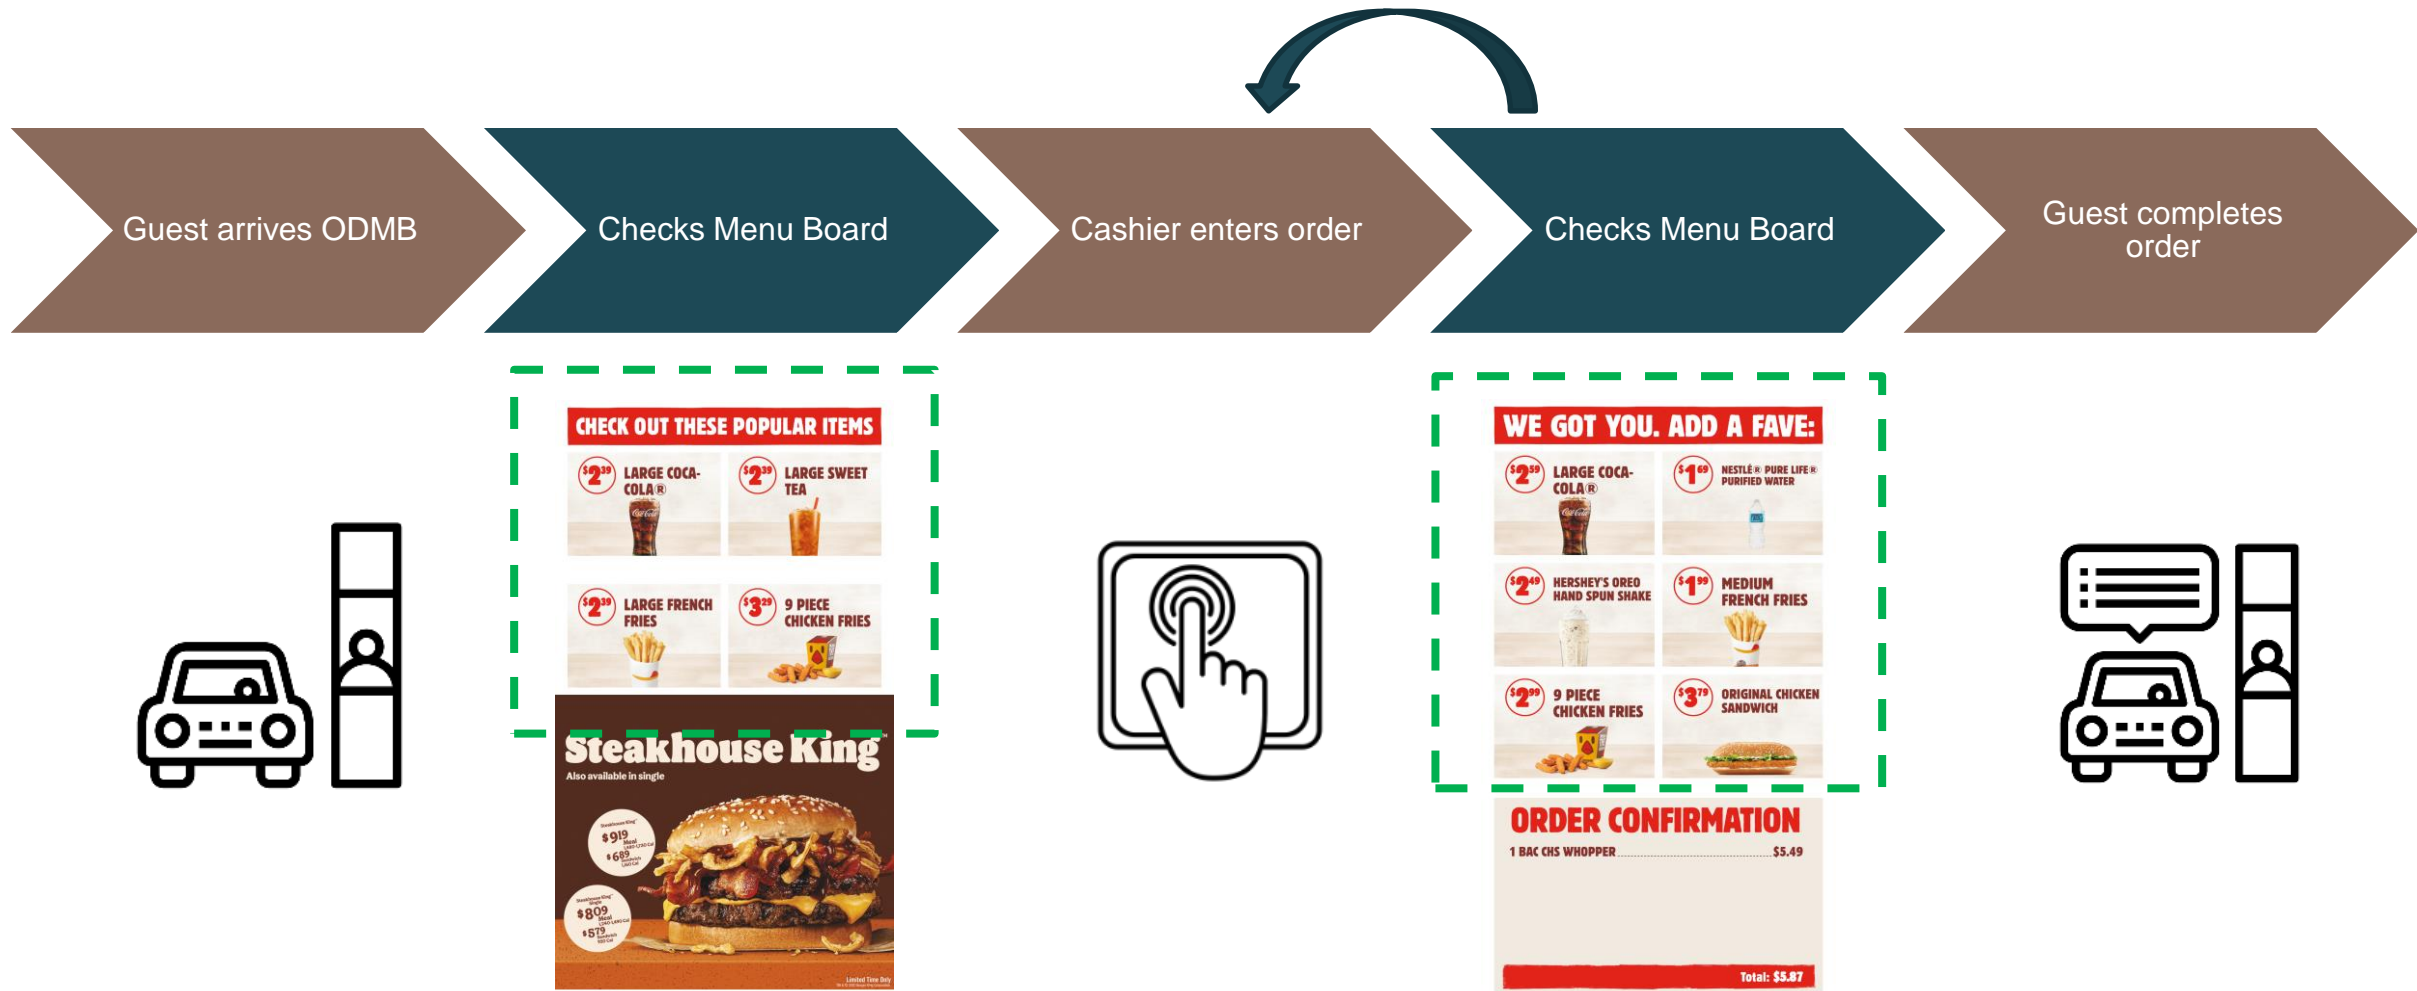


# Food Recommendation Use Case





# Food Recommendation Use Case



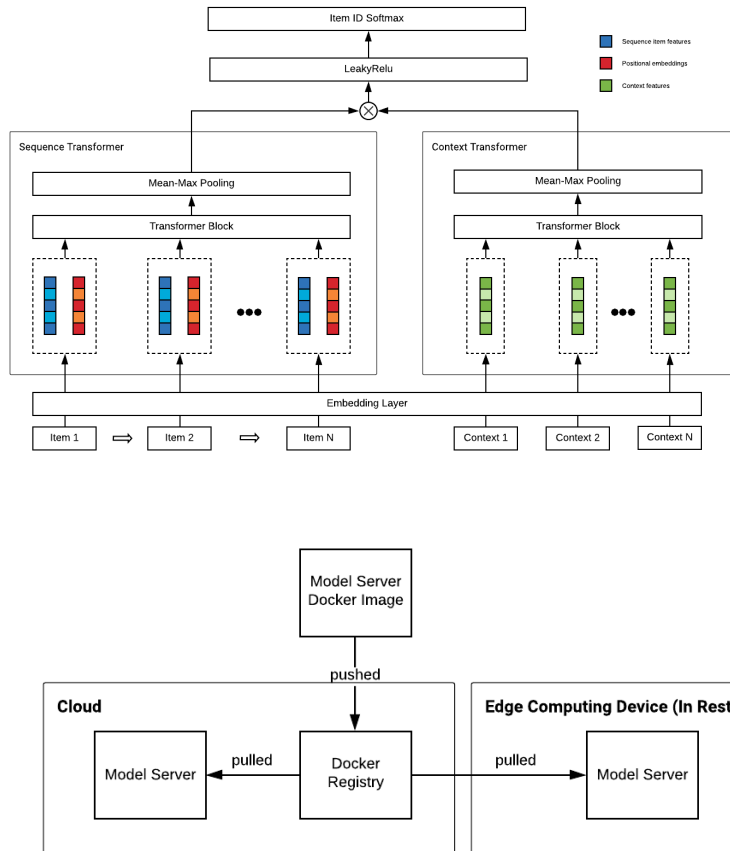
# Use Case Challenges



## Challenges

- Lack of user identifiers
- Same session food compatibilities
- Other variables in our use case: locations, weathers, time, etc.
- Deployment challenges

# Use Case Challenges



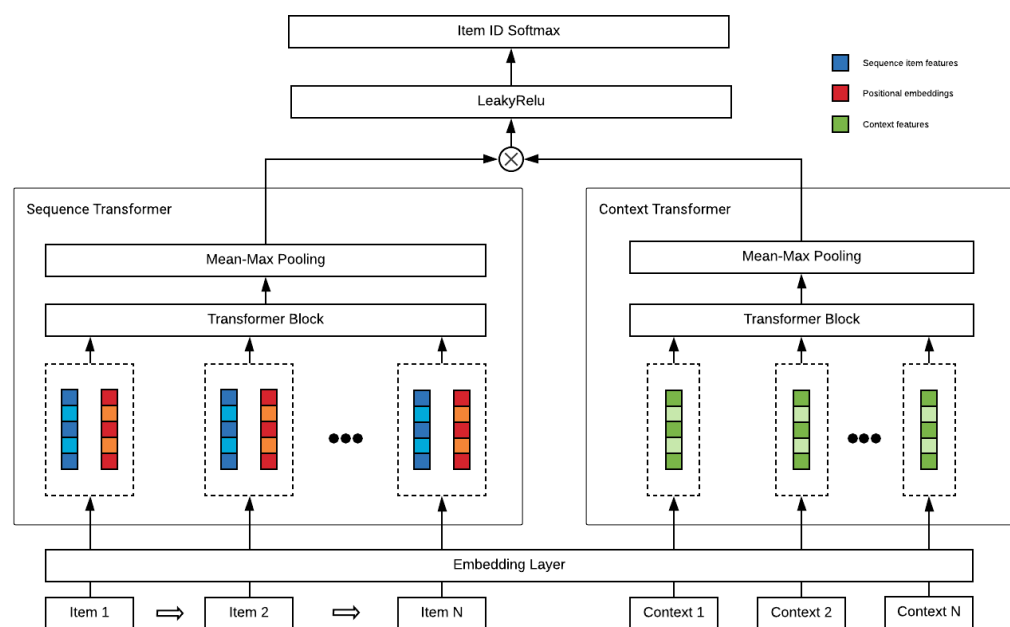
## Solutions

- Session based recommendation model
- Able to take complex context features into consideration
- Able to be deployed anywhere, both edge / cloud

# Transformer Cross Transformer (TxT)



# TxT Model Overview

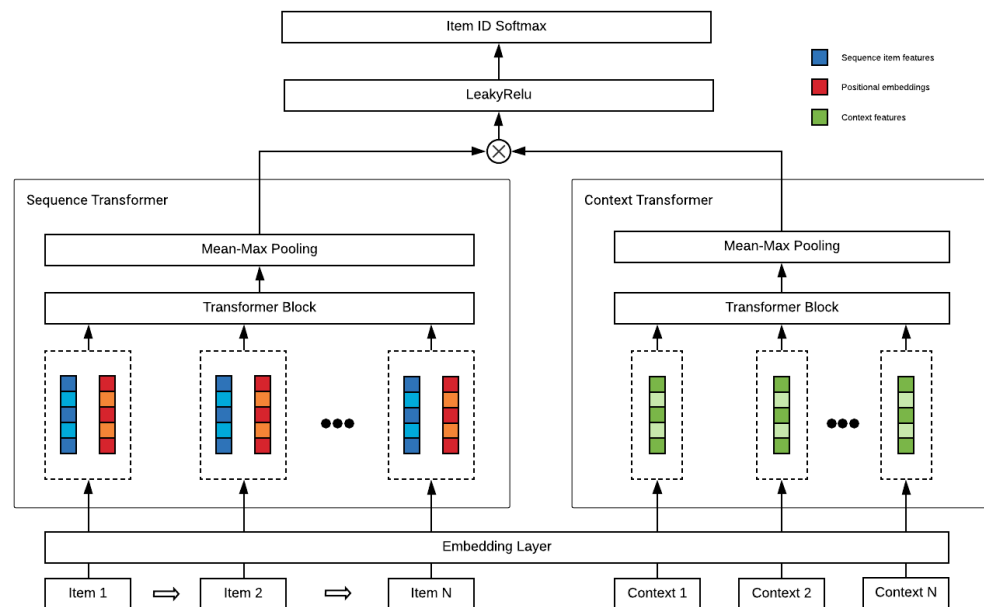


## Model Components


- **Sequence Transformer**
  - Taking item order sequence as input
- **Context Transformer**
  - Taking multiple context features as input
- **Latent Cross Joint Training**
  - Element-wise product for both transformer outputs

# Model Comparison

TxT

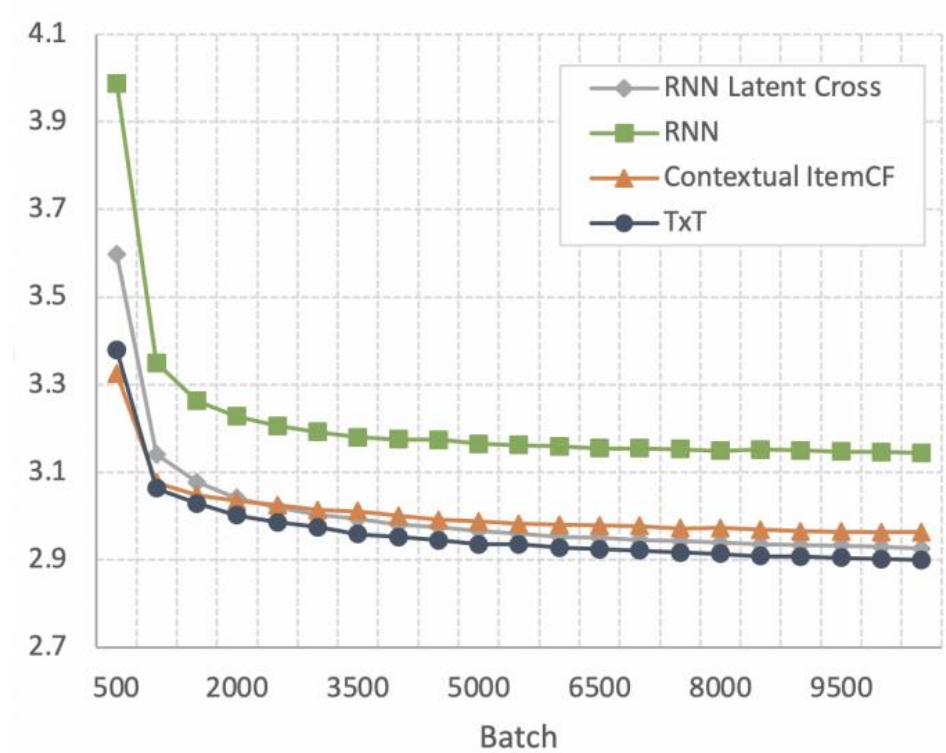


RNN Latent Cross



# Offline Evaluation

## Offline Training Loss




## Offline Training Result

Model	Top1 Accuracy	Top3 Accuracy
RNN	29.98%	46.24%
Contextual ItemCF	32.18%	48.37%
RNN Latent Cross	33.10%	49.98%
TxT	34.52%	52.37%

# Online Performance

## Inference Performance

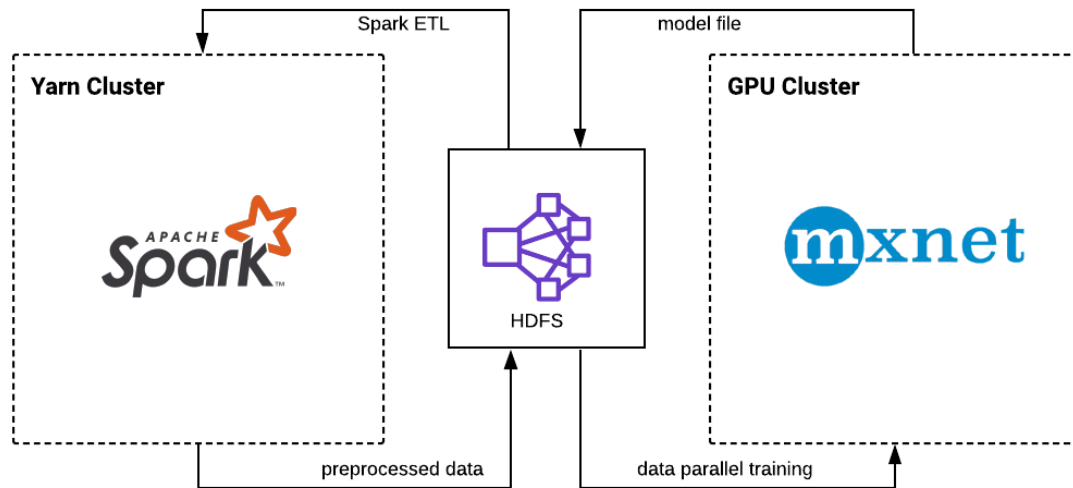


## A/B Testing Result

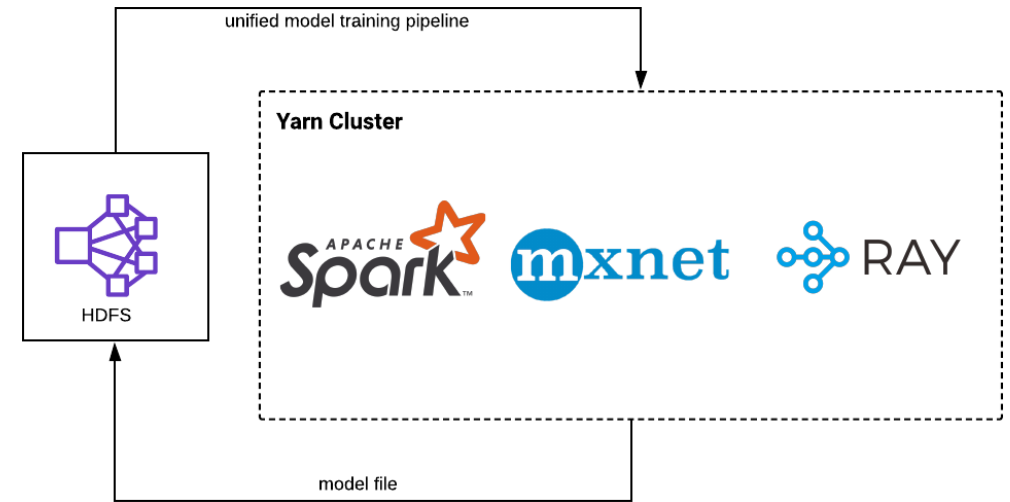
Model	Conversation Rate Gain	Add-on Sales Gain
RNN Latent Cross (control)	-	-
TxT	+7.5%	+4.7%

# Model Training Architecture

## Previous



## Current





# AI on Big Data





**DATA+AI SUMMIT EUROPE**

**#DataTeams #DataAISummit**

# AI on Big Data



## Accelerating Data Analytics + AI Solutions At Scale

-  BigDL: Distributed Deep Learning Framework for Apache Spark  
<https://github.com/intel-analytics/BigDL>
-  Analytics Zoo: Distributed TensorFlow, Keras and PyTorch on Apache Spark/Flink & Ray <https://github.com/intel-analytics/analytics-zoo>
- We develop *Project Orca* in Analytics Zoo based on Spark and Ray to allow users to easily scale out single node Python notebook across large clusters, by providing:
  - Data-parallel preprocessing for Python AI (supporting common Python libraries such as Pandas, Numpy, PIL, TensorFlow Dataset, PyTorch DataLoader, etc.)
  - Sklearn-style APIs for transparently distributed training and inference (supporting TensorFlow, PyTorch, Keras, MXNet, Horovod, etc.)<https://github.com/intel-analytics/analytics-zoo/tree/master/pyzoo/zoo/orca>

# Ray

Ray is a fast and simple framework for building and running distributed applications.

- Ray Core provides easy Python interface for parallelism by using remote functions and actors.

Ray is packaged with several high-level libraries to accelerate machine learning workloads.

- Tune: Scalable Experiment Execution and Hyperparameter Tuning
- RLlib: Scalable Reinforcement Learning
- RaySGD: Distributed Training Wrappers
- <https://github.com/ray-project/ray/>



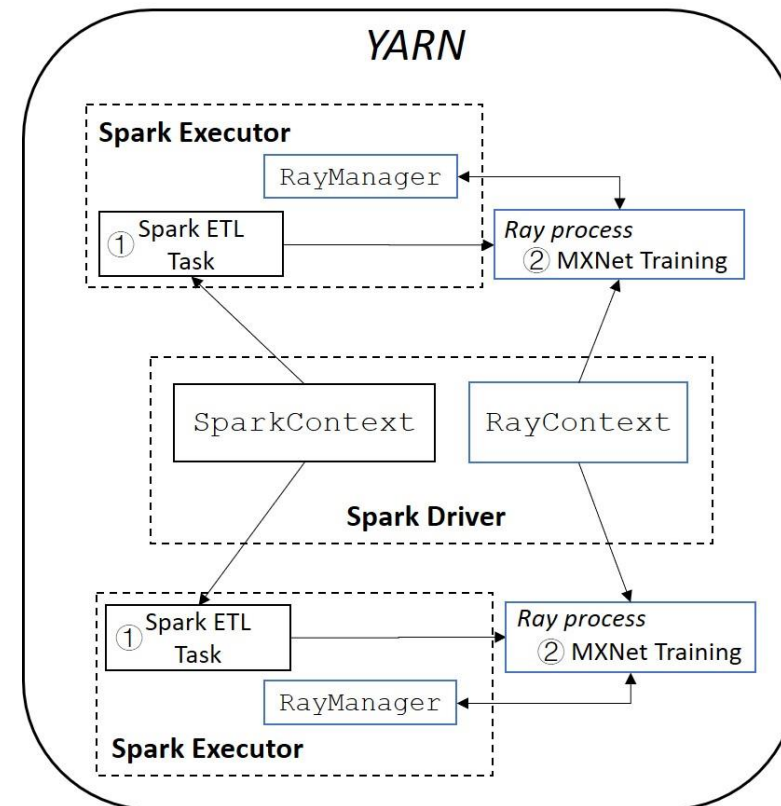
# Distributed Training Pipeline on Big Data



# RayOnSpark

Seamlessly integrate Ray applications into Spark data processing pipelines.

- Runtime cluster environment preparation.
- Create a *SparkContext* on the driver node and use Spark to perform data cleaning, ETL, and preprocessing tasks.
- *RayContext* on Spark driver launches Ray across the cluster.
- Similar to RaySGD, we implement a lightweight shim layer around native MXNet modules for easy deployment on YARN cluster.
- Each MXNet worker takes the local data partition of Spark RDD or DataFrame from the plasma object store used by Ray.





# End-to-end Distributed Training Pipeline

*Project Orca provides a user-friendly interface for the pipeline.*

- Minimum code changes and learning efforts are needed to scale the training from single node to big data clusters.
- The entire pipeline runs on a single cluster. No extra data transfer needed.

```
from zoo.orca import init_orca_context
from zoo.orca.learn.mxnet import Estimator

# init_orca_context unifies SparkContext and RayContext
sc = init_orca_context(cluster_mode="yarn", num_nodes, cores, memory)
# Use sc to load data and do data preprocessing.

mxnet_estimator = Estimator(train_config, model=txt, loss=SoftmaxCrossEntropyLoss(),
                             metrics=[mx.metric.Accuracy(), mx.metric.TopKAccuracy(3)])

mxnet_estimator.fit(data=train_rdd, validation_data=val_rdd, epochs=..., batch_size=...)
```

# Conclusion

- Context-Aware Fast Food Recommendation at Burger King with RayOnSpark

<https://arxiv.org/abs/2010.06197>

<https://medium.com/riselab/context-aware-fast-food-recommendation-at-burger-king-with-rayonspark-2e7a6009dd2d>

- For more details of RayOnSpark: [https://databricks.com/session\\_na20/running-emerging-ai-applications-on-big-data-platforms-with-ray-on-apache-spark](https://databricks.com/session_na20/running-emerging-ai-applications-on-big-data-platforms-with-ray-on-apache-spark)

- More information for Analytics Zoo at:

<https://github.com/intel-analytics/analytics-zoo>

<https://analytics-zoo.github.io/>



# Feedback

Your feedback is important to us.

Don't forget to rate  
and review the sessions.

