

使用 Analytics-Zoo 构建统一的大数据 AI 应用的架构实践

史栋杰

资深软件架构师

自我介绍

史栋杰，英特尔资深软件架构师。多年从事企业级计算、风控、大数据分析、云计算容器编排、数据分析与人工智能领域的研发，英特尔开源框架 BigDL 与 Analytics-Zoo 的贡献者之一。

目录

1. 大规模人工智能应用面临的挑战
2. 统一的大数据分析及人工智能
3. 跨行业的端到端客户案例实践

以数据为中心的世界

全球超过 OVER

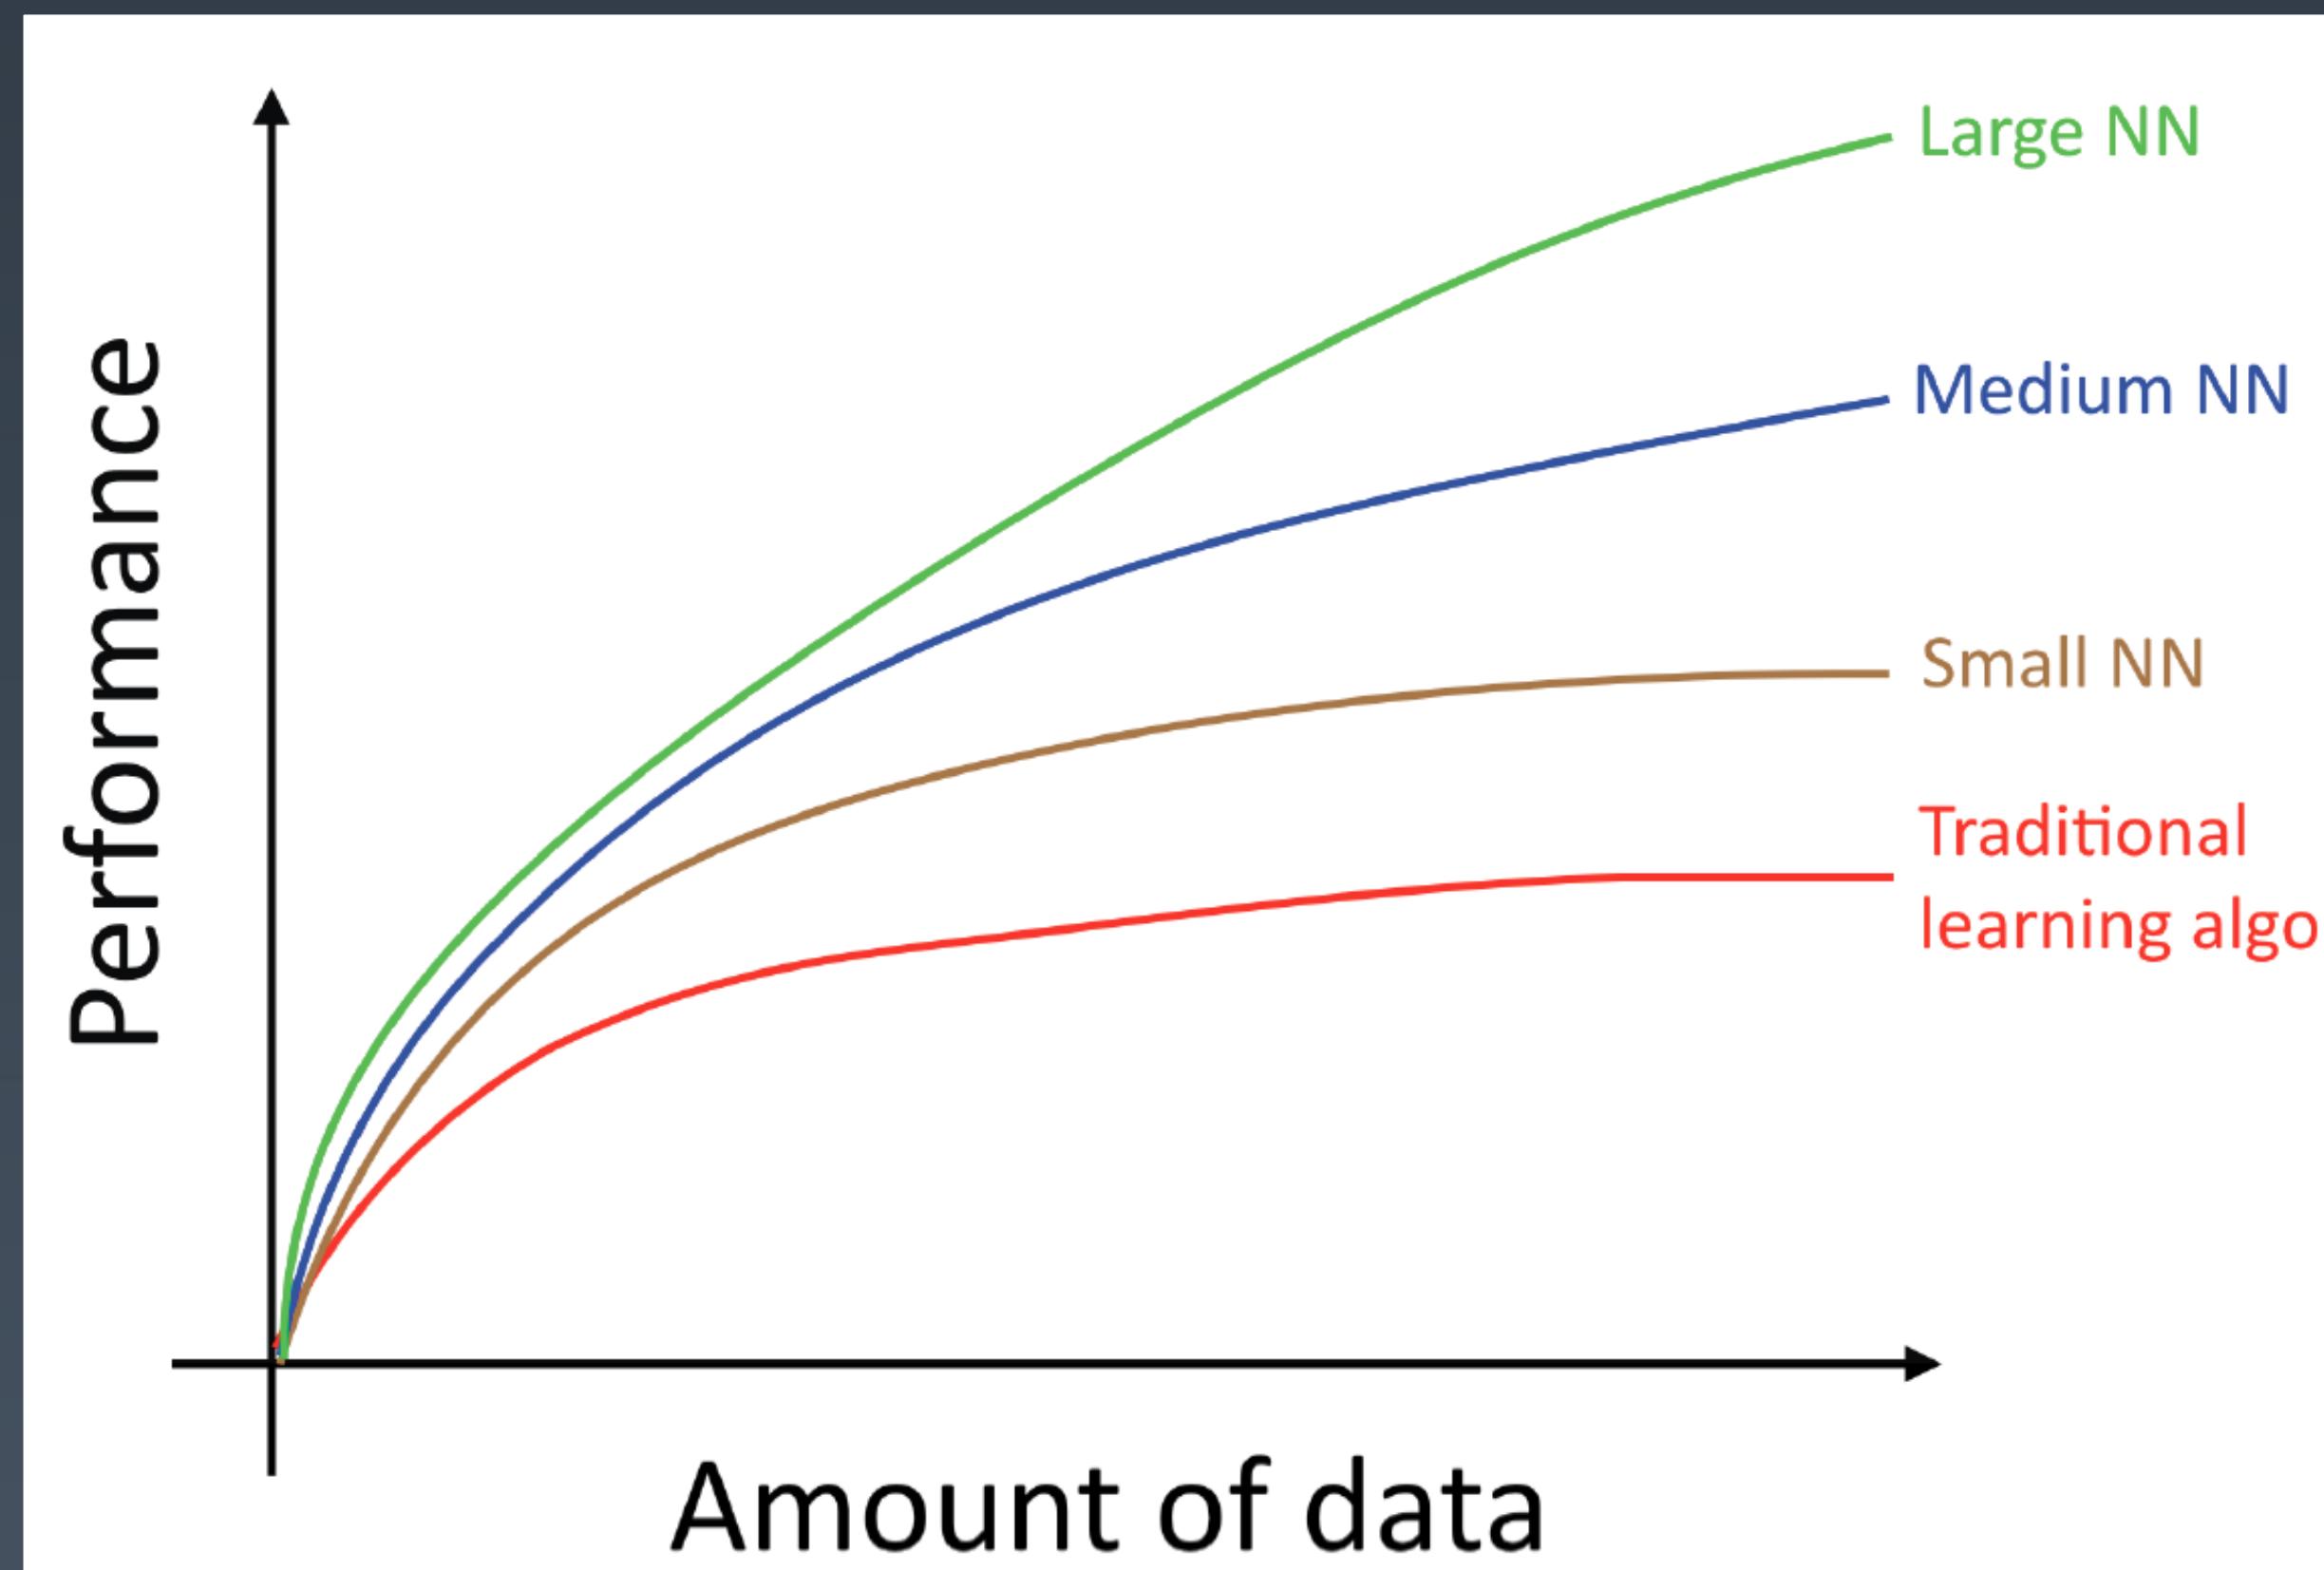
一半 HALF 数据 OF THE
WORLD'S DATA

创建于过去
WAS CREATED IN THE LAST
两年 2 YEARS

其中只有不到
2% 的数据
经过了分析
LESS THAN
HAS BEEN
ANALYZED

大规模人工智能应用

数据驱动深度学习和人工智能应用



“Machine Learning Yearning”,
Andrew Ng, 2016

大规模人工智能应用

正面临巨大的挑战



复杂性
Complexity

成本
Cost

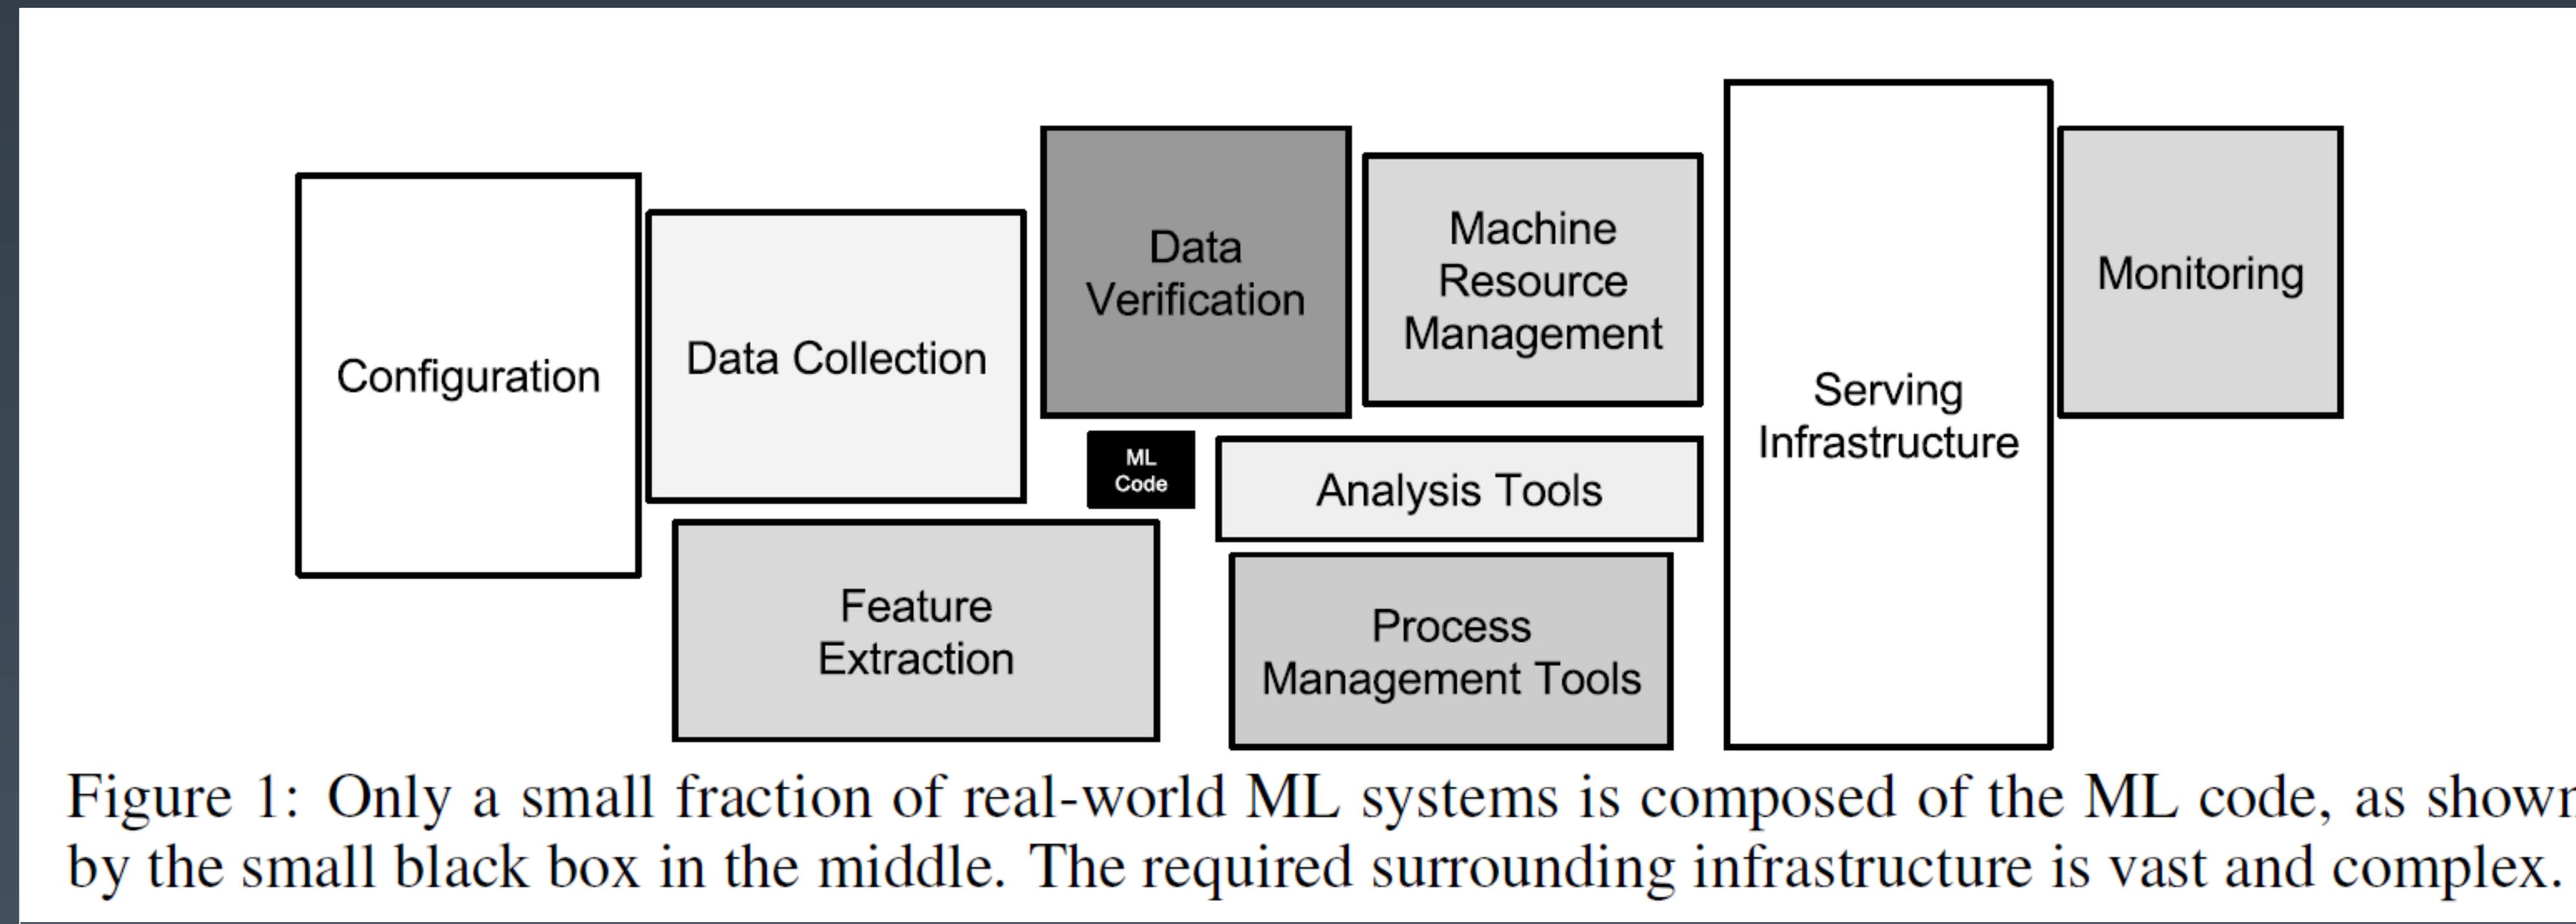
可扩展性
Scalability

专有接口
Proprietary Interfaces

数据隐私
Data Privacy

大规模人工智能应用

正面临巨大的挑战



“Hidden Technical Debt in Machine Learning Systems”, Sculley et al., Google, NIPS 2015

统一的大数据分析及人工智能

获取 / 存储

Source/Store

清洗 / 准备

Clean/Prepare

分析 / 建模

Analyze / Modeling

部署 / 可视化

Deploy / Visualize

集成的数据流水线

Unified Data Pipeline



数据管理
Data Management



数据分析
Data Analysis



数据科学及人工智能
Data Science and AI

大数据上的人工智能



基于 Apache Spark 的高性能
深度学习框架

High-Performance Deep Learning
Framework for Apache Spark

software.intel.com/bigdl



统一的分析 + 人工智能平台
Integrated Analytics + AI Toolkit

基于 Apache Spark 的分布式
TensorFlow、PyTorch、Keras 和 BigDL

高级流水线、参考用例、人工智能模型、特征工程等

<https://github.com/intel-analytics/analytics-zoo>

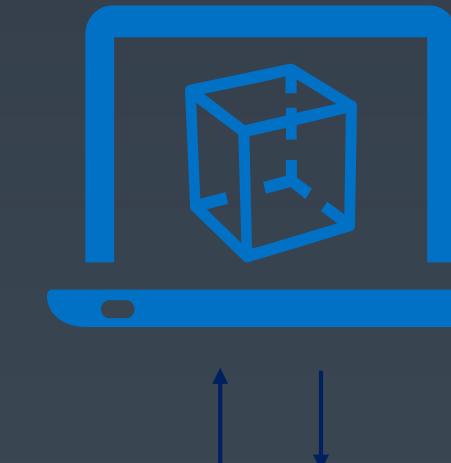
加快数据分析及人工智能大规模应用

Accelerating DATA Analytics + AI Solutions DEPLOYMENT At SCALE

统一的数据分析和AI流水线

端到端、从原型到生产化部署的无缝扩展
Seamless Scaling from Laptop to Production

在笔记本电脑上使用样本数据构建原型
Prototype on laptop using sample data



在集群上使用历史数据运行模型试验
Experiment on clusters with history data



在分布式生产环境中部署
Production deployment w/ distributed data pipeline



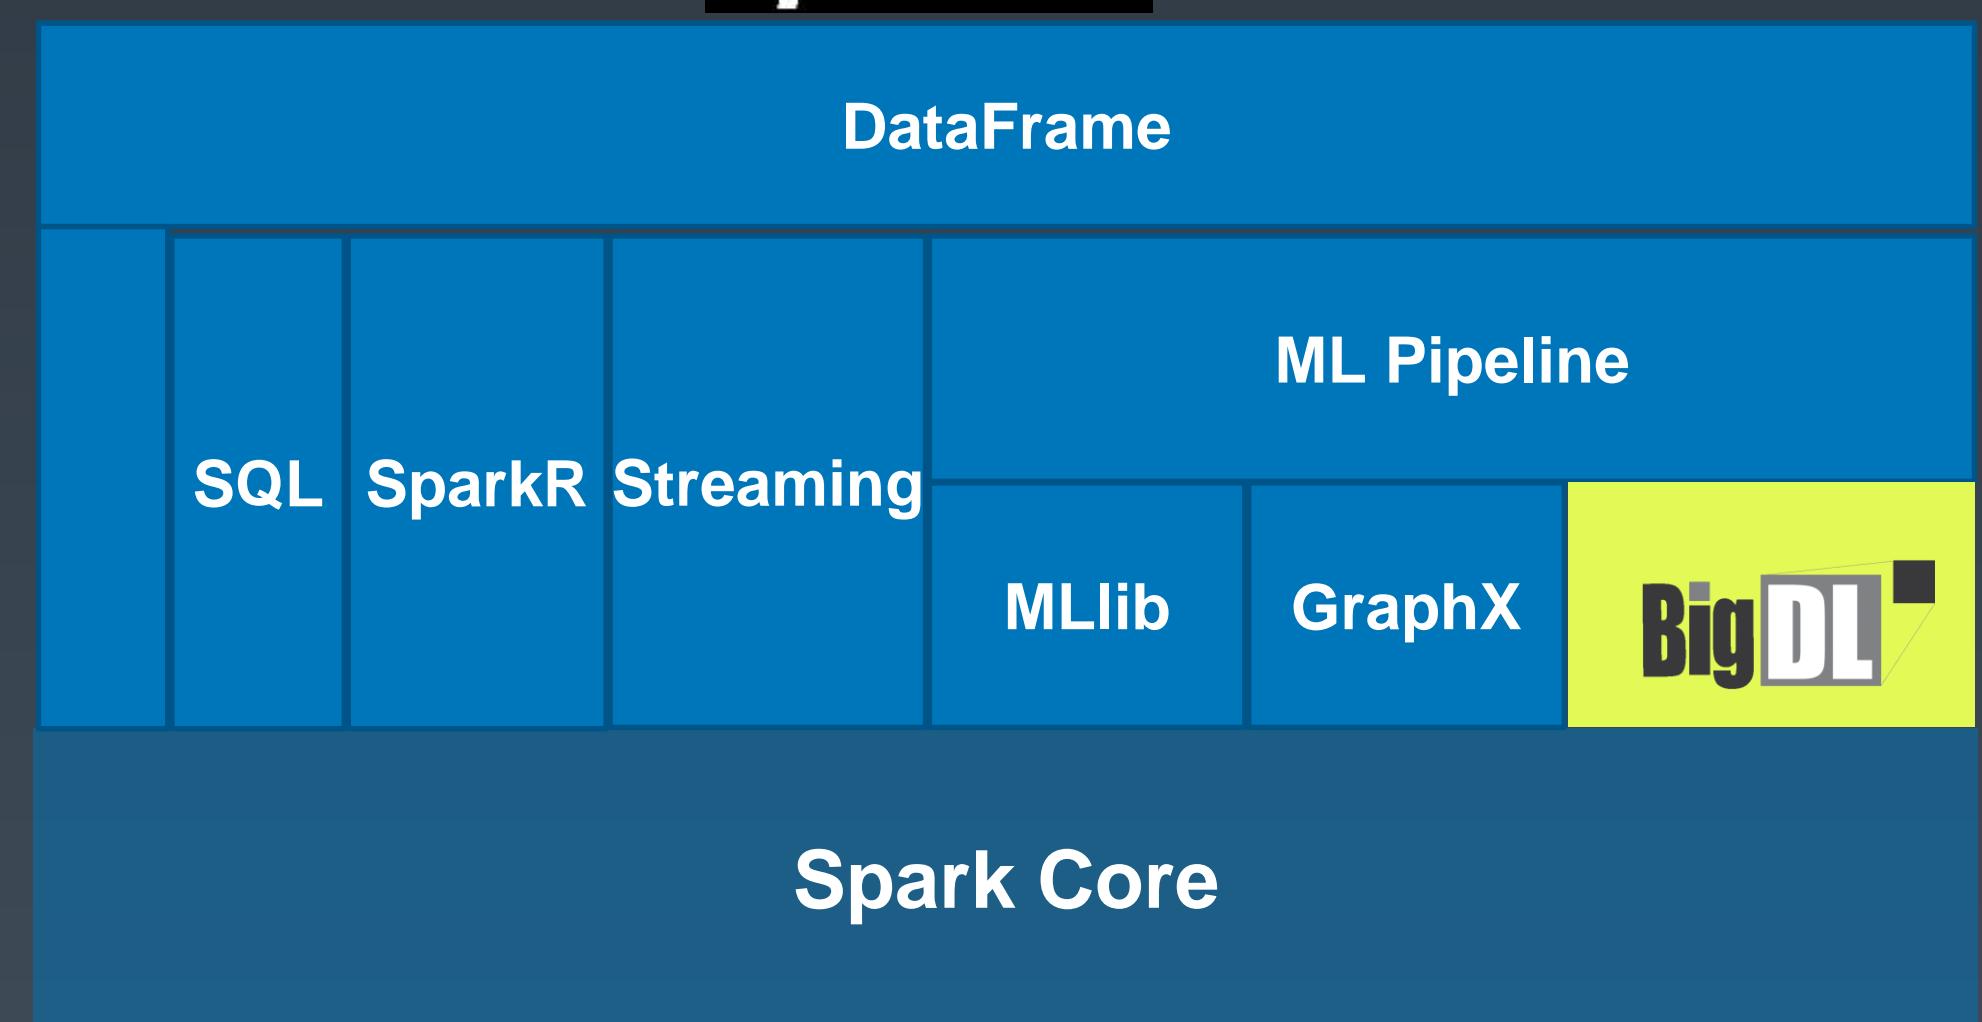
生产数据流水线



- 从笔记本电脑到分布式集群几乎无需任何代码更改
- 无需数据拷贝，直接访问生产大数据系统
- 高效构建端到端的数据分析+ AI 流水线原型
- 无缝扩展部署到大数据集群及生产环境

BigDL

- 构建于Apache Spark*之上的深度学习框架
- 将深度学习带给广大的大数据用户和大数据数据科学家
 - 以标准Spark*程序方式创建深度学习应用
 - 运行于现有的YARN*/Spark*/K8S*集群
- 与流行的深度学习框架具有同等功能
 - Caffe*, Torch*, Tensorflow*等
- 高性能
 - 使用Intel MKL, MKLDNN加速
- 有效的扩展性
 - 使用Spark*作分布式训练及推理

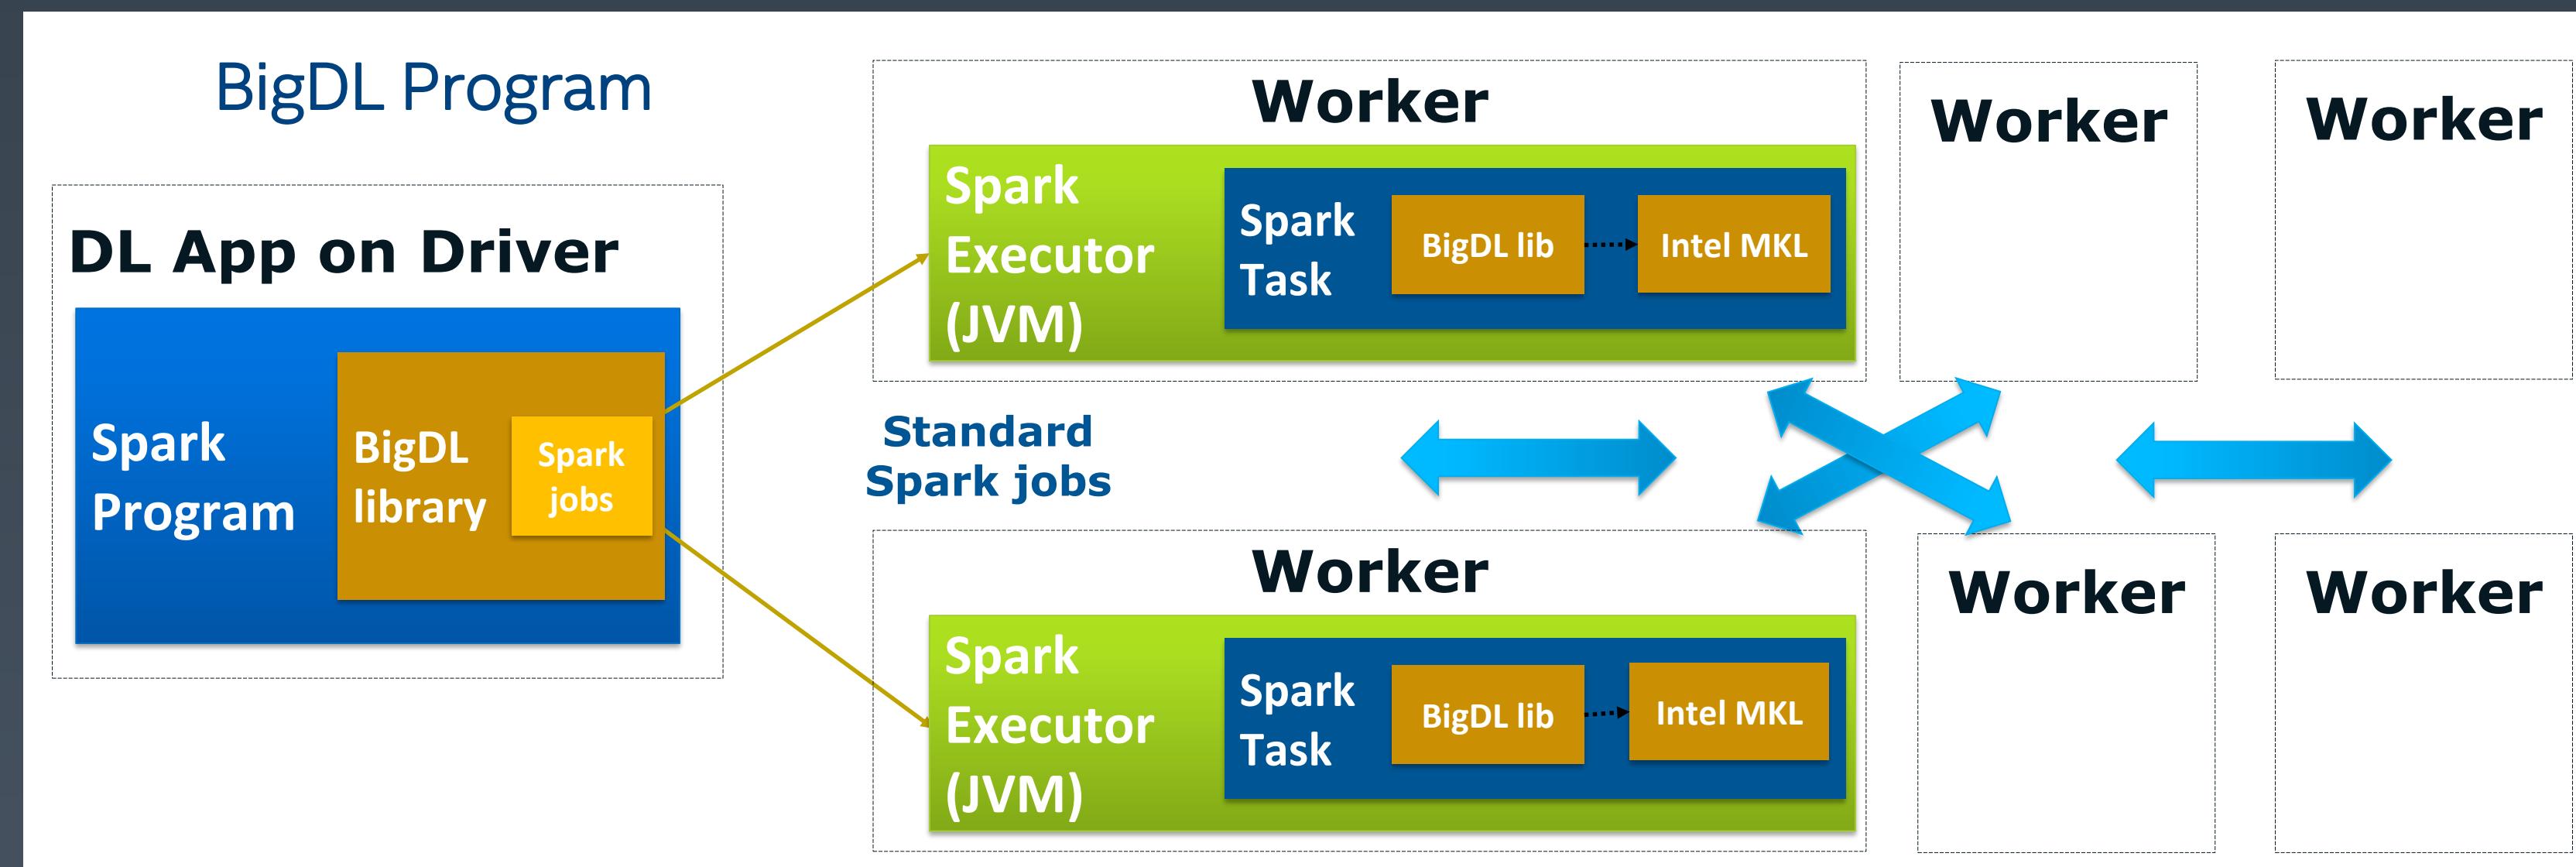


<https://github.com/intel-analytics/BigDL>

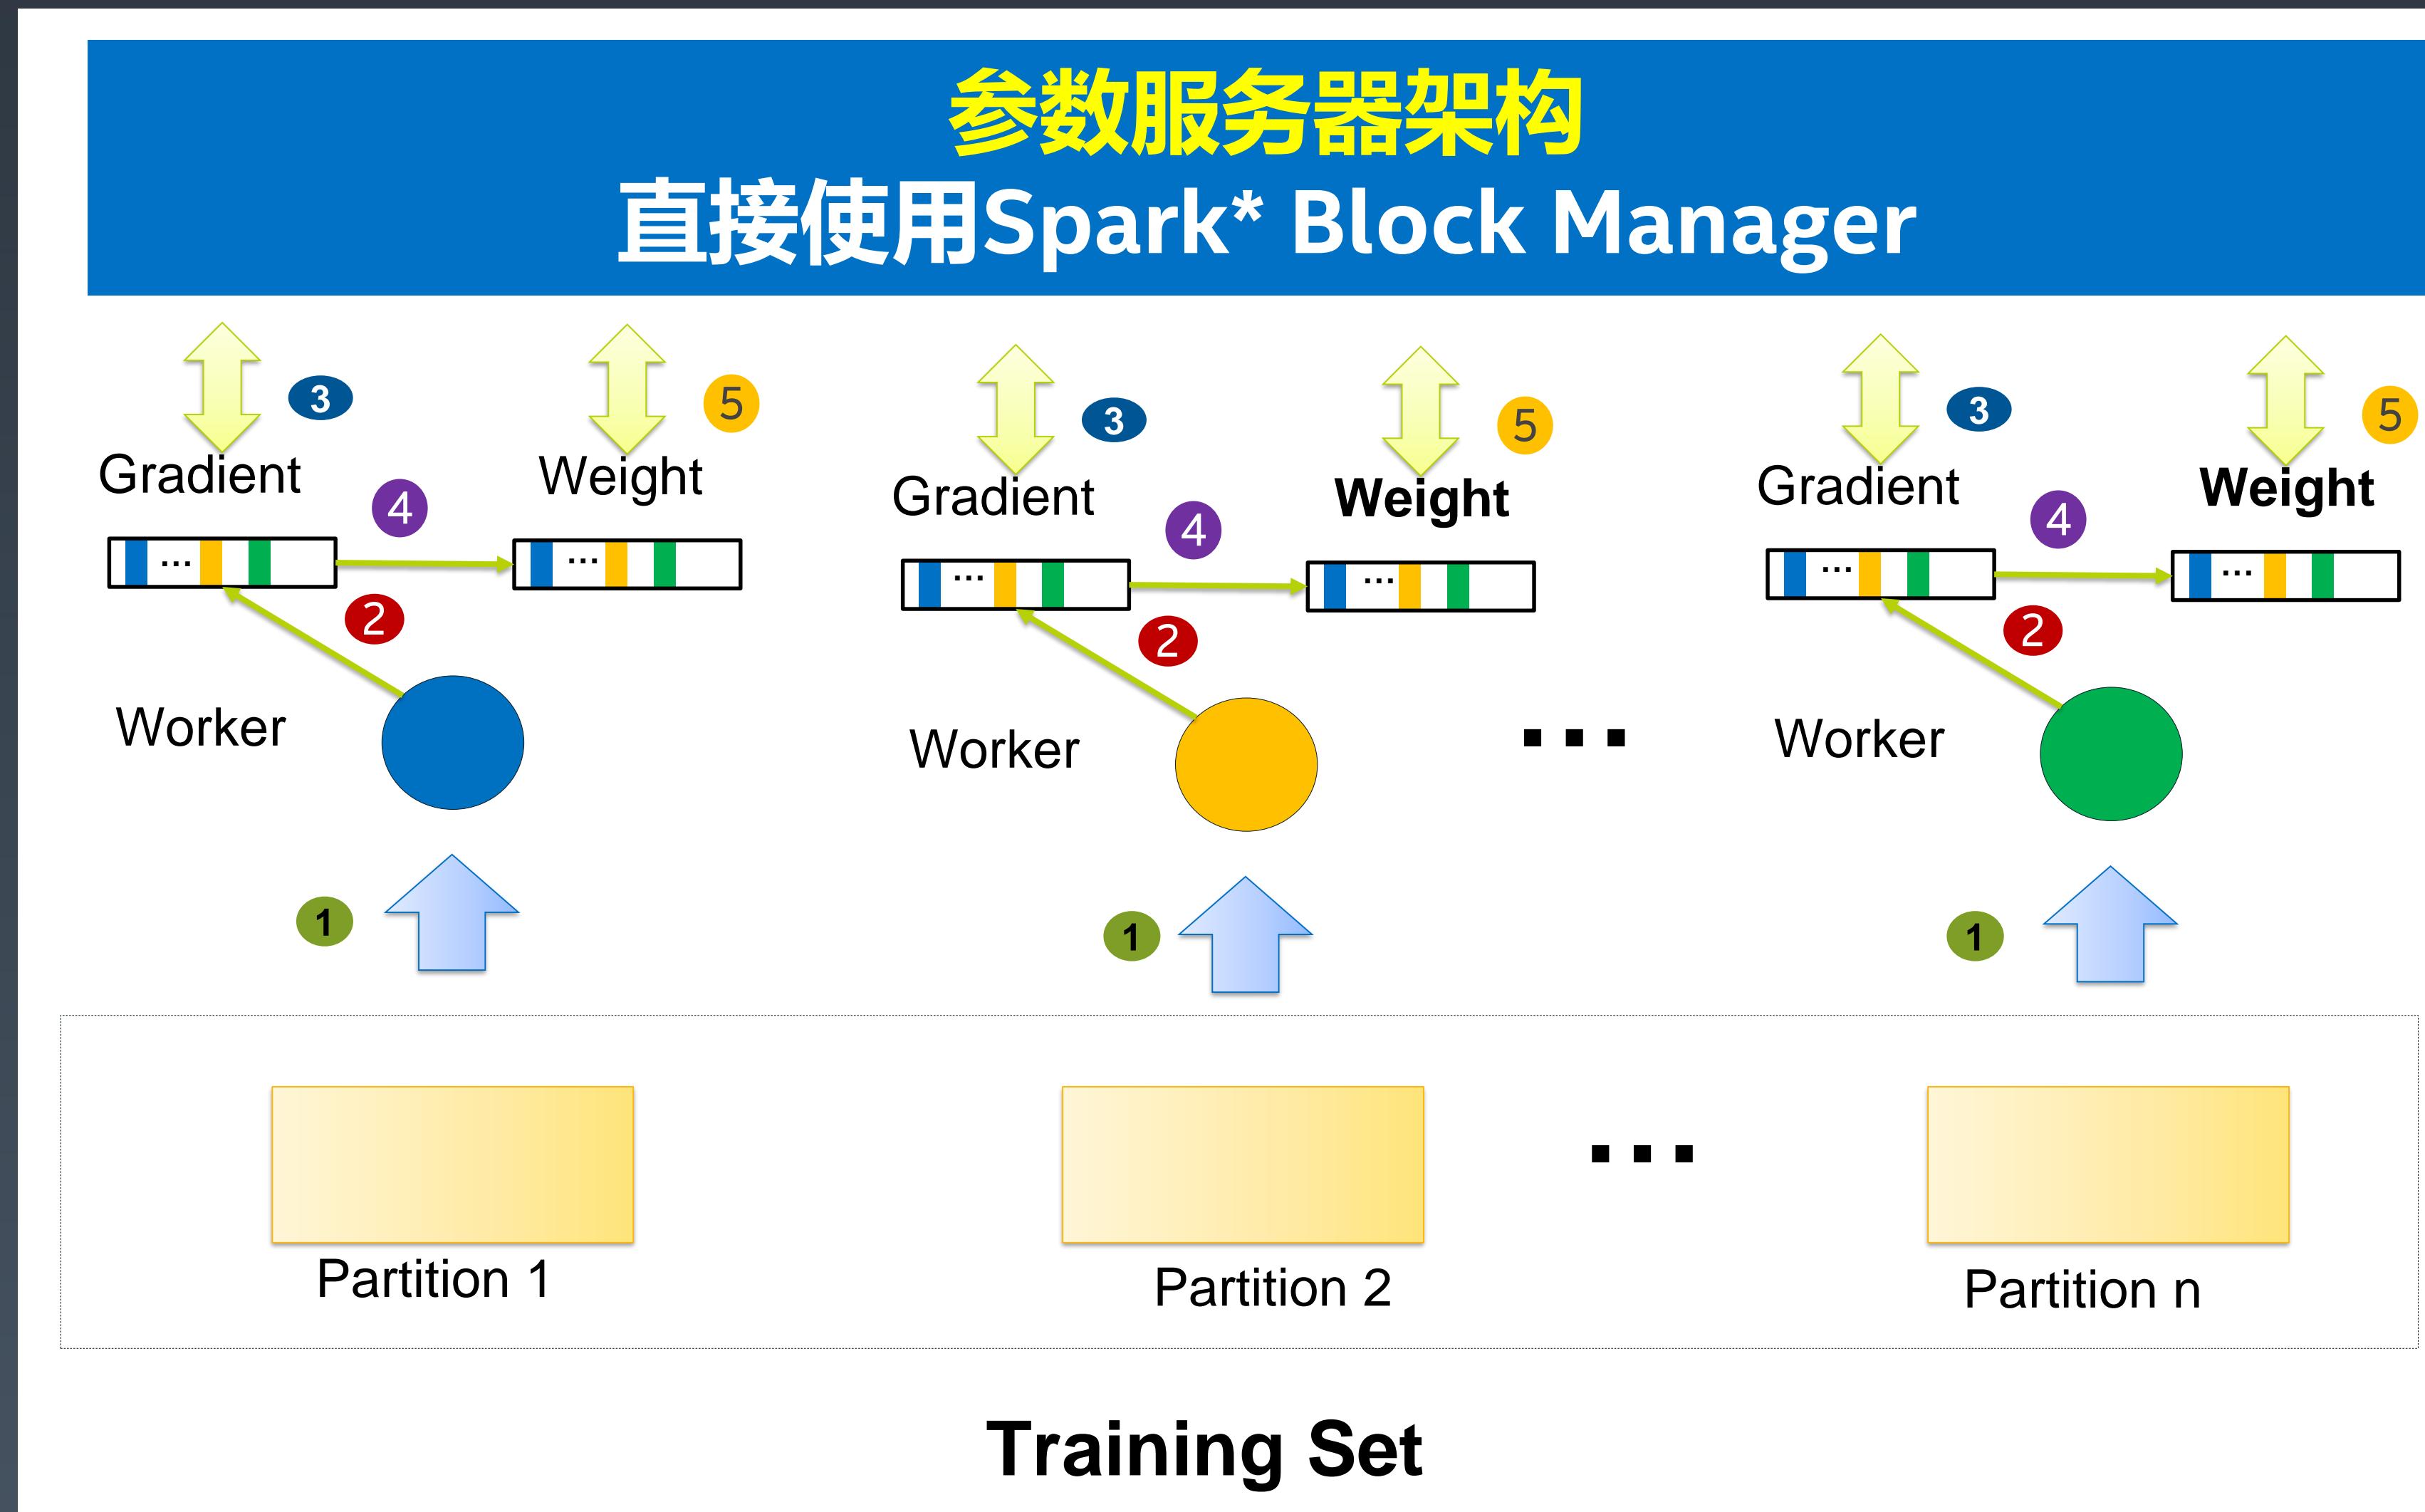
<https://bigdl-project.github.io/>

BigDL 以标准Spark*程序方式创建深度学习应用

- 标准的Spark* Job
 - 无需改变YARN*/Spark*/K8S*集群
- 数据并发
 - 每一个Spark Task在全量数据集的一个子集（Batch）上运行统一的模型



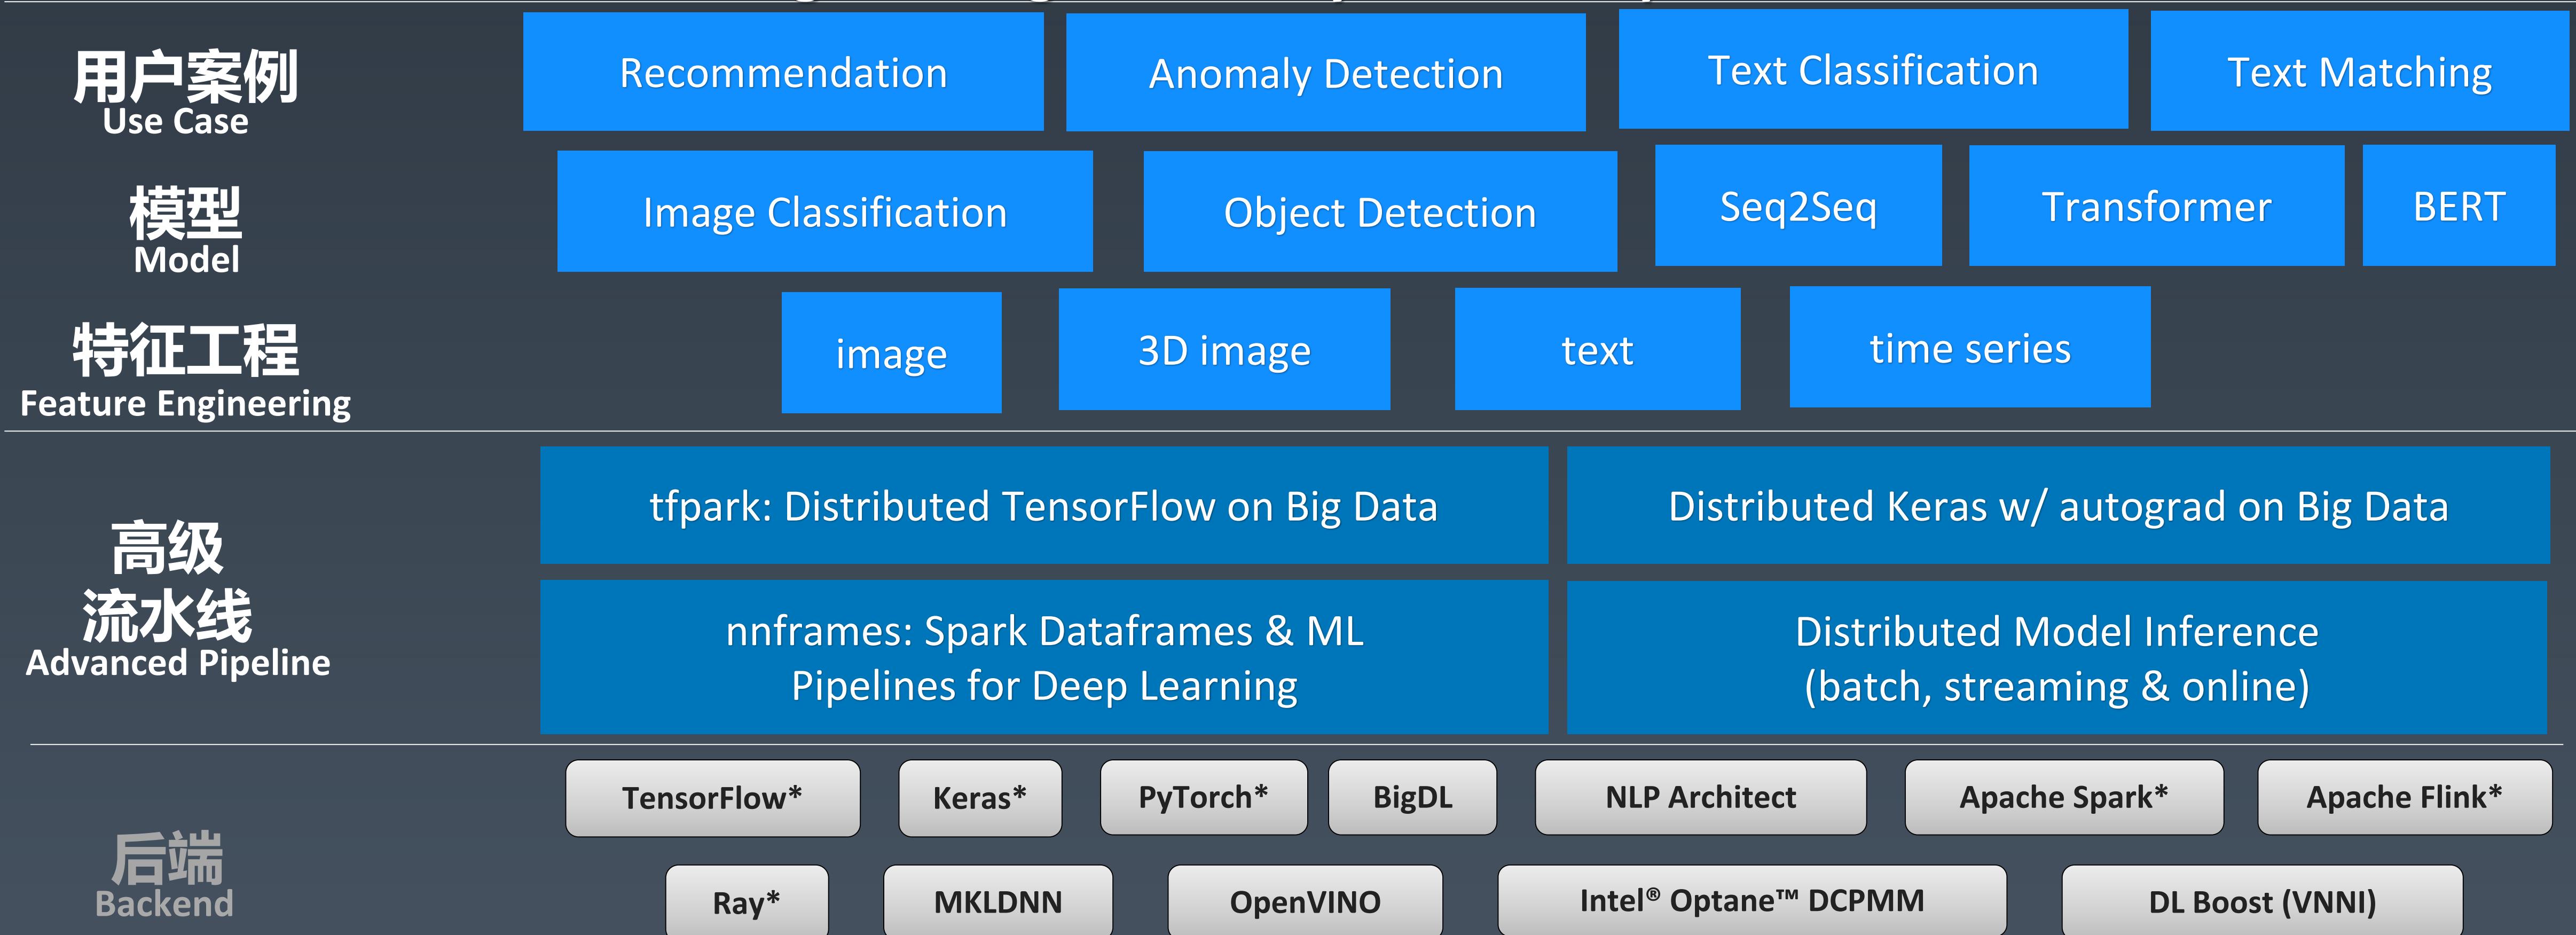
BigDL 分布式训练



Analytics Zoo

统一的大数据分析+人工智能平台

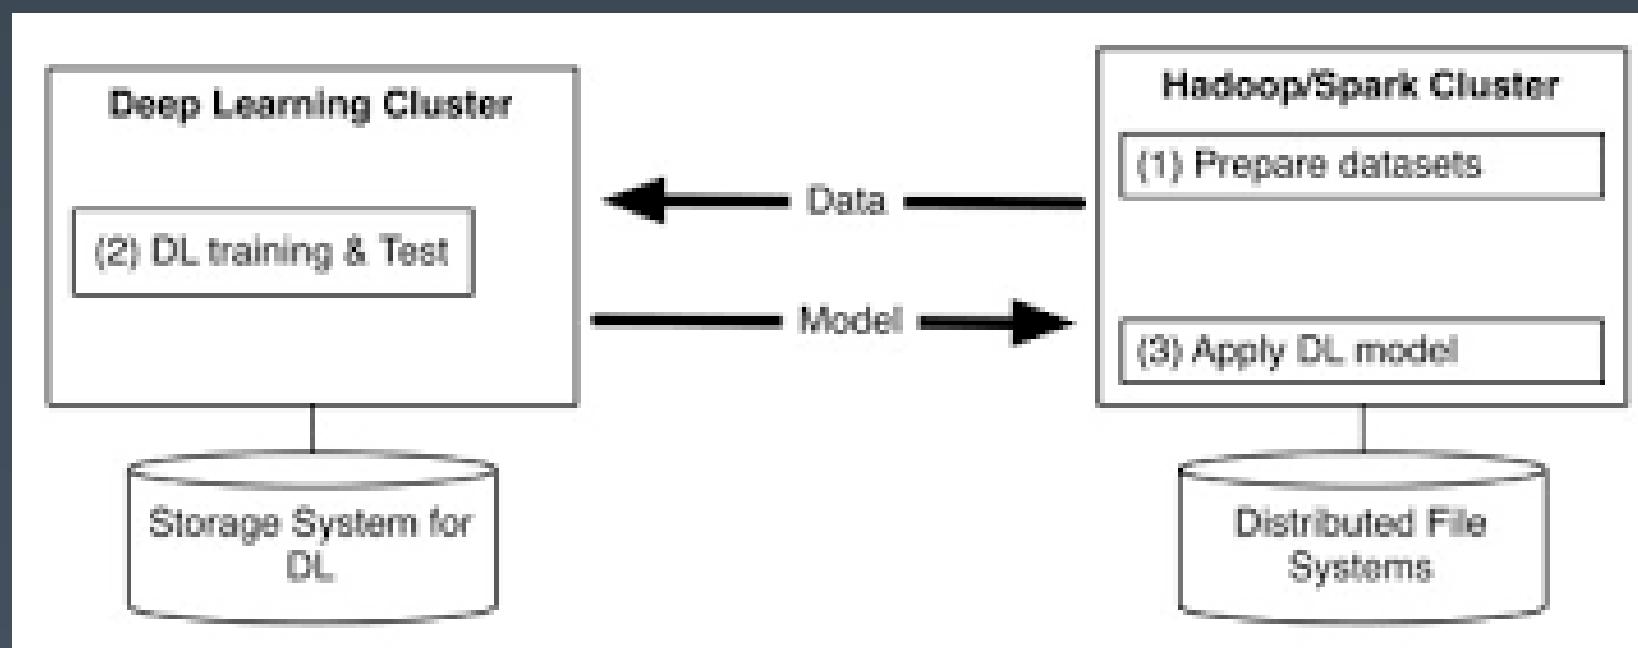
Integrated Big Data Analytics and AI platform



<https://github.com/intel-analytics/analytics-zoo>

基于 Spark* 的分布式 TensorFlow* 流水线

- 无缝集成 TensorFlow* 与 Spark*
- & BigDL*
- 比 TensorFlow * built with MKL 更好的性能
- 易于使用的分布式API，很少的代码改动



在两个集群中迁移数据和模型的不便

```
#pyspark code  
train_rdd = spark.hadoopFile(...).map(...)  
dataset = TFDataset.from_rdd(train_rdd,...)
```

用 PySpark* 载入数据以及预处理数据或特征工程

```
#tensorflow code  
import tensorflow as tf  
slim = tf.contrib.slim  
images, labels = dataset.tensors  
with slim.arg_scope(lenet.lenet_arg_scope()):  
    logits, end_points = lenet.lenet(images, ...)  
loss = tf.reduce_mean( \  
    tf.losses.sparse_softmax_cross_entropy( \  
        logits=logits, labels=labels))
```

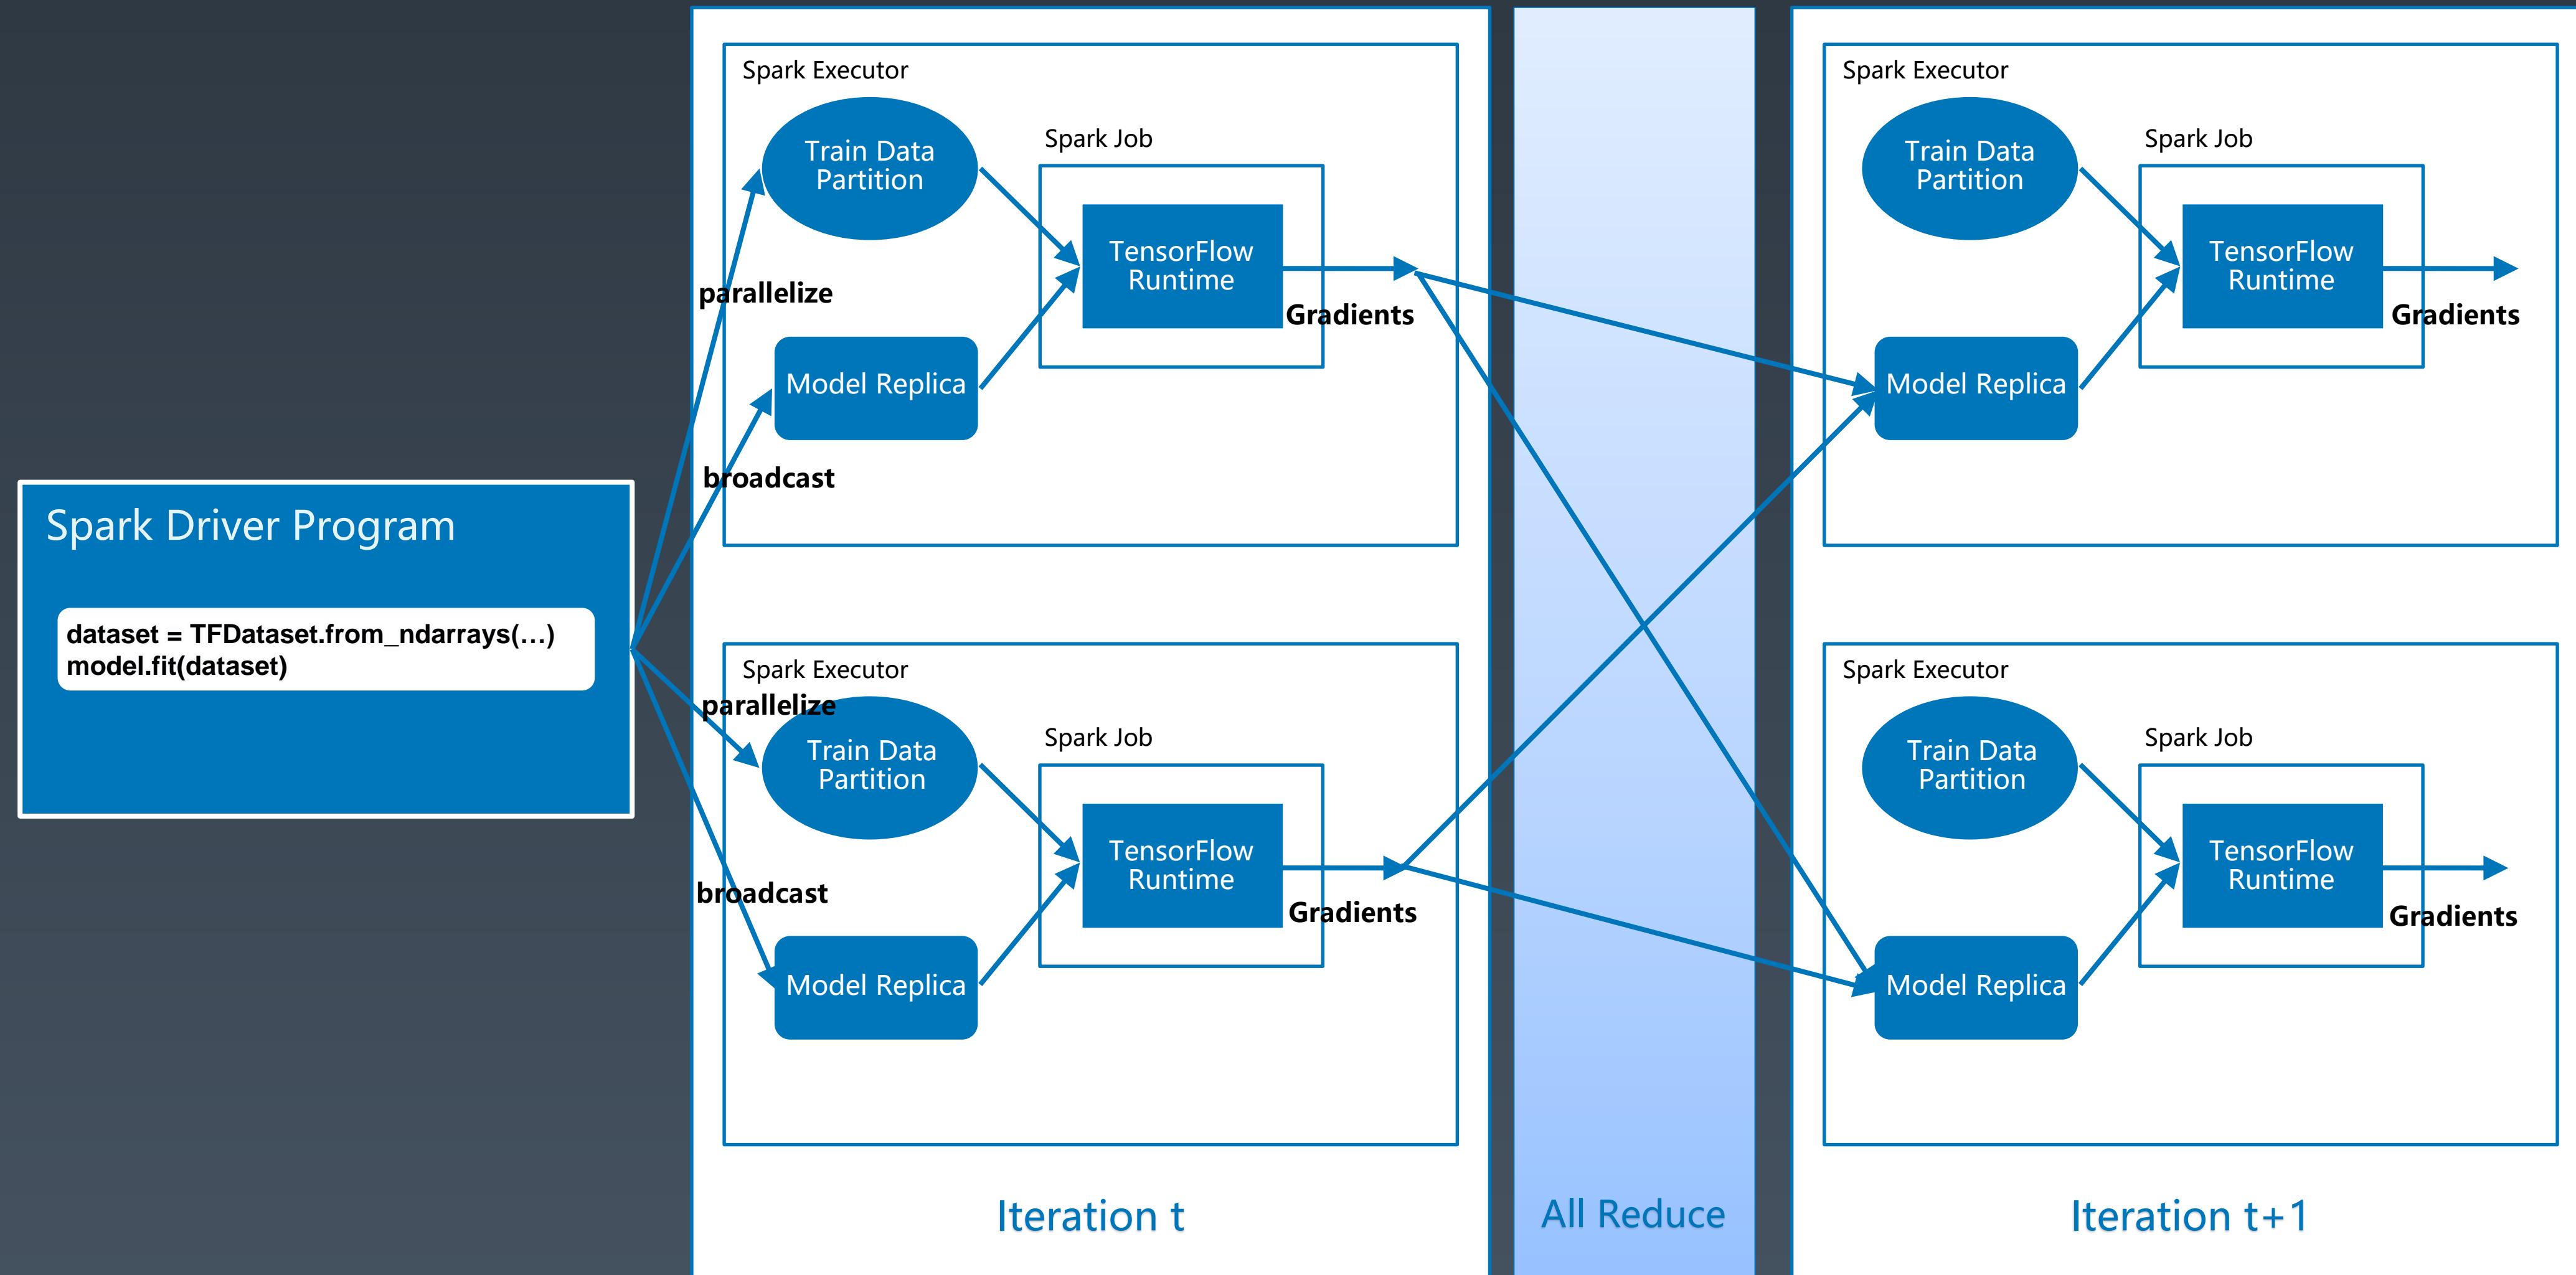
用 TensorFlow* 或 Keras* 定义深度学习模型

```
#distributed training on Spark  
optimizer = TFOptimizer.from_loss(loss, Adam(...))  
optimizer.optimize(end_trigger=MaxEpoch(5))
```

在 Spark* 上分布式训练或者推理

在 PySpark 程序中内嵌 TensorFlow 代码

基于 Spark* 的分布式 TensorFlow* 流水线



基于Spark* Dataframe & ML 流水线的深度学习

```
#Spark dataframe transformations  
parquetfile = spark.read.parquet(...)  
train_df = parquetfile.withColumn(...)
```

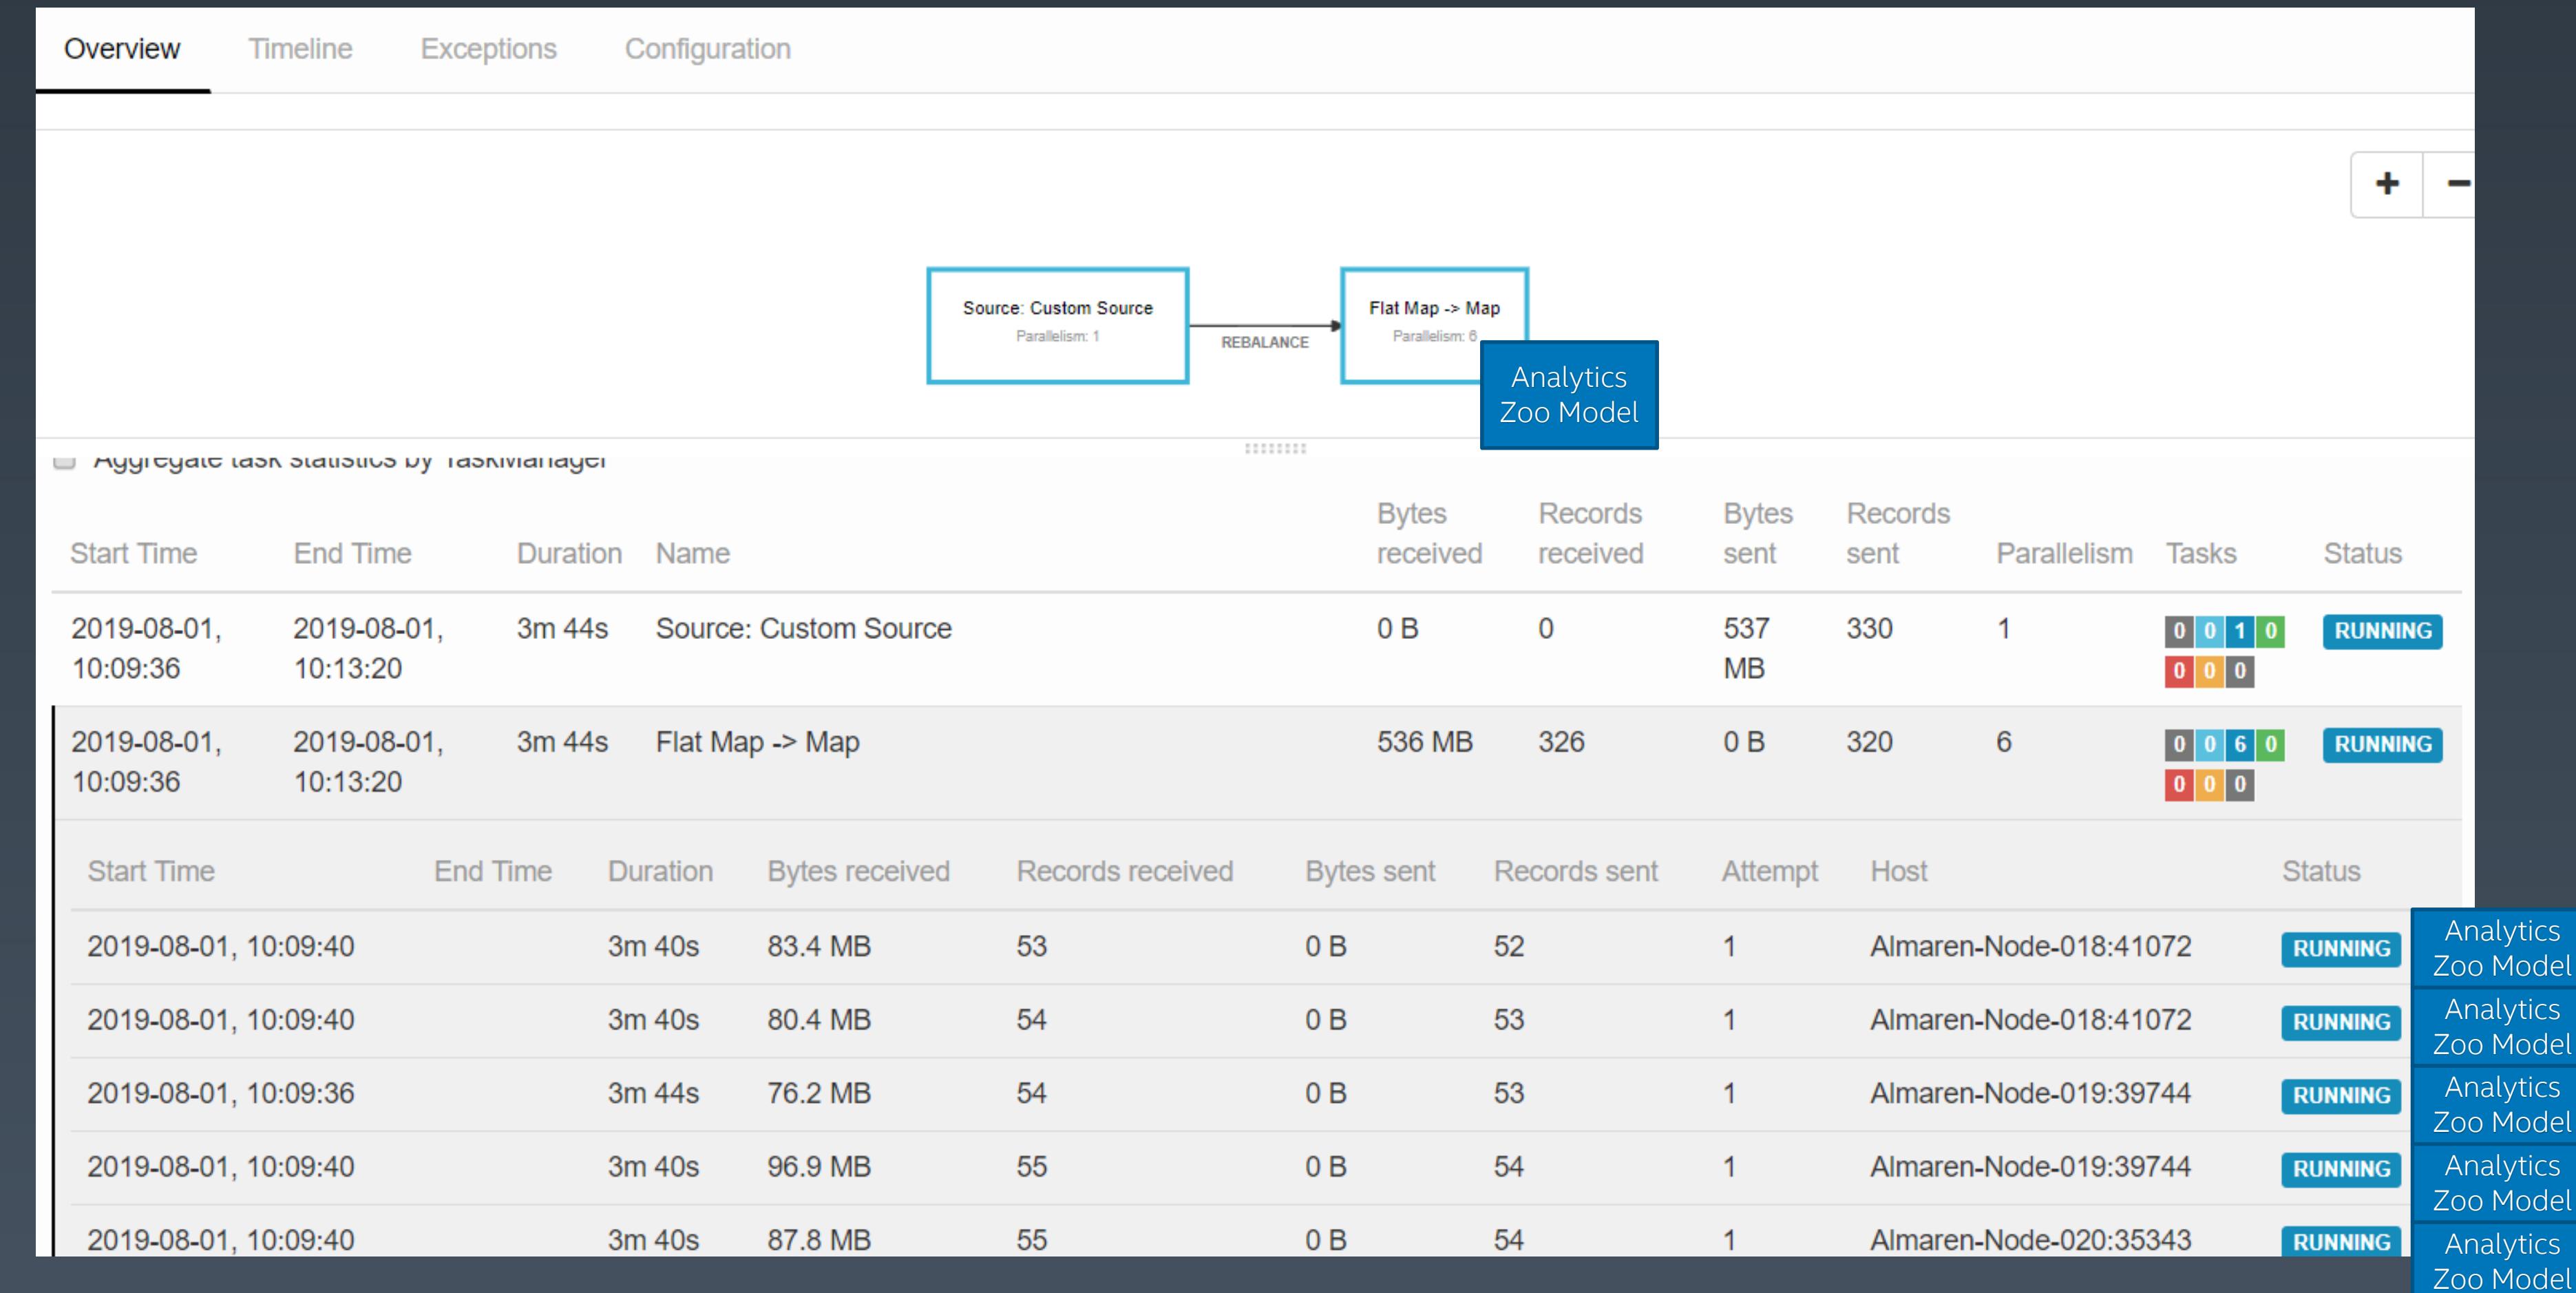
```
#Keras API  
model = Sequential()  
    .add(Convolution2D(32, 3, 3, activation='relu', input_shape=...)) \  
    .add(MaxPooling2D(pool_size=(2, 2))) \  
    .add(Flatten()).add(Dense(10, activation='softmax'))
```

```
#Spark ML pipeline  
Estimator = NNEstimator(model, CrossEntropyCriterion()) \  
    .setLearningRate(0.003).setBatchSize(40).setMaxEpoch(5) \  
    .setFeaturesCol("image")  
nnModel = estimator.fit(train_df)
```

在 Spark* Dataframe 和 ML 流水线中，直接支持深度神经网络模型

分布式、实时(流式)模型推理流水线

- 纯Java或Python API
- 支持Spark*
Streaming, Flink*,
Storm*, Kaka*等
- 支持Web Services
- 使用OpenVINO和DL
Boost(VNNI) 加速



POJO风格的模型推理API

- 支持加载各种框架的模型
 - BigDL
 - Caffe*
 - TensorFlow*
 - OpenVINO
- 线程安全及自动扩展
 - 多工作模型共享参数

```
import com.intel.analytics.zoo.pipeline.inference.AbstractInferenceModel;

public class TextClassification extends AbstractInferenceModel {
    public RankerInferenceModel(int concurrentNum) {
        super(concurrentNum);
    }
    ...
}

public class ServingExample {
    public static void main(String[] args) throws IOException {
        TextClassification model = new TextClassification();
        model.load(modelPath, weightPath);

        texts = ...
        List<JTensor> inputs = preprocess(texts);
        for (JTensor input : inputs) {
            List<Float> result =
                model.predict(input.getData(), input.getShape());
            ...
        }
    }
}
```

使用OpenVINO加速模型推理

- 支持Image Classification 和 Object Detection等
- 支持加载TensorFlow*模型
- 支持直接加载OpenVINO IR
- 支持模型动态Optimize及 Calibrate

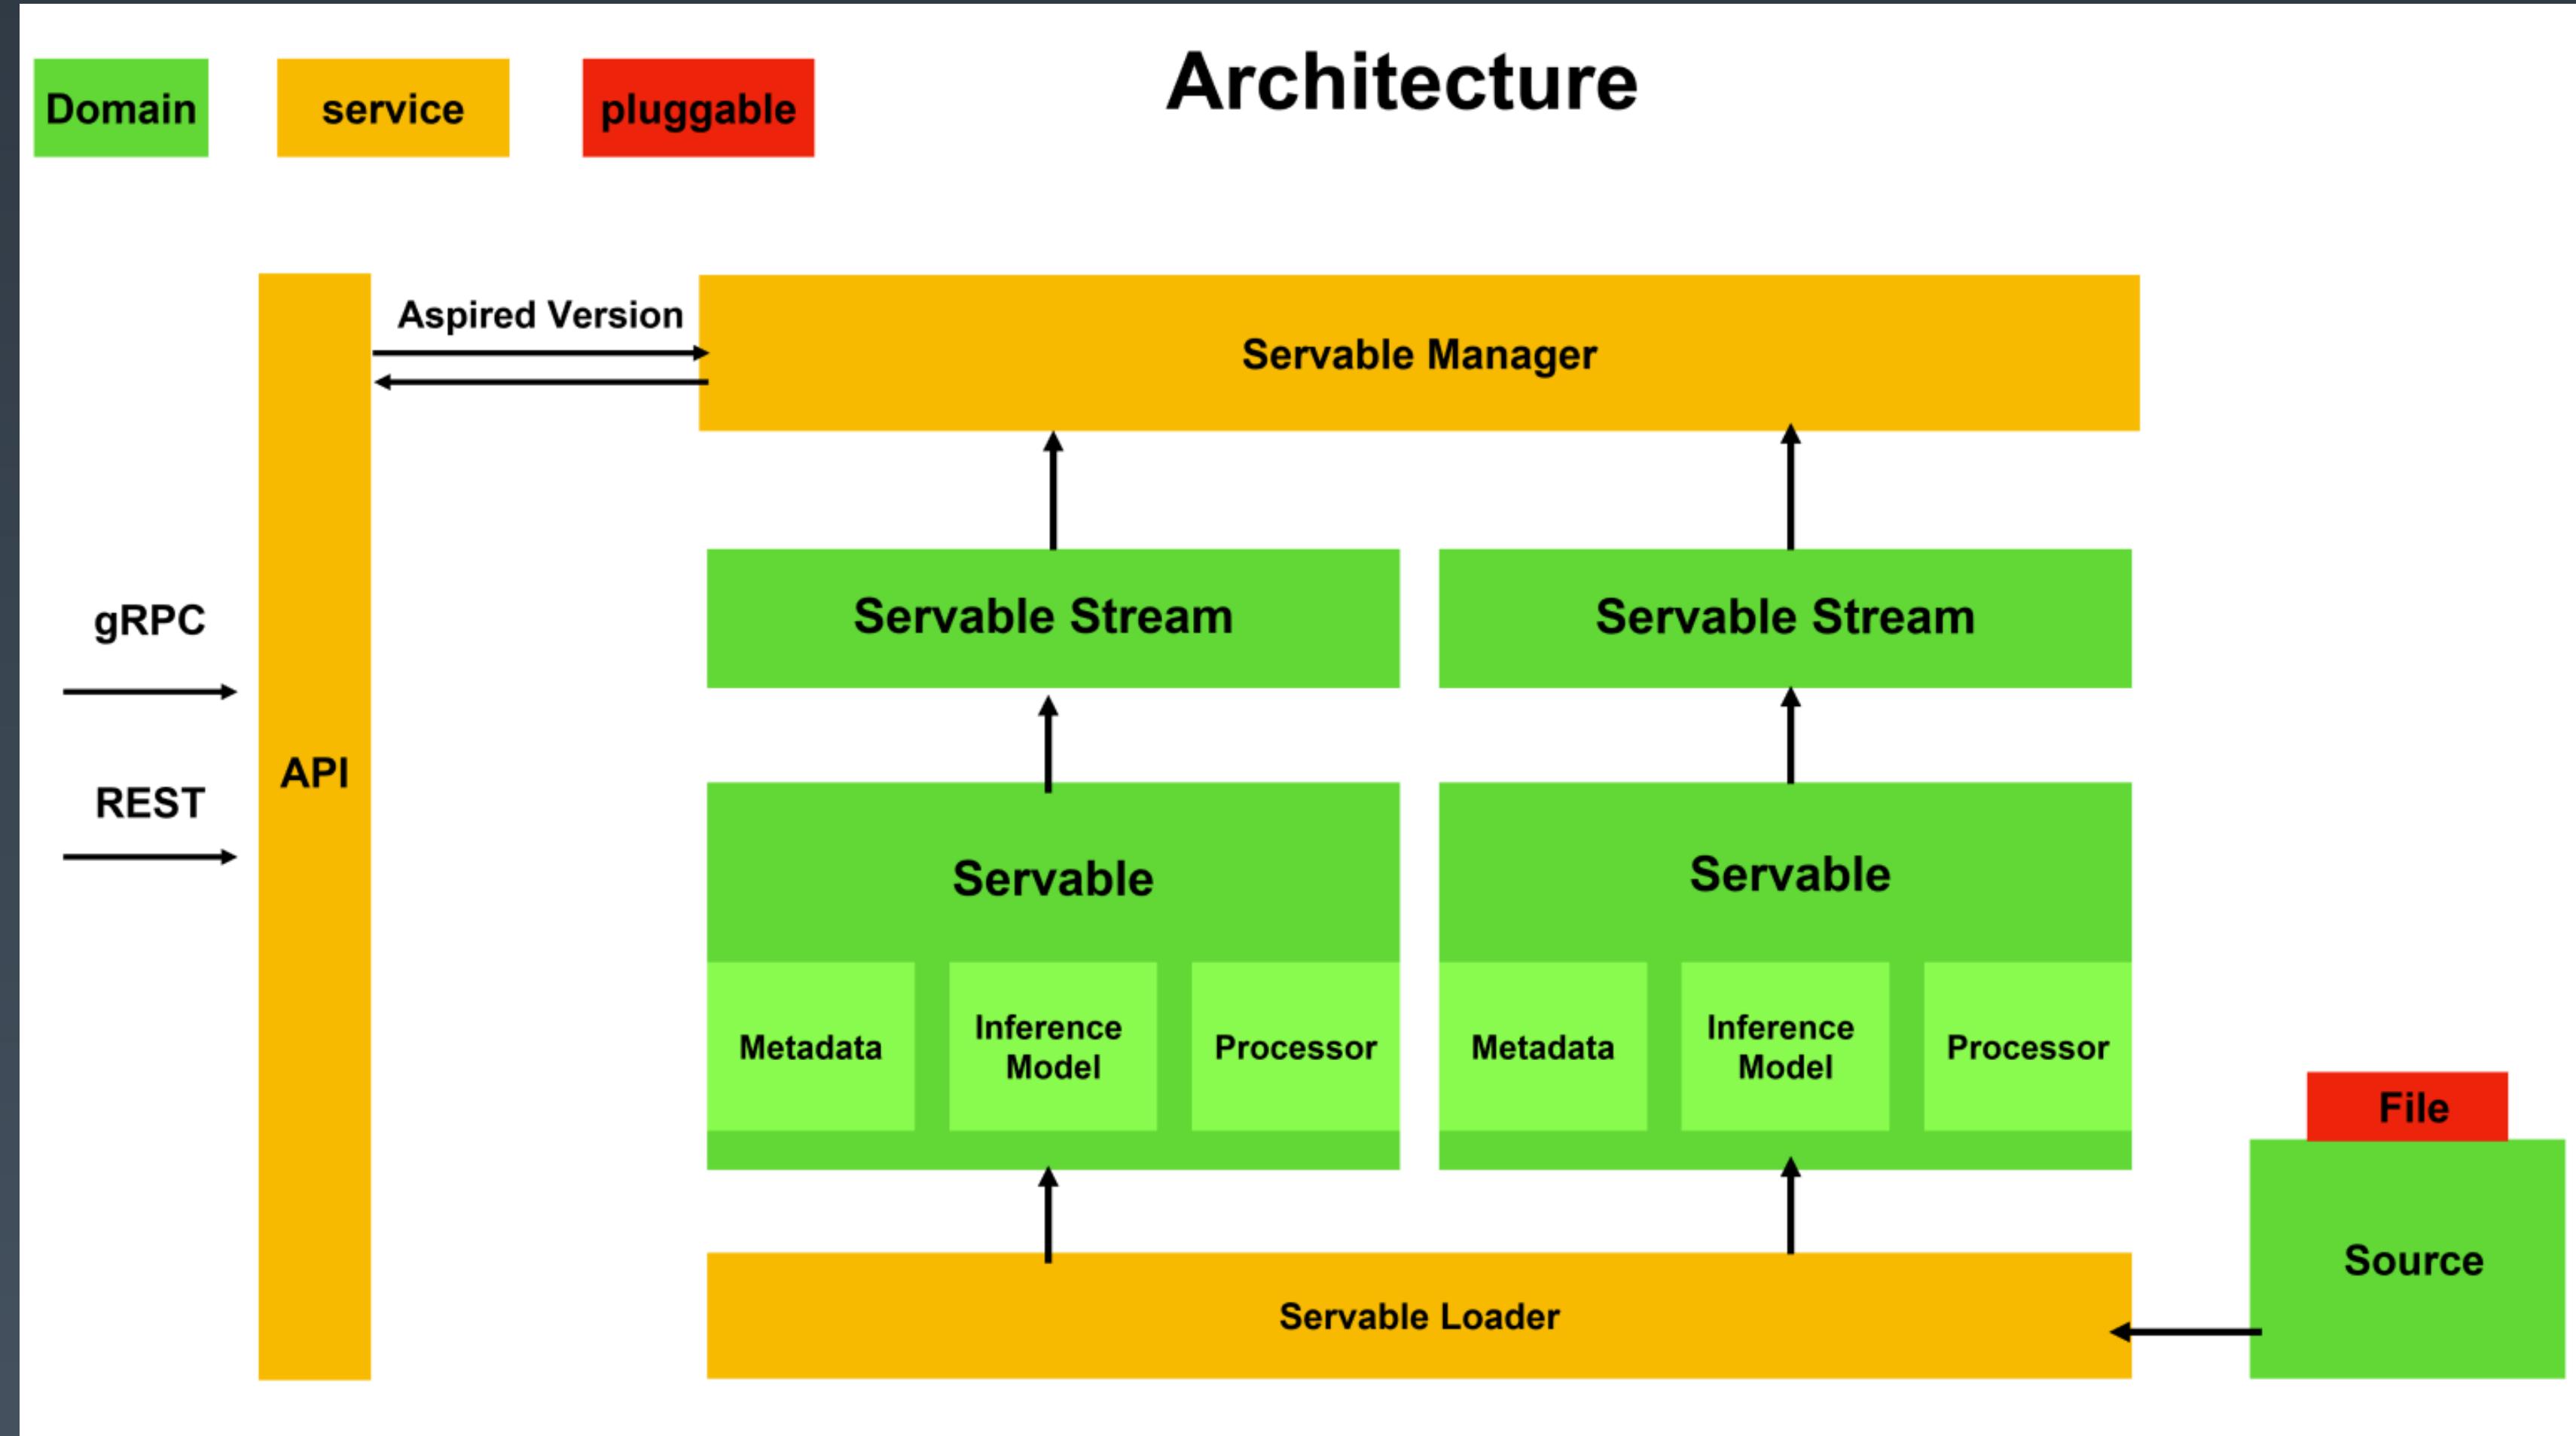
```
from zoo.common.nncontext import init_nncontext
from zoo.feature.image import ImageSet
from zoo.pipeline.inference import InferenceModel

sc = init_nncontext("OpenVINO Object Detection Inference Example")
images = ImageSet.read(options.img_path, sc,
                      resize_height=600,
                      resize_width=600).get_image().collect()
input_data = np.concatenate([image.reshape((1, 1) + image.shape) for
                            image in images], axis=0)

model = InferenceModel()
model.load_tf(options.model_path, backend="openvino",
              model_type=options.model_type)
predictions = model.predict(input_data)

# Print the detection result of the first image.
print(predictions[0])
```

Analytics Zoo Web Serving



Analytics Zoo Cluster Serving 使分布式推理更加简单

部署

- ✓ 一个本地节点或者一个Docker容器
- ✓ 已有的 Yarn/Spark/Flink (or K8s) 集群

使用



1 一条命令:

- 启动Docker容器以及Zoo Cluster Serving

- 此命令指定:
 - 输入 和 输出 的队列名字
 - 模型 的文件路径
 - 预/后处理 的文件路径
 - 集群 的访问路径



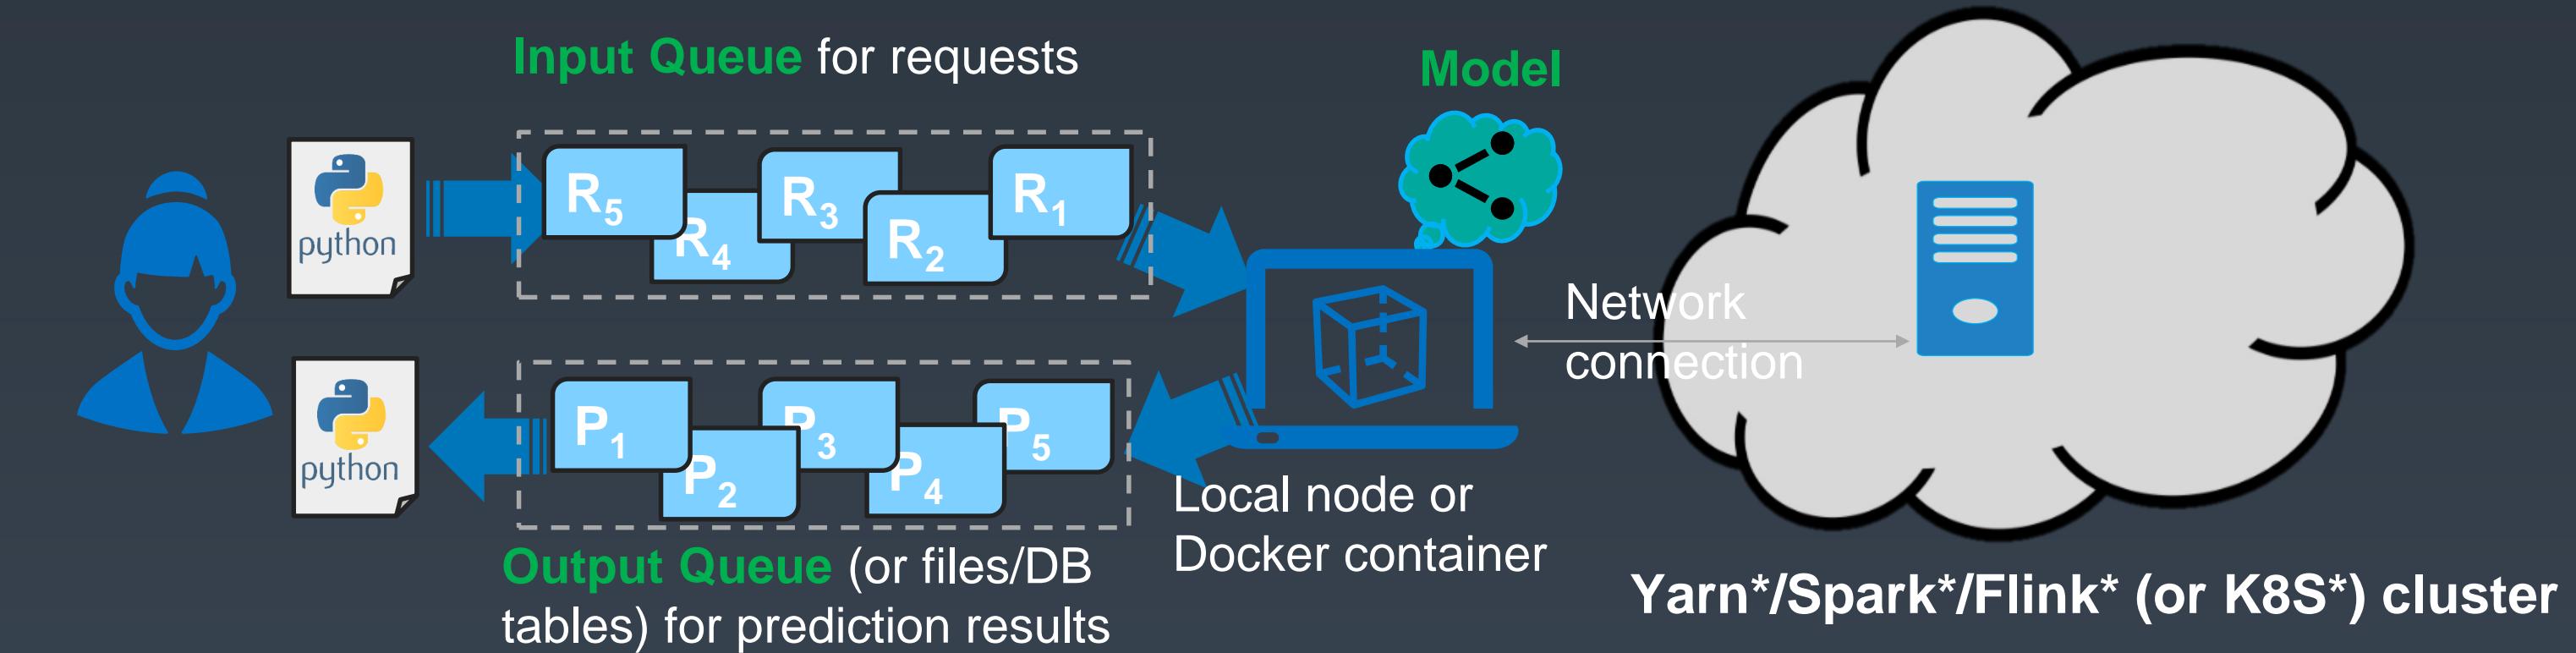
2 一个简单的Python脚本:

- 将请求数据发送到 Input Queue
- 从 Output Queue (或文件/数据库)获得推理结果



3 Analytics Zoo 在集群上自动执行分布式、实时 (流式) 模型推理

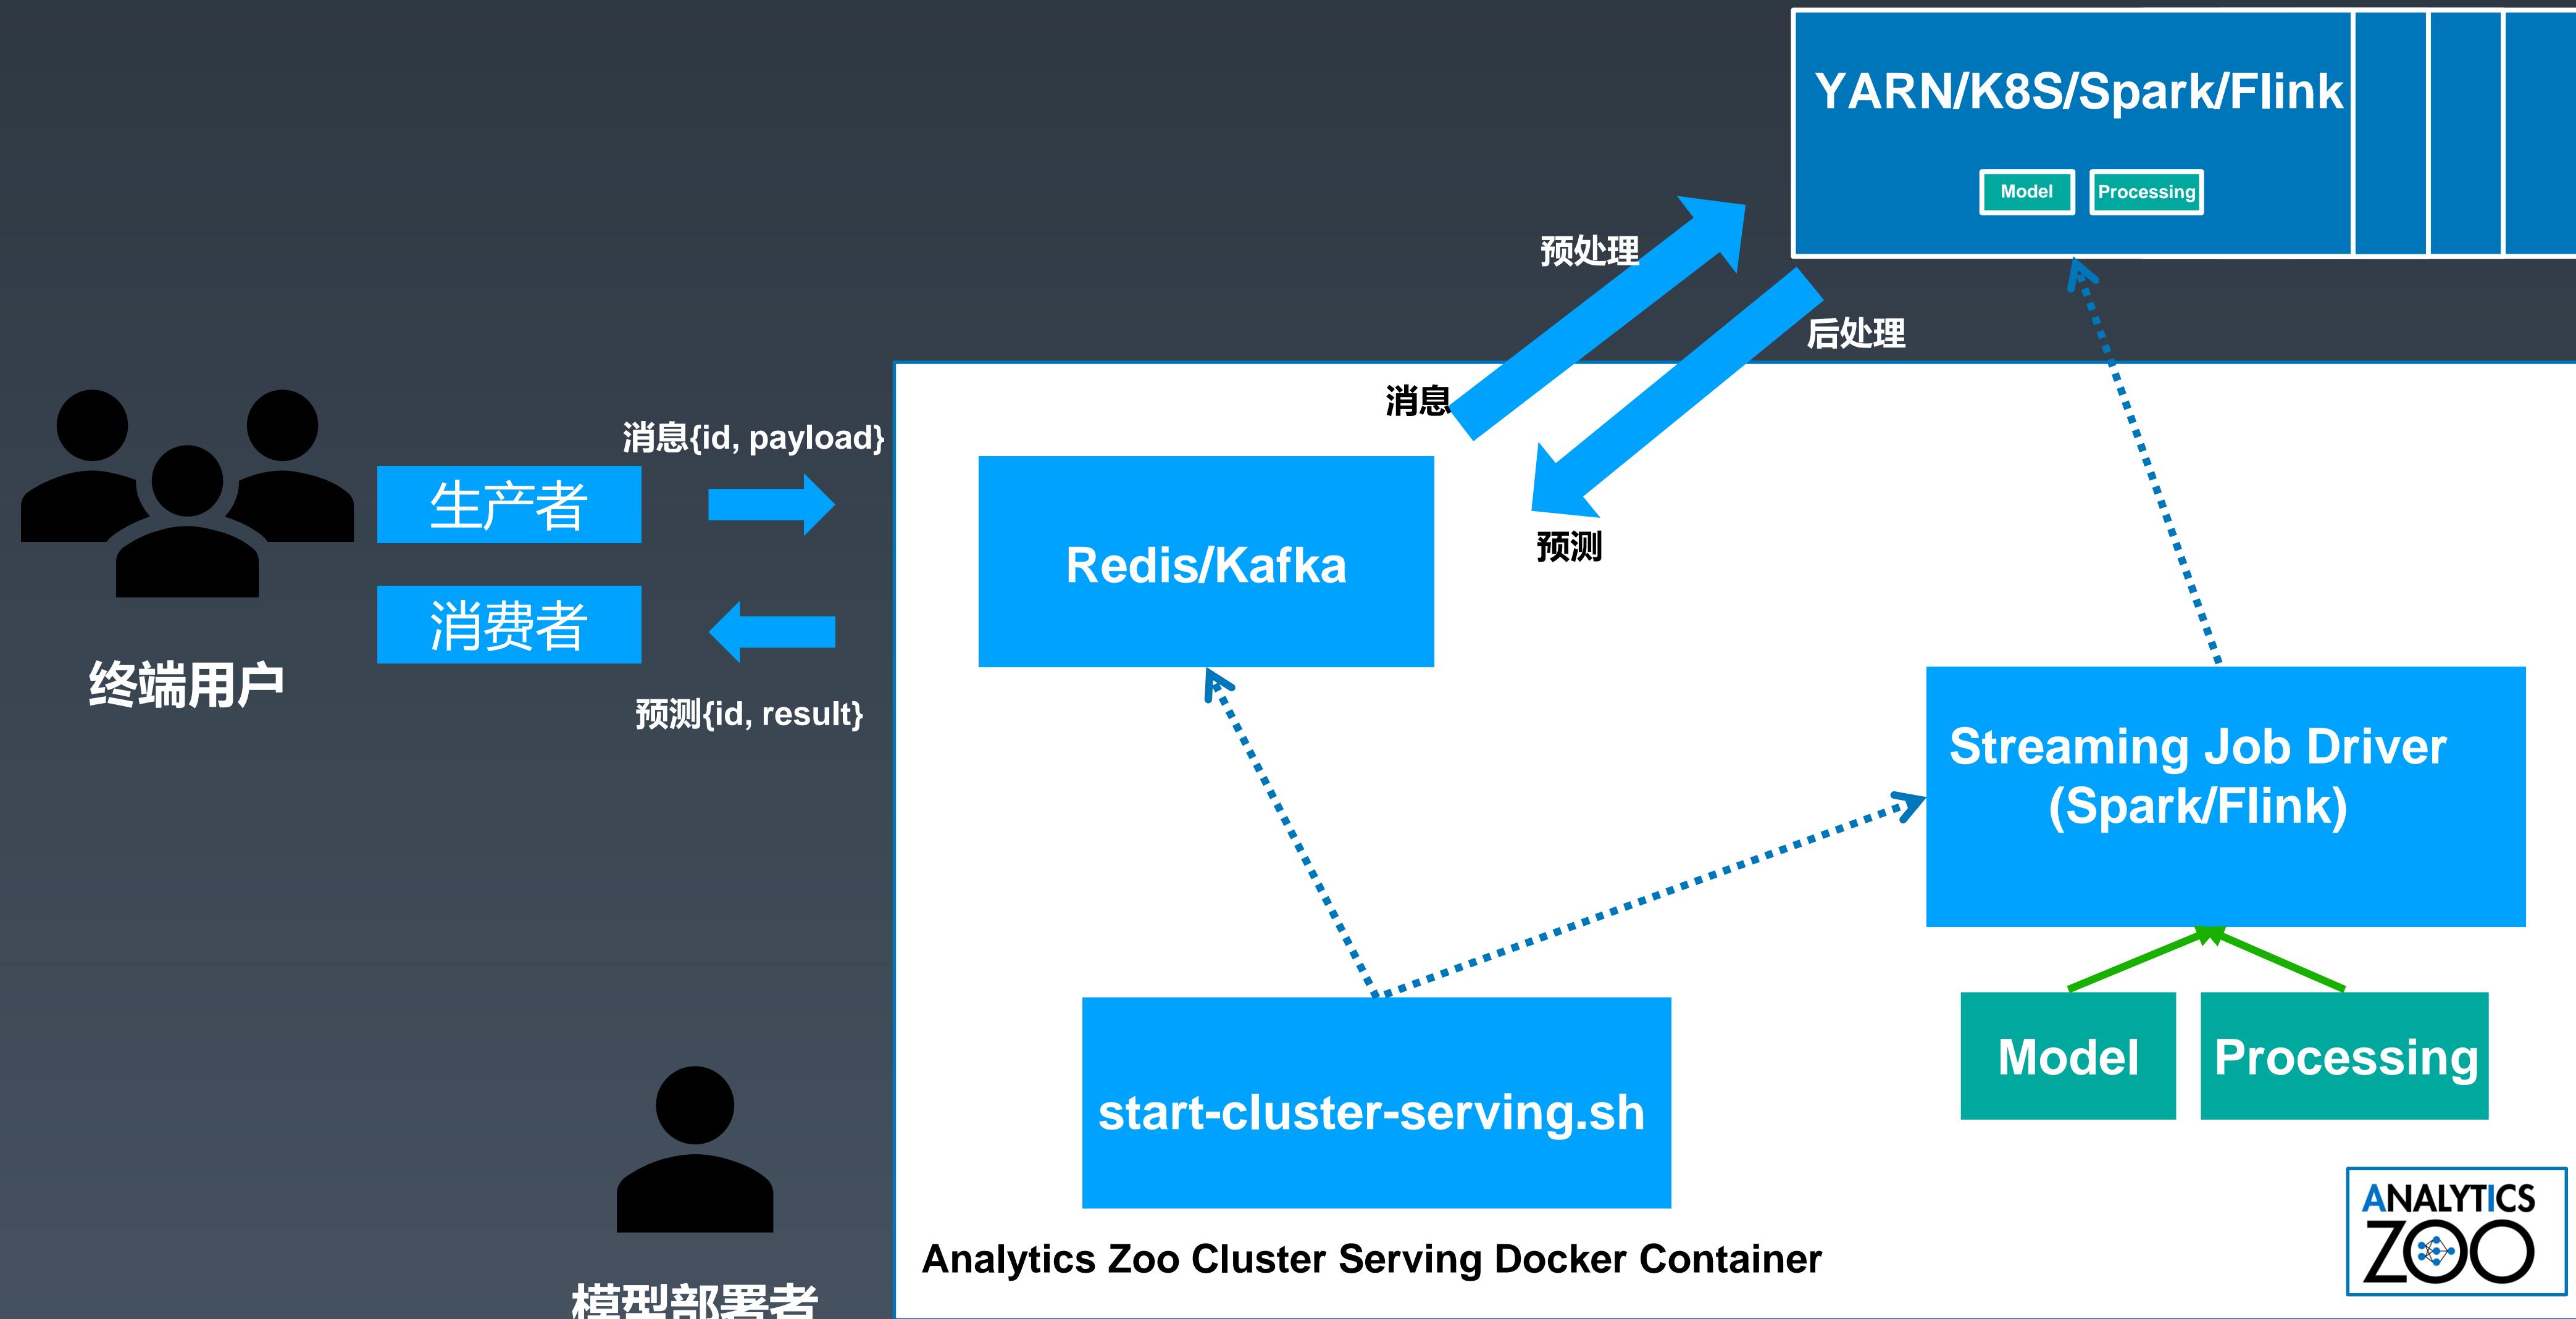
- 支持 TensorFlow*, Keras*, PyTorch*, Caffe*, BigDL 和 OpenVINO 的模型, 可使用 Int8 加速
- 通过Spark* Streaming 或 Flink* 线性扩展



✓ 可扩展的分布式推理由Analytics Zoo托管

✓ 用户无需为开发和部署复杂的分布式推理方案而费心

Analytics Zoo Cluster Serving 架构



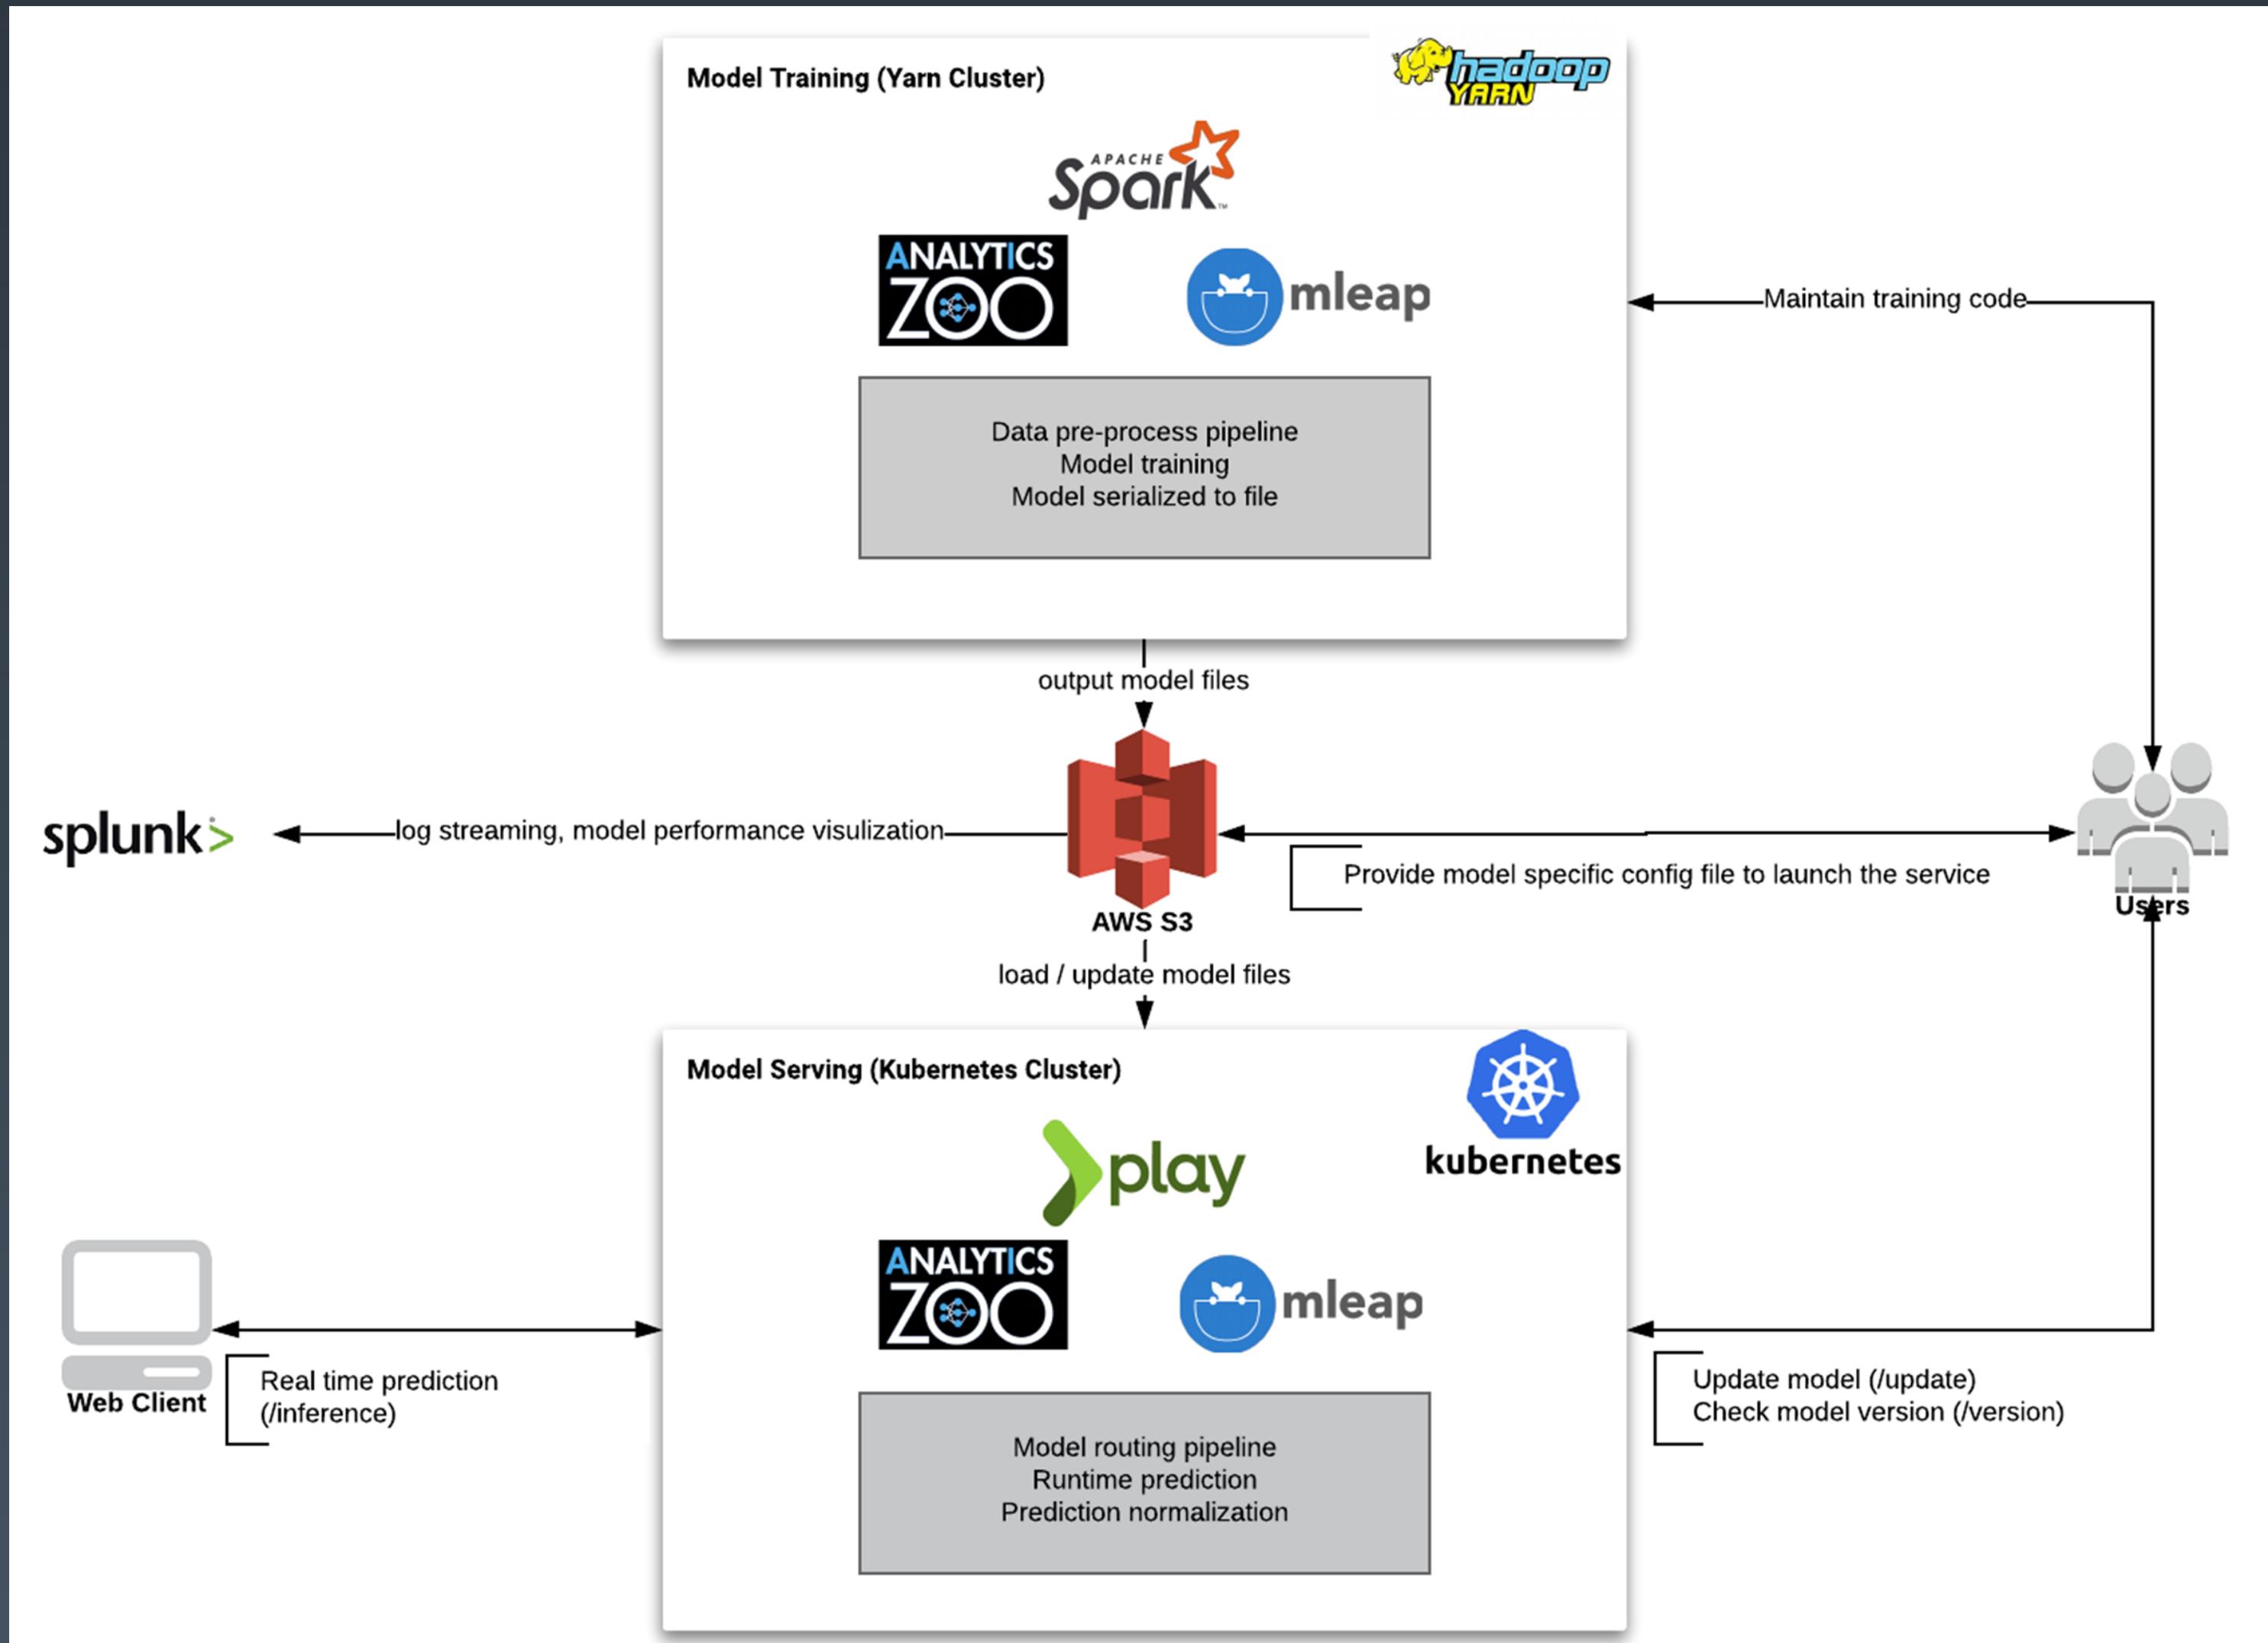
跨行业的端到端客户案例实践



Office DEPOT.

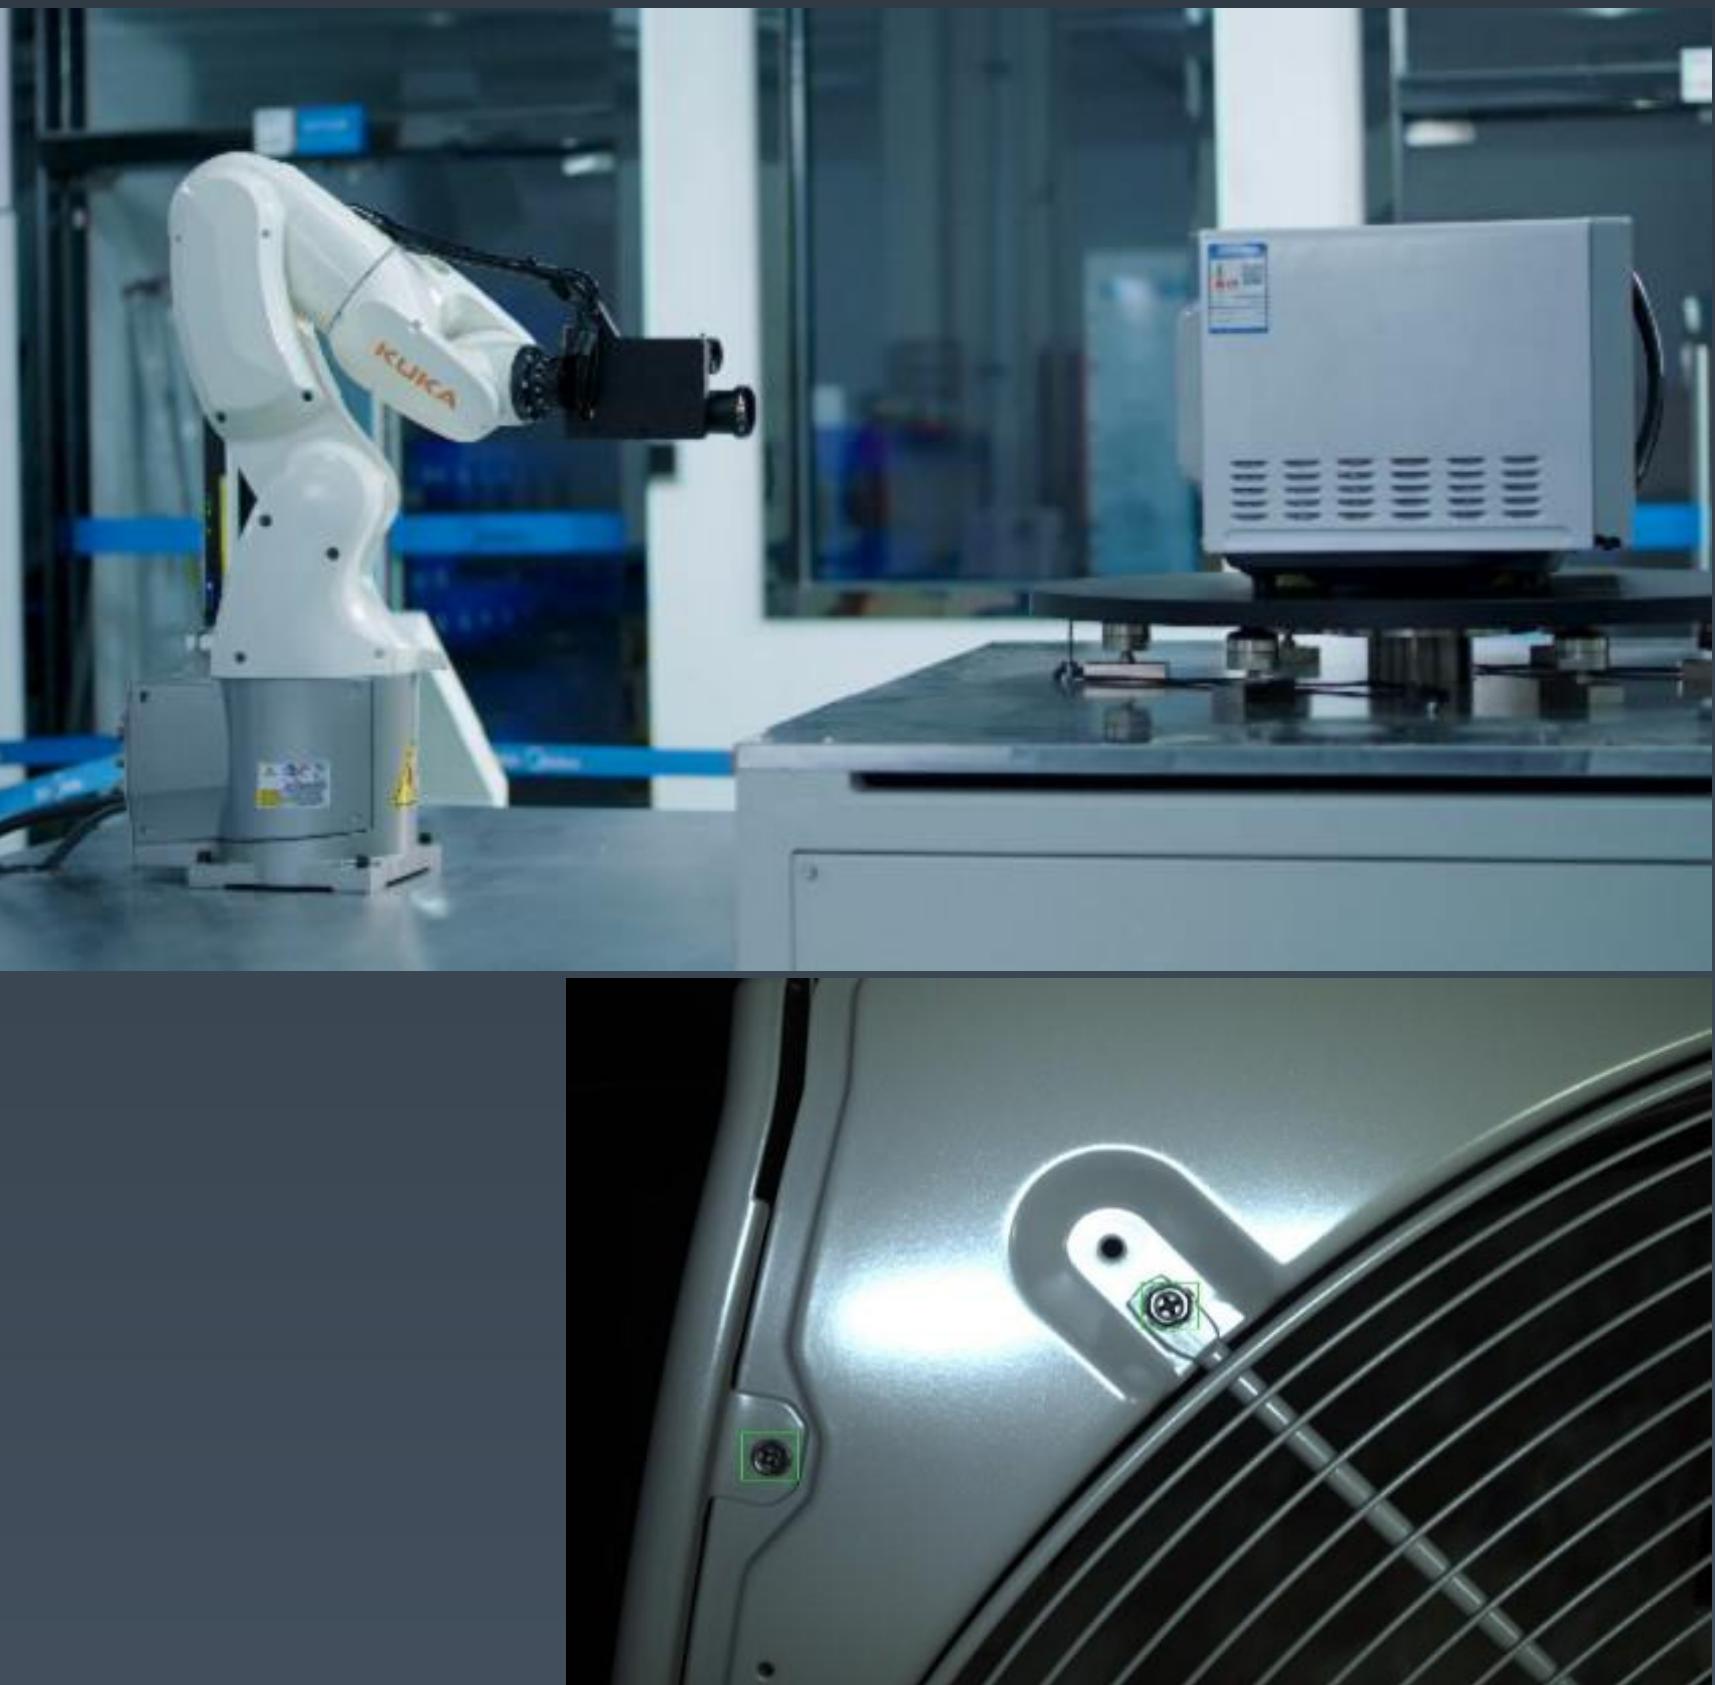
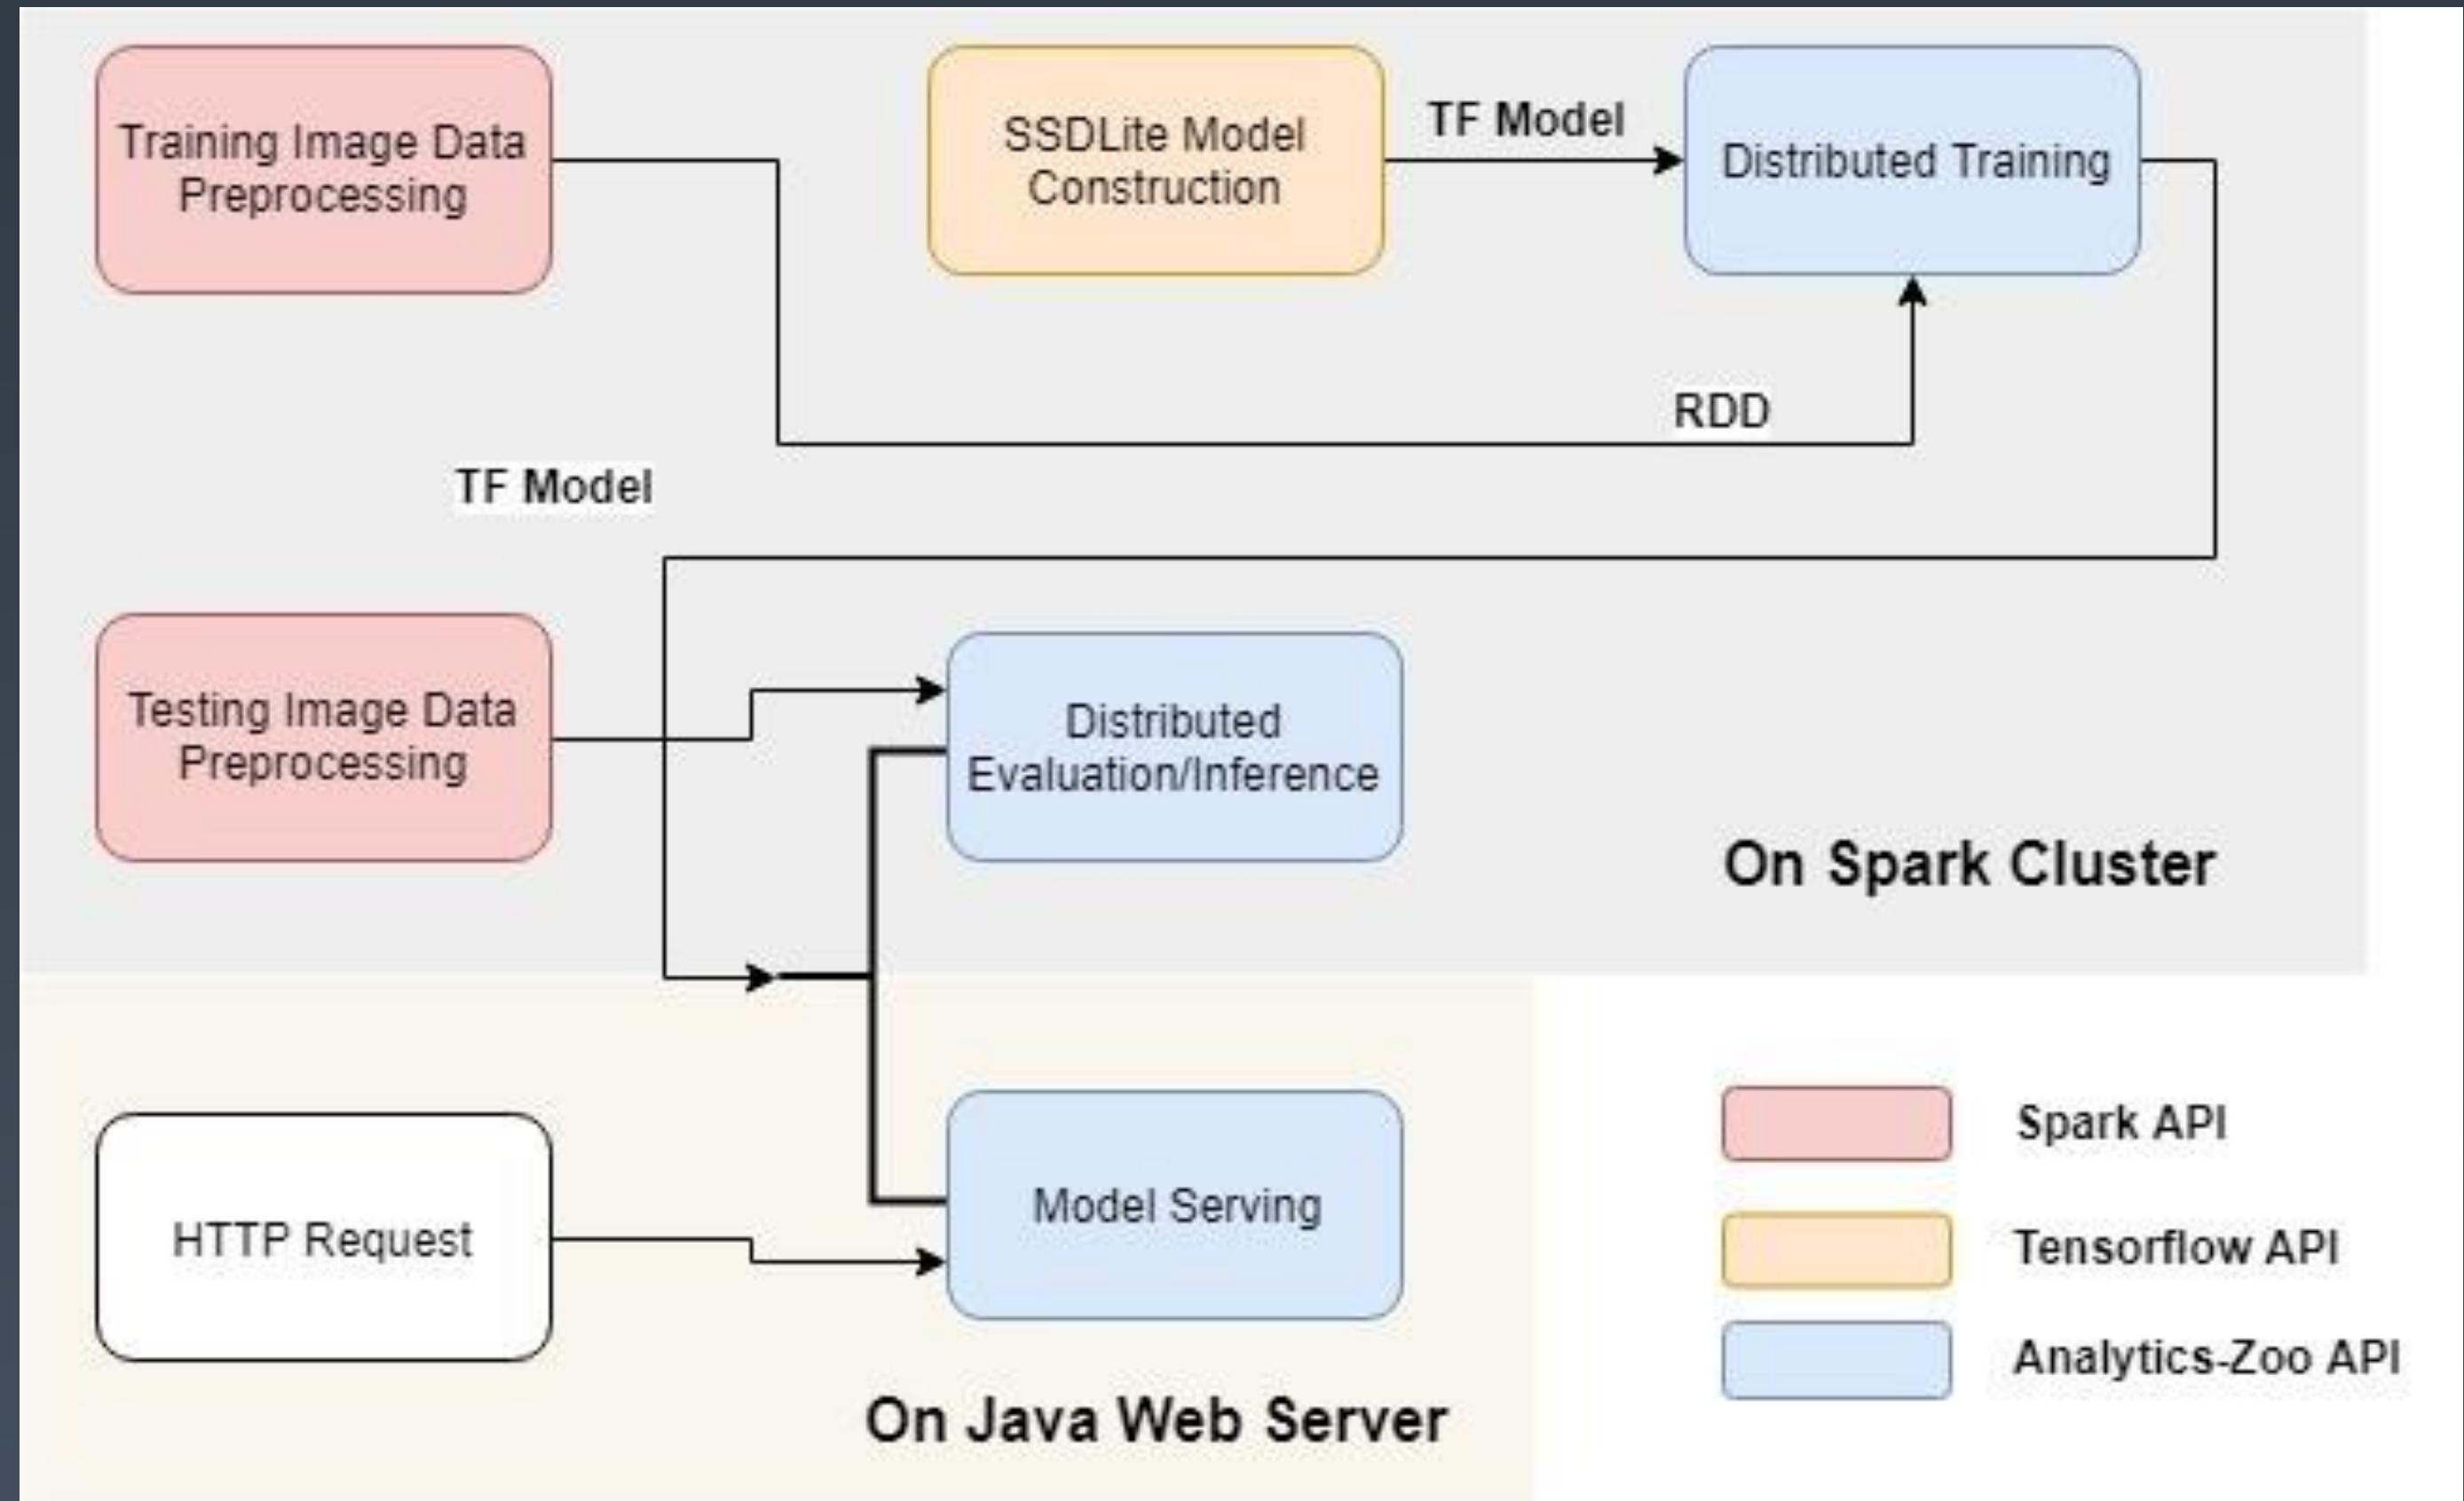


Office Depot*: 基于用户 Session 行为的产品推荐



<https://software.intel.com/en-us/articles/real-time-product-recommendations-for-office-depot-using-apache-spark-and-analytics-zoo-on>
<https://conferences.oreilly.com/strata/strata-ca-2019/public/schedule/detail/73079>

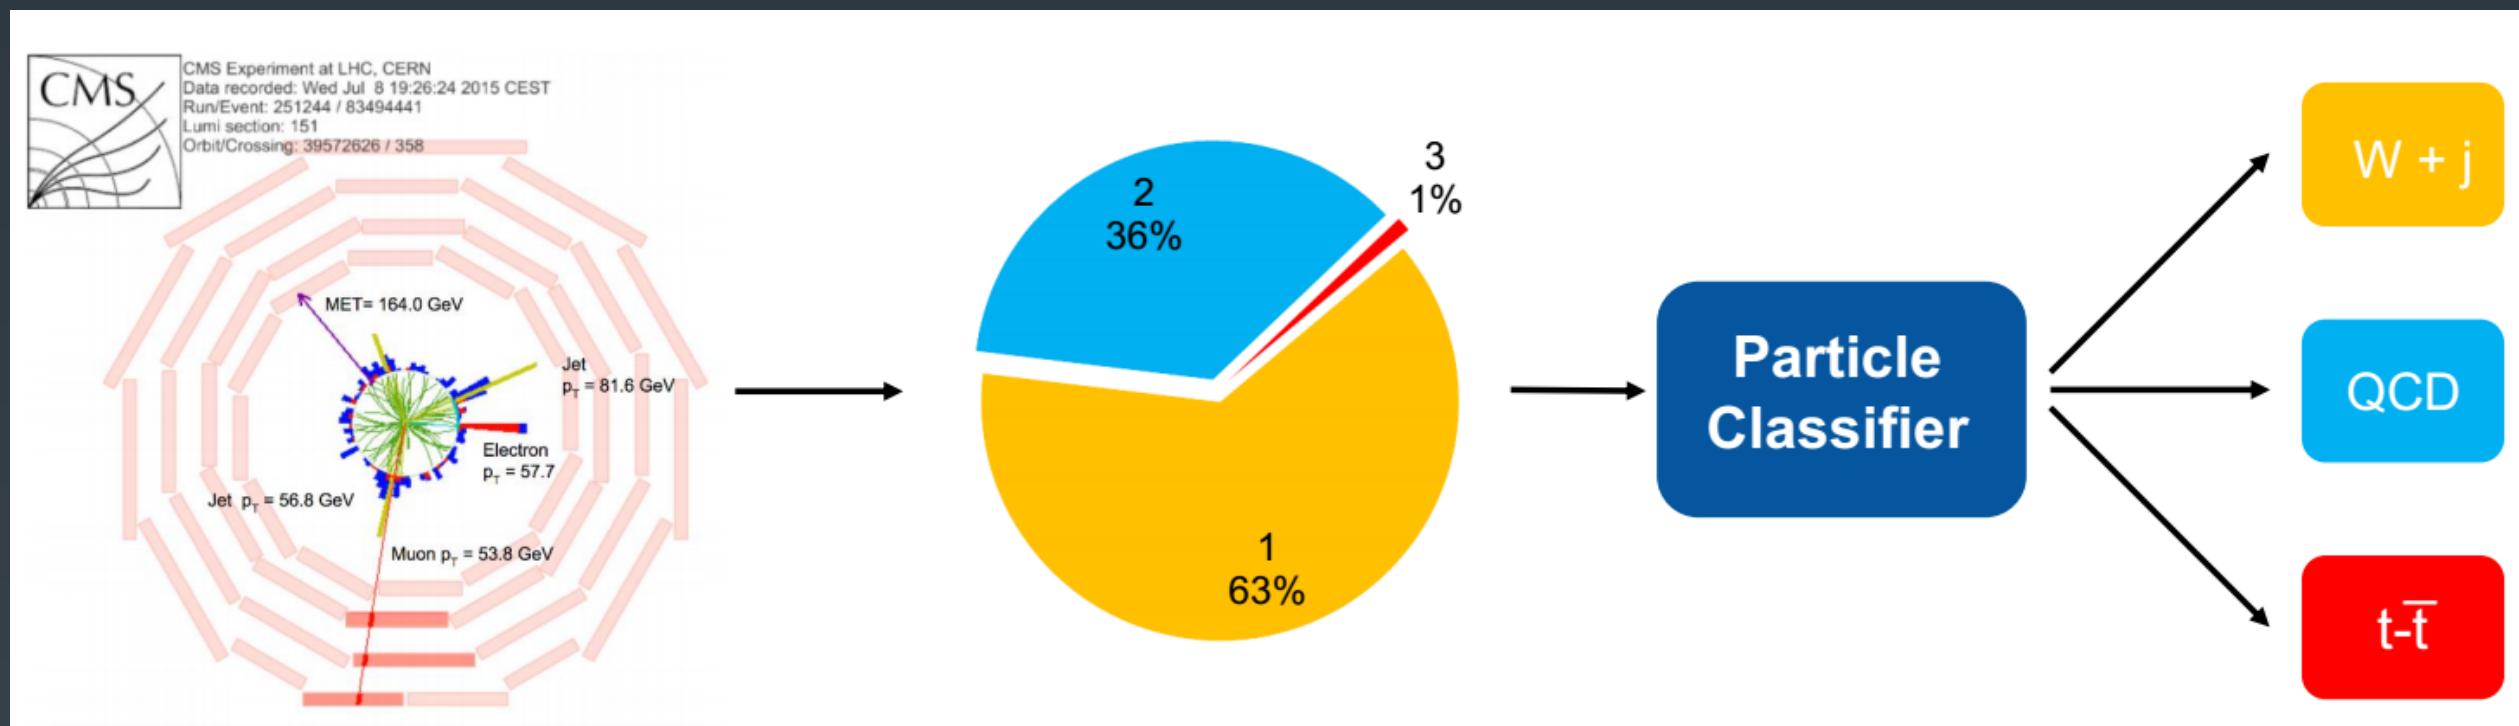
美的*: 工业视觉检测云平台



<https://software.intel.com/en-us/articles/industrial-inspection-platform-in-midea-and-kuka-using-distributed-tensorflow-on-analytics>

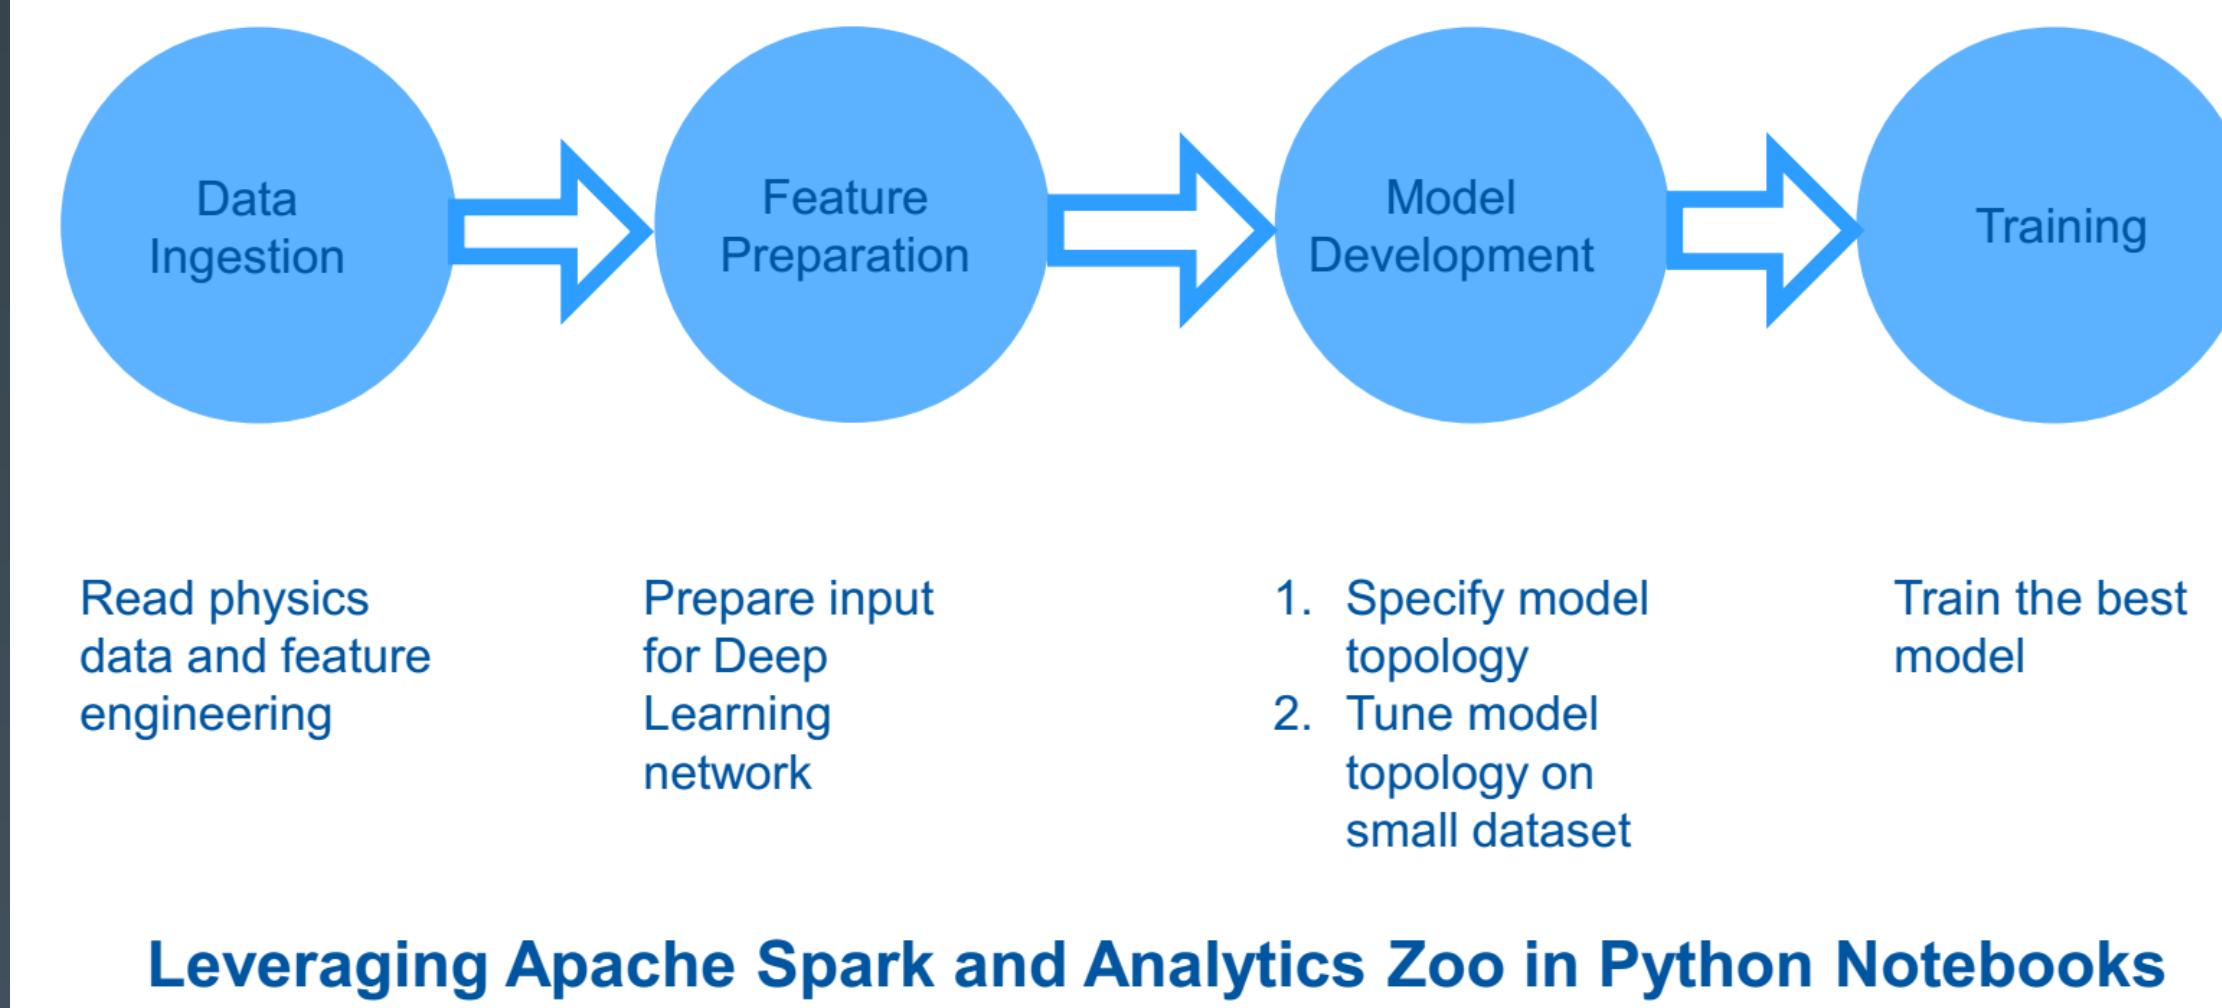
<https://www.intel.cn/content/www/cn/zh/analytics/artificial-intelligence/midea-case-study.html>

CERN*: 基于深度学习的高能物理粒子事件分类

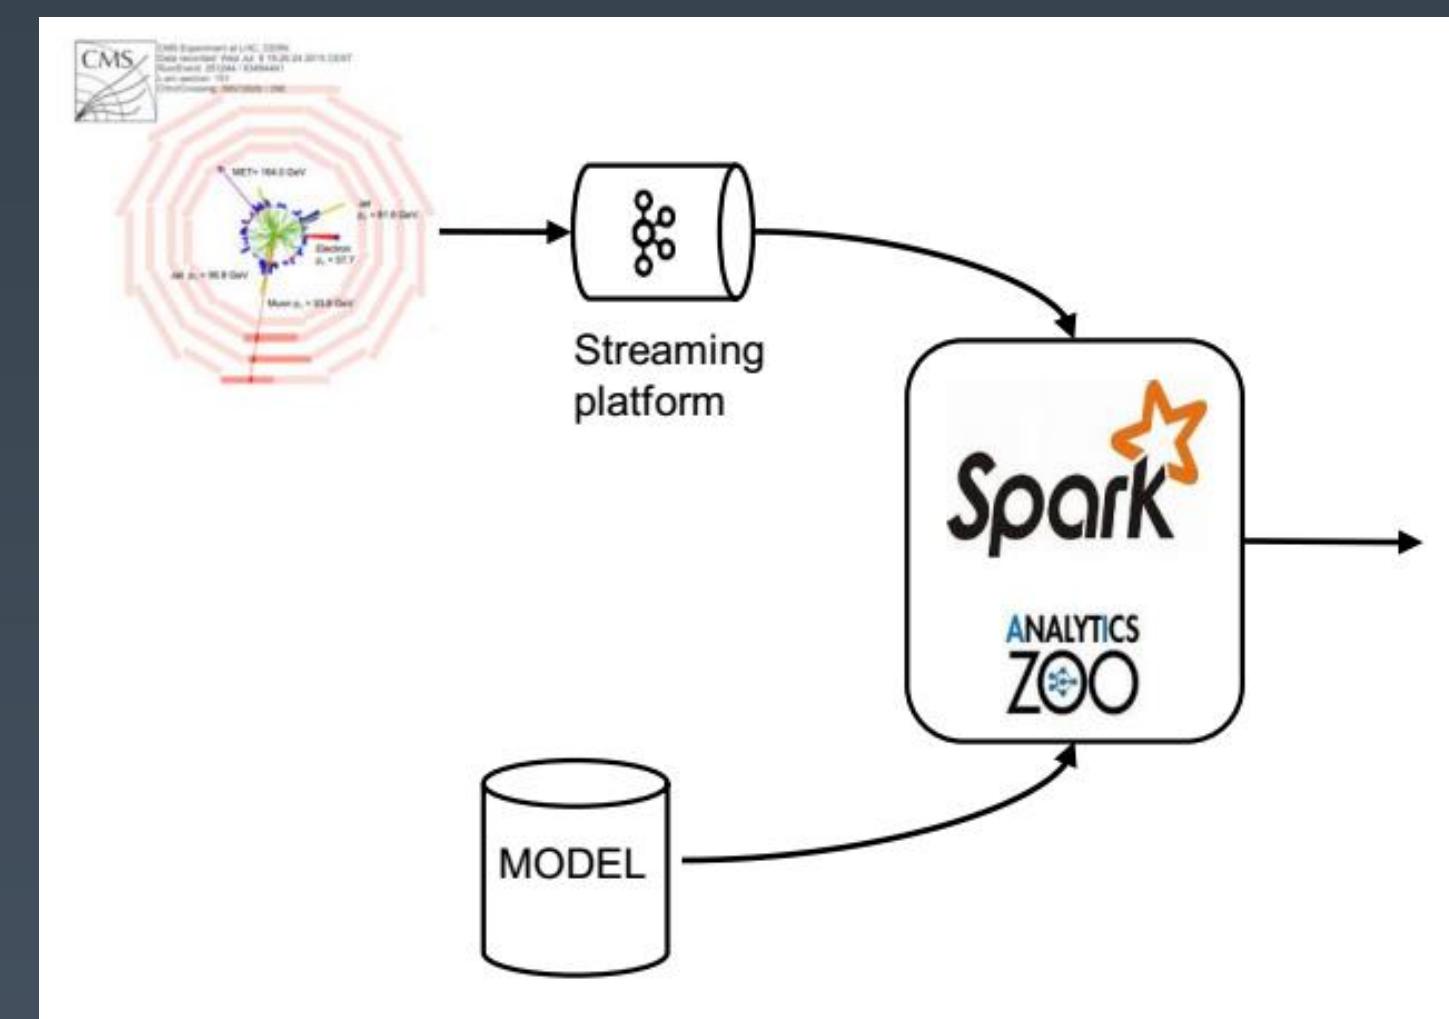


高能物理数据的深度学习流水线
Deep Learning Pipeline for High Energy Physics

Data Pipeline



使用 Apache Kafka* 和 Spark* 构建模型服务
Model serving using Apache kafka and Spark Streaming



<https://db-blog.web.cern.ch/blog/luca-canali/machine-learning-pipelines-high-energy-physics-using-apache-spark-bigdl>

<https://databricks.com/session/deep-learning-on-apache-spark-at-cerns-large-hadron-collider-with-intel-technologies>

¹ 该CERN的解决方案为原型实现，尚未投入生产部署。

THANKS! | QCon th

LEGAL NOTICES AND DISCLAIMERS

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations, and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products. For more complete information visit intel.com/performance.

Intel does not control or audit the design or implementation of third-party benchmark data or websites referenced in this document. Intel encourages all of its customers to visit the referenced websites or others where similar performance benchmark data are reported and confirm whether the referenced benchmark data are accurate and reflect performance of systems available for purchase.

Optimization notice: Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software, or service activation. Performance varies depending on system configuration. No computer system can be absolutely secure. Check with your system manufacturer or retailer or learn more at intel.com/benchmarks.

Intel, the Intel logo, Intel Inside, the Intel Inside logo, Intel Atom, Intel Core, Iris, Movidius, Myriad, Intel Nervana, OpenVINO, Intel Optane, Stratix, and Xeon are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries.

*Other names and brands may be claimed as the property of others.

© Intel Corporation