

Building Deep Reinforcement Learning Applications on Apache Spark with Analytics Zoo using BigDL

Yuhao Yang

Intel Data Analytics Technologies

Agenda

Analytics Zoo overview

Reinforcement learning overview

Reinforcement learning with Analytics zoo

future directions

Analytics Zoo

- ***Analytics + AI Platform for Apache Spark and BigDL***
 - Open source, Scala/Python, Spark 1.6 and 2.X

Analytics Zoo

High level API, Industry pipelines, App demo & Util

BigDL

MKL, Tensors, Layers, optim Methods, all-reduce

Apache Spark

RDD, DataFrame, Scala/Python

Analytics Zoo

High level pipeline APIs

nnframes: Spark DataFrames and ML Pipelines for DL

Keras-style API

autograd: custom layer/loss using auto differentiation

Transfer learning

Analytics Zoo

Built-in deep learning pipelines & models

Object detection: API and pre-trained SSD and Faster-RCNN

Image classification: API and pre-trained VGG, Inception, ResNet, MobileNet, etc.

Text classification API with CNN, LSTM and GRU

Recommendation API with NCF, Wide and Deep etc.

Analytics Zoo

End-to-end reference use cases

reinforcement learning

anomaly detection

sentiment analysis

fraud detection

image augmentation


object detection

variational autoencoder

...

Reinforcement Learning (RL)

- RL is for Decision-making




Examples of RL applications


- Play: Atari, poker, Go, ...
- Interact with users: recommend, Healthcare, chatbot, personalize, ..
- Control: auto-driving, robotics, finance, ...

Deep Reinforcement Learning (DRL)

Agents take **actions** (a) in **state** (s) and receives **rewards** (R)
Goal is to find the **policy** (π) that maximized future rewards



Cartpole



Observation

Type: Box(4)

Num	Observation	Min	Max
0	Cart Position	-2.4	2.4
1	Cart Velocity	-Inf	Inf
2	Pole Angle	$\sim -41.8^\circ$	$\sim 41.8^\circ$
3	Pole Velocity At Tip	-Inf	Inf

Actions

Type: Discrete(2)

Num	Action
0	Push cart to the left
1	Push cart to the right

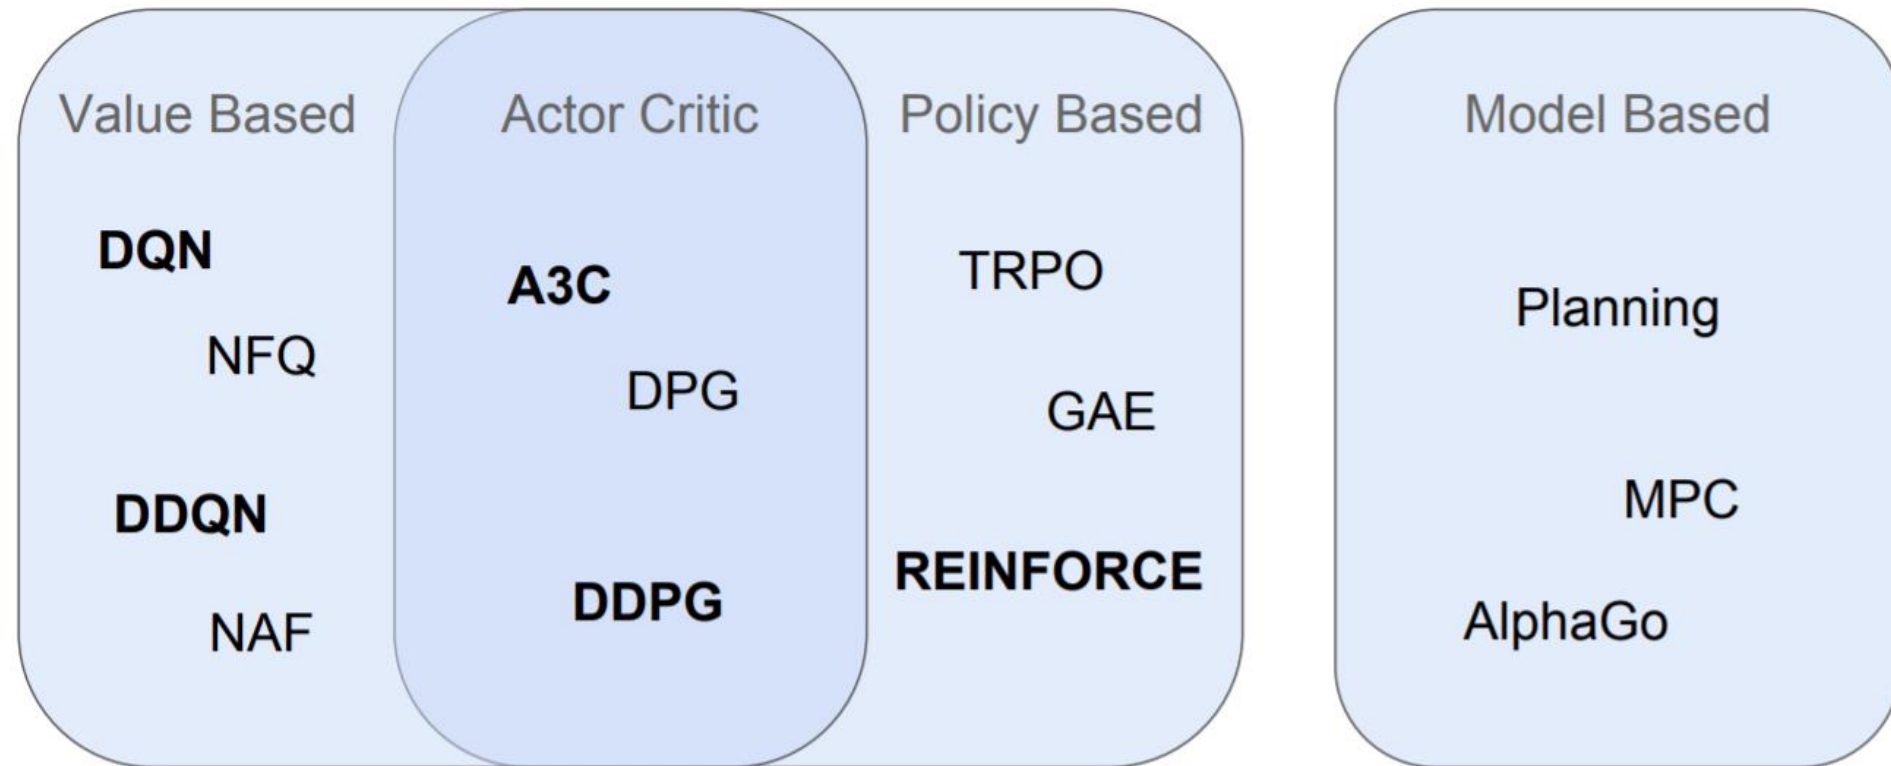
Reward

Reward is 1 for every step taken,

Approaches to Reinforcement Learning

- **Value-based** RL
 - Estimate the optimal value function $Q^*(S,A)$
 - Output of the Neural network is the value for $Q(S, A)$
- **Policy-based** RL
 - Search directly for the optimal policy π^*
 - Output of the neural network is the probability of each action.
- Model-based RL


DRL algo



Examples

- 1. Simple DQN to demo API and train with Spark RDD.
- 2. Distributed REINFORCE

Q-network



Bellman Equation

- ▶ Optimal value maximises over all decisions. Informally:

$$\begin{aligned} Q^*(s, a) &= r_{t+1} + \gamma \max_{a_{t+1}} r_{t+2} + \gamma^2 \max_{a_{t+2}} r_{t+3} + \dots \\ &= r_{t+1} + \gamma \max_{a_{t+1}} Q^*(s_{t+1}, a_{t+1}) \end{aligned}$$

- ▶ Value iteration algorithms solve the Bellman equation

$$Q_{i+1}(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q_i(s', a') \mid s, a \right]$$

DQN critical routines

```
for e in range(EPISODES):
    state = env.reset()
    state = np.reshape(state, [1, state_size])
    for time in range(500):
        action = agent.act(state)     $\epsilon$ -greedy action selection
        next_state, reward, done, _ = env.step(action)
        reward = reward if not done else -10
        next_state = np.reshape(next_state, [1, state_size])
        agent.remember(state, action, reward, next_state, done)
        state = next_state
        if len(agent.memory) > batch_size:
            agent.replay(batch_size)
```


Parallelize the neural network training

```
def replay(self, batch_size):
    X_batch = np.array([0,0,0,0])
    y_batch = np.array([0,0])
    minibatch = random.sample(self.memory, batch_size)
    for state, action, reward, next_state, done in minibatch:
        target = reward
        if not done:
            target = (reward + self.gamma *
                     np.amax(self.model.predict_local(next_state)[0]))
        target_f = self.model.predict_local(state)
        target_f[0][action] = target
        X_batch = np.vstack((X_batch, state))
        y_batch = np.vstack((y_batch, target_f))

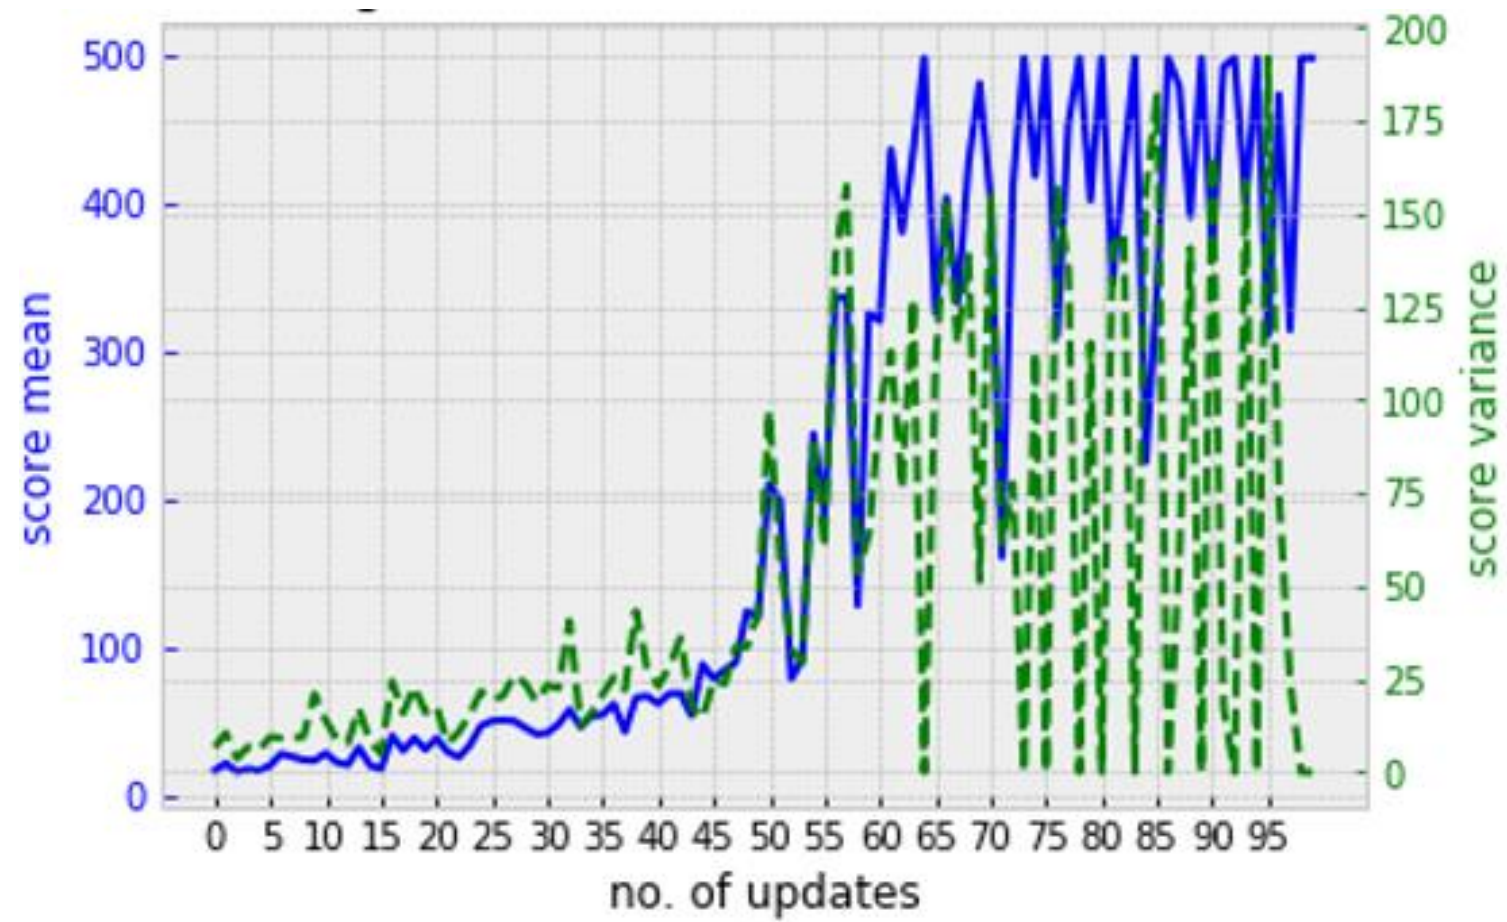
    rdd_sample = to_RDD(X_batch, y_batch)
    self.model.fit(rdd_sample, None, nb_epoch=10, batch_size=batch_size)
```

experience replay

Analytics Zoo Keras-style Model

```
def _build_model(self):  
    # Neural Net for Deep-Q Learning Model  
    model = Sequential()  
    model.add(Dense(24, input_dim=self.state_size, activation='relu'))  
    model.add(Dense(24, activation='relu'))  
    model.add(Dense(self.action_size, activation='linear'))  
    model.compile(loss='mse',  
                  optimizer=Adam(learningrate=self.learning_rate))  
    return model
```

Vanilla DQN



Policy gradients

- In Policy Gradients, we usually use a neural network (or other function approximators) to directly model the action probabilities.
- we tweak the parameters θ of the neural network so that “good” actions will be sampled more likely in the future.

$$\nabla_{\theta} E[R_t] = E[\nabla_{\theta} \log P(a) R_t]$$

REINFORCE

function REINFORCE

 Initialise θ arbitrarily

for each episode $\{s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T\} \sim \pi_\theta$ **do**

for $t = 1$ to $T - 1$ **do**

$\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) v_t$

end for


end for

return θ

end function

Time breakdown

- Game playing takes the most time in each iteration



Distributed REINFORCE

create and cache several agents on each partition as specified by parallelism

and cache it

with DistributedAgents(sc, create_agent=create_agent, parallelism=parallelism) as a:

```
agents = a.agents # a.agents is a RDD[Agent]
```

optimizer = None

```
num_trajs_per_part = int(math.ceil(15.0 / parallelism))
```

```
mean_std = []
```

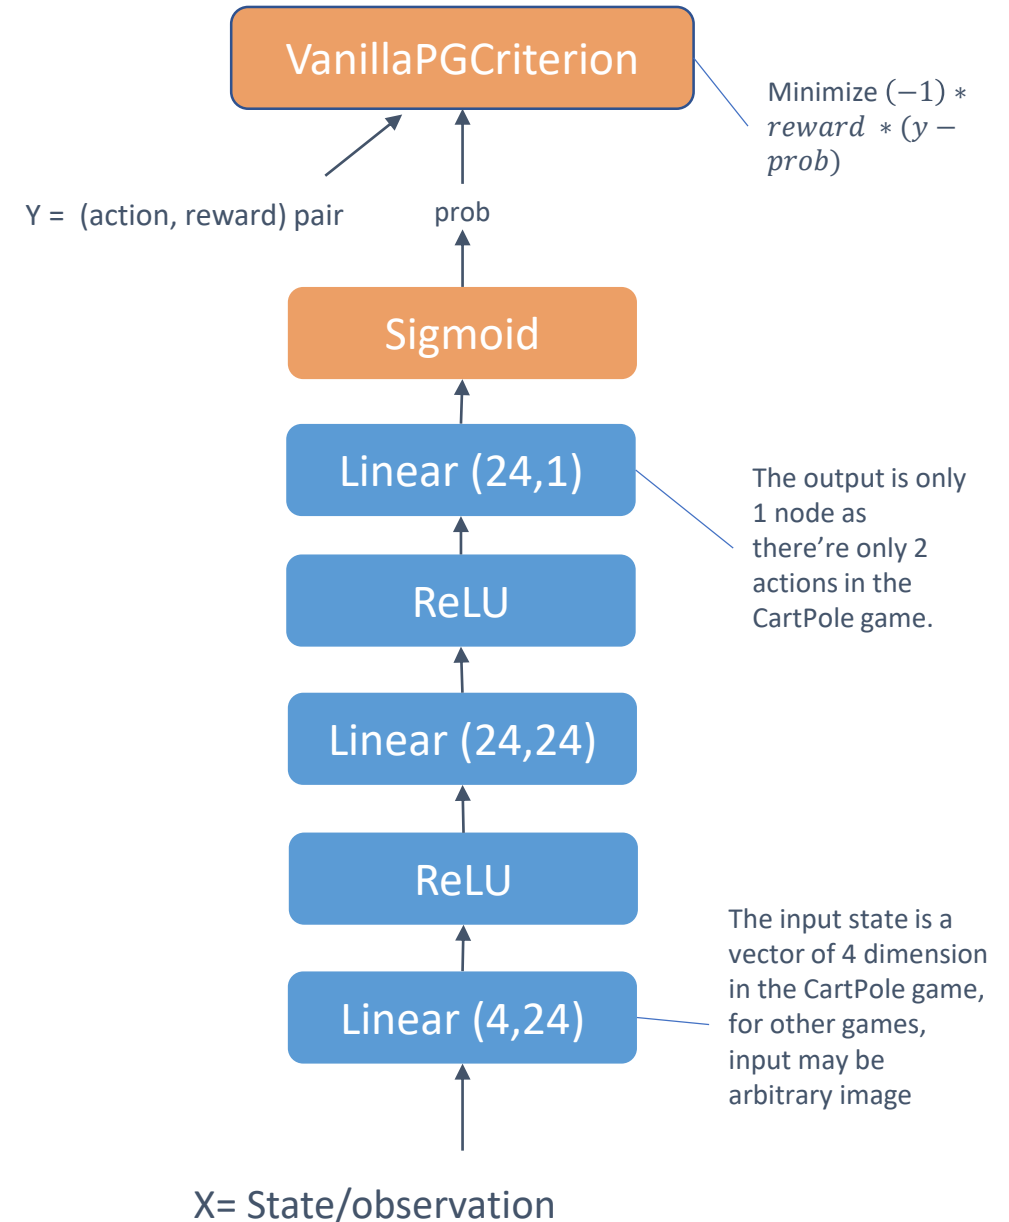
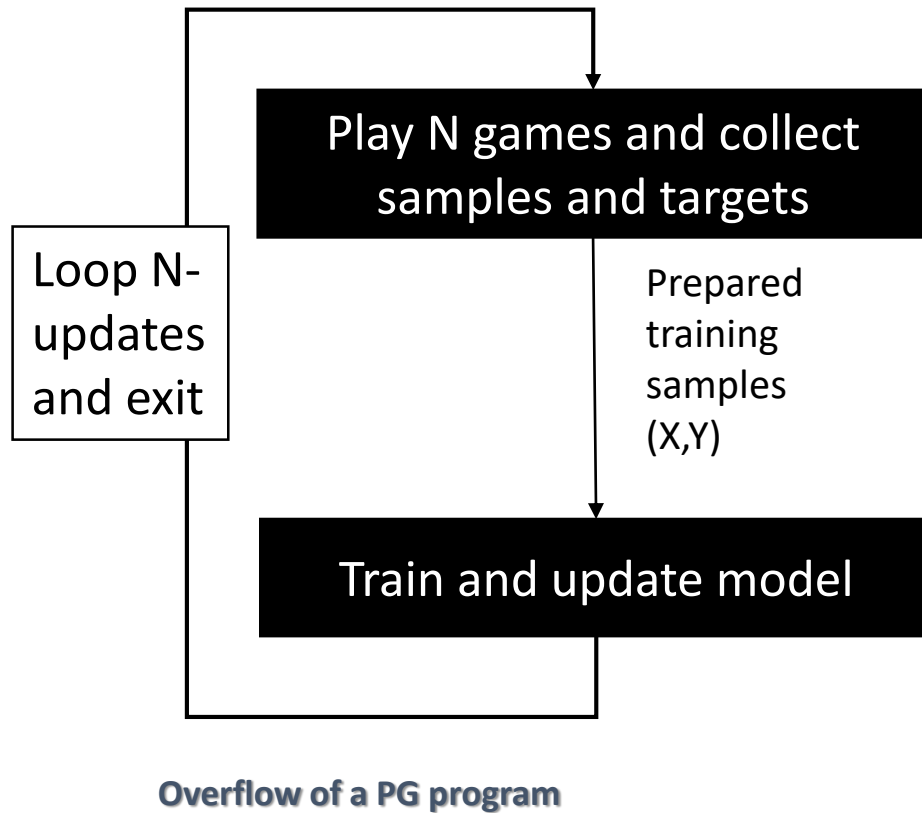
```
for i in range(60):
```

with `SampledTrajs(sc, agents, model, num_trajs_per_part=num_trajs_per_part)` as `trajs`:


```
trajs = trajs.samples \ # samples is a RDD[Trajectory]
```

```
.map(lambda traj: (traj.data["observations"],
                    traj.data["actions"],
                    traj.data["rewards"]))
```

REINFORCE algorithm



Distributed REINFORCE



Other RL algorithms

- Flappy bird with DQN
- Discrete and continuous PPO
- A2C (in roadmap)



Q & A

Analytics Zoo

High level API, Industry pipelines, App demo & Util

<https://github.com/intel-analytics/analytics-zoo>

Thanks Shane Huang and Yang Wang for working on RL implementations.