

A large, abstract graphic occupies the right side of the image. It consists of a dark blue background with a dense, glowing network of interconnected nodes in shades of purple, teal, and light blue. Interspersed among the nodes are binary code sequences (0s and 1s) in various colors, suggesting a digital or data-oriented theme.

Build. Unify. Scale.

WIFI SSID:SparkAISummit | Password: UnifiedAnalytics

ORGANIZED BY
 databricks®

Game Playing Using AI on Apache Spark

Yuhao Yang, Intel

Acknowledgement: Huang, Shengsheng; Yu, Shan

Agenda

- Analytics Zoo
- Problem statement
- Results and Observations

Analytics Zoo

Distributed TensorFlow, Keras and BigDL on Apache Spark

High-Level Pipeline APIs

Distributed TensorFlow and Keras on Spark
Native deep learning support in Spark DataFrames and ML Pipelines
Model serving API for model serving/inference pipelines

Reference Use Cases

Anomaly detection, sentiment analysis, fraud detection, chatbot, sequence prediction, etc.

Built-In Deep Learning Models

Image classification, object detection, text classification, recommendations, sequence-to-sequence, anomaly detection, etc.

Feature Engineering

Feature transformations for

- Image, text, 3D imaging, time series, speech, etc.

Backends

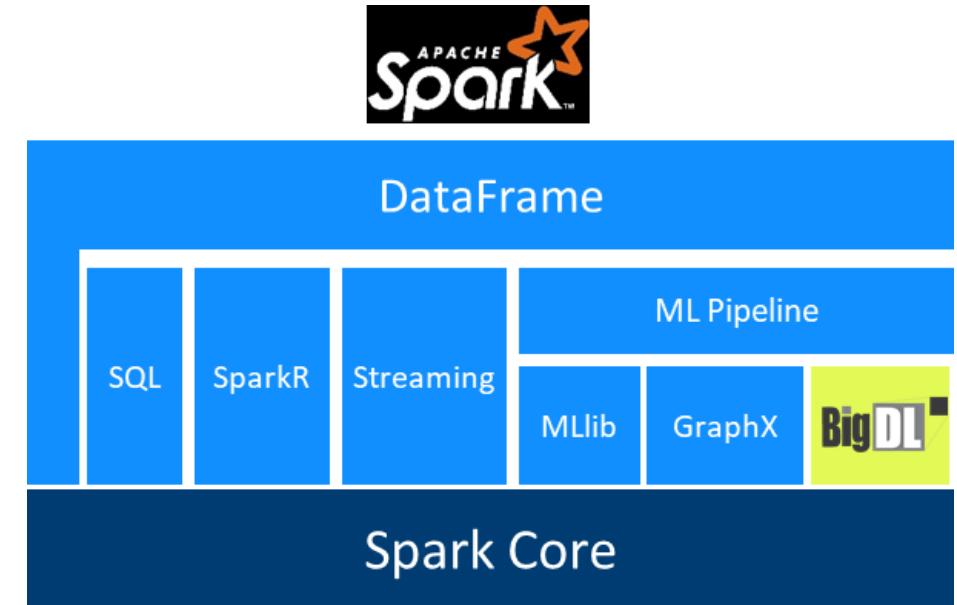
Spark, TensorFlow, Keras, BigDL, OpenVINO, MKL-DNN, etc.

<https://github.com/intel-analytics/analytics-zoo/>



Bringing Deep Learning To Big Data Platform

- Distributed deep learning framework for Apache Spark*
- Make deep learning more accessible to big data users and data scientists
 - Write deep learning applications as [standard Spark programs](#)
 - Run on existing Spark/Hadoop clusters ([no changes needed](#))
- Feature parity with popular deep learning frameworks
 - E.g., Caffe, Torch, Tensorflow, etc.
- High performance (on CPU)
 - Built-in [Intel MKL](#) and multi-threaded programming
- Efficient scale-out
 - Leveraging Spark for distributed training & inference

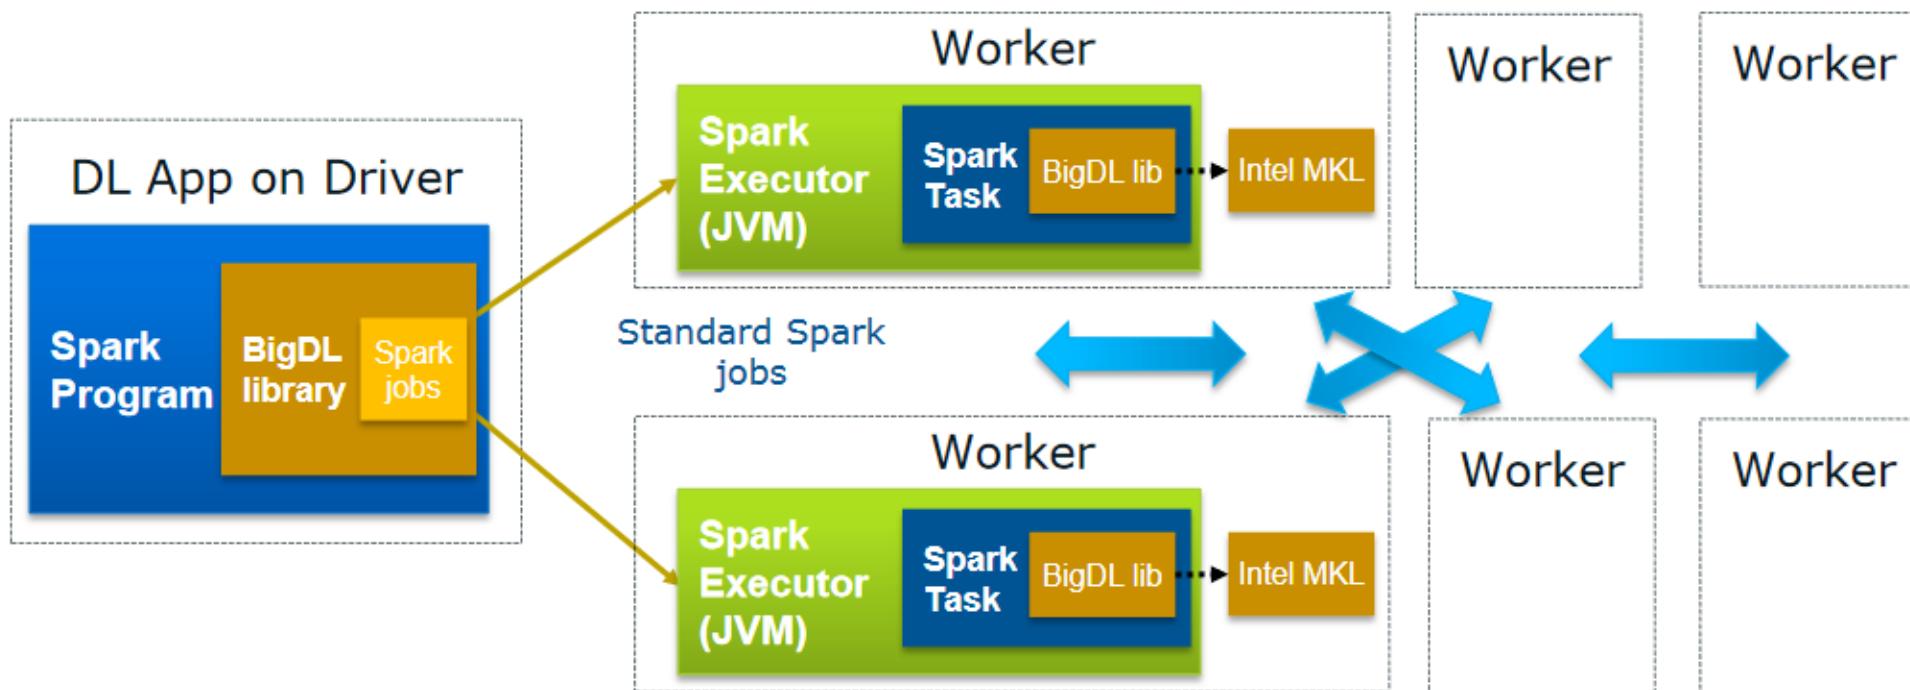


<https://github.com/intel-analytics/BigDL>

<https://bigdl-project.github.io/>

Architecture

- Training using synchronous mini-batch
- Distributed training iterations run as standard Spark jobs
- Data parallelism
- All-reduce



AI-based Game Playing

- Why build AI agents to play game
- Important work/playground for AI game playing:
 - Go: AlphaGo/AlphaZero
 - Atari Games
 - StarCraft/StarCraft II
- Out attempts to build AI to play game
 - Use AI to play a game that has better connection with real world, hopefully transferable to general AI task.
 - Use Spark/distributed technology for high demands of computation and hybrid and complex computation

We choose to play FIFA

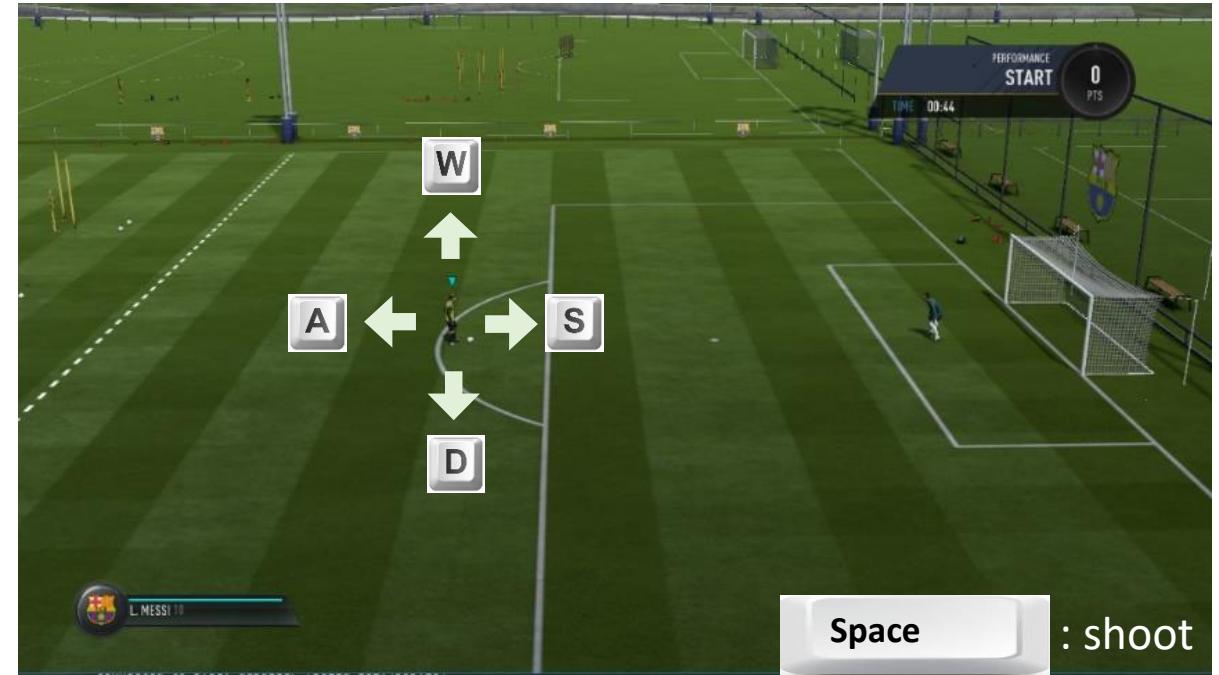
- FIFA18: a real-time 3d modern football simulation video game.
- Why FIFA
 - High resolution screen (1920*1280)
 - Multiple game modes and difficulty levels.
 - Connections to real-world
- DeepMind's recent attempt on 2v2 soccer game on MuJoCo. Next target of full-court FIFA.



Introduction to Shooting Bronze

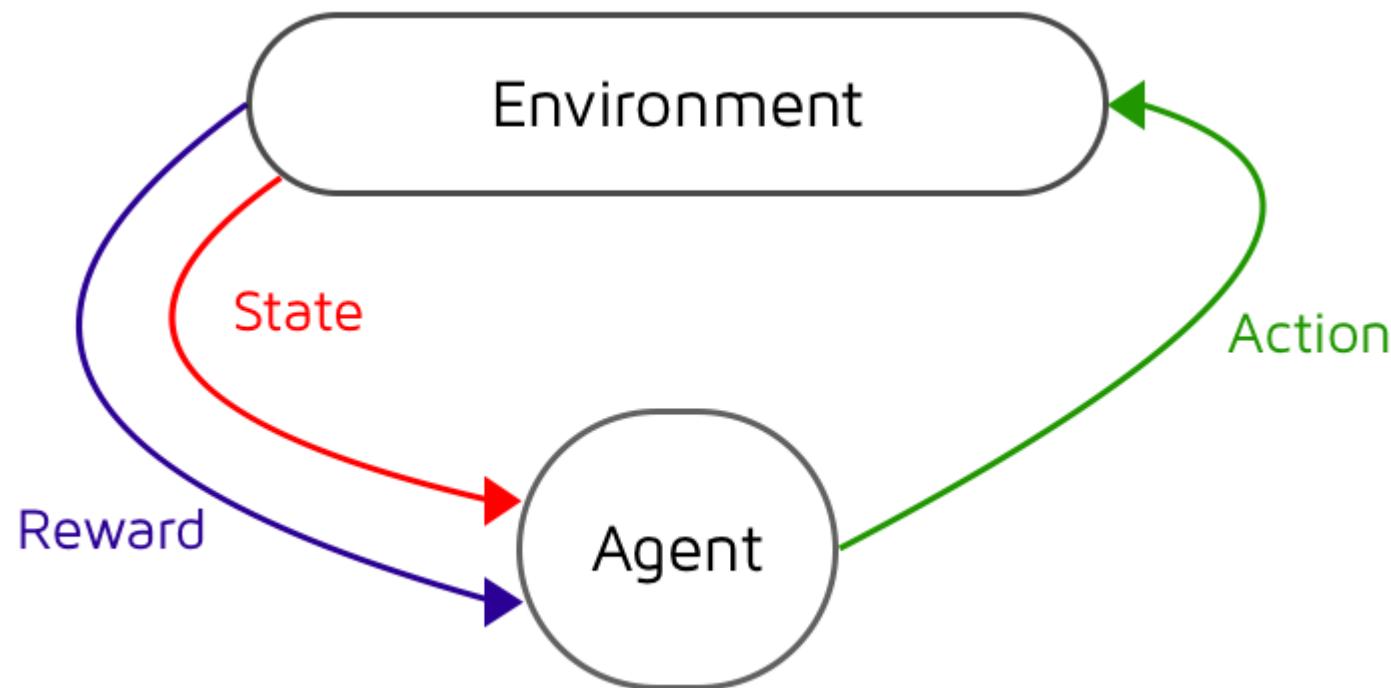
– our major experiment environment

- **Shooting** is one of the mini-games in FIFA, **Bronze** is the easy level. There're also *silver* and *gold* levels.
- Game mode
 - Player & guard 1v1
 - Control player to shoot
 - Game reset after shoot (different start location)
 - Goal: get more score in 44s
- Evaluation
 - Single shoot:
 - Score \leq 200: miss
 - $200 < \text{score} < 1200$: hit
 - Total score:
 - E (<3000) / D (<4000) / C (<6000) / B (<8000) / A (>8000)
- Keyboard control
 - A/S/W/D: left/right/up/down
 - Space: shoot



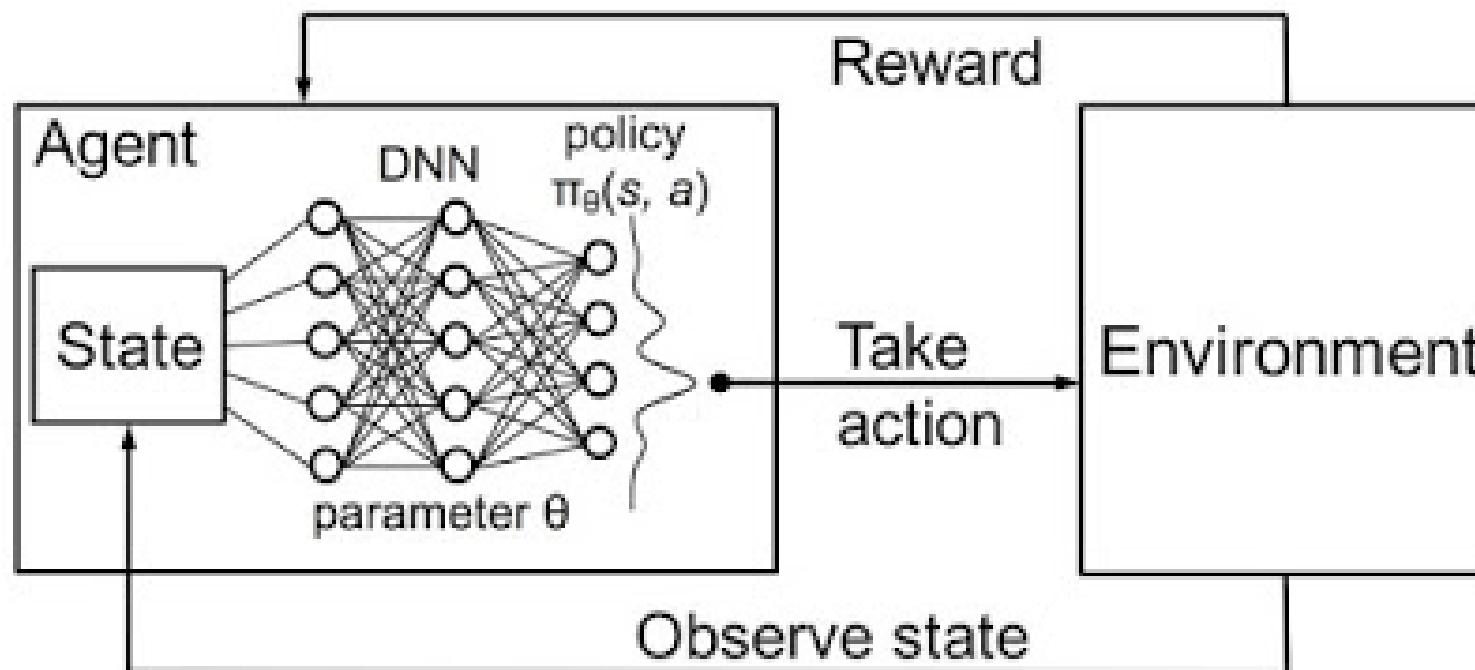
Reinforcement Learning (RL)

- RL is for Decision-making



Deep Reinforcement Learning (DRL)

Agents take **actions** (a) in **state** (s) and receives **rewards** (R)
Goal is to find the **policy** (π) that maximized future rewards



Approaches to Reinforcement Learning

- **Value-based** RL
 - Estimate the optimal value function $Q^*(S,A)$
 - Output of the Neural network is the value for $Q(S, A)$
- **Policy-based** RL
 - Search directly for the optimal policy π^*
 - Output of the neural network is the probability of each action.
- Model-based RL

AI on FIFA platform

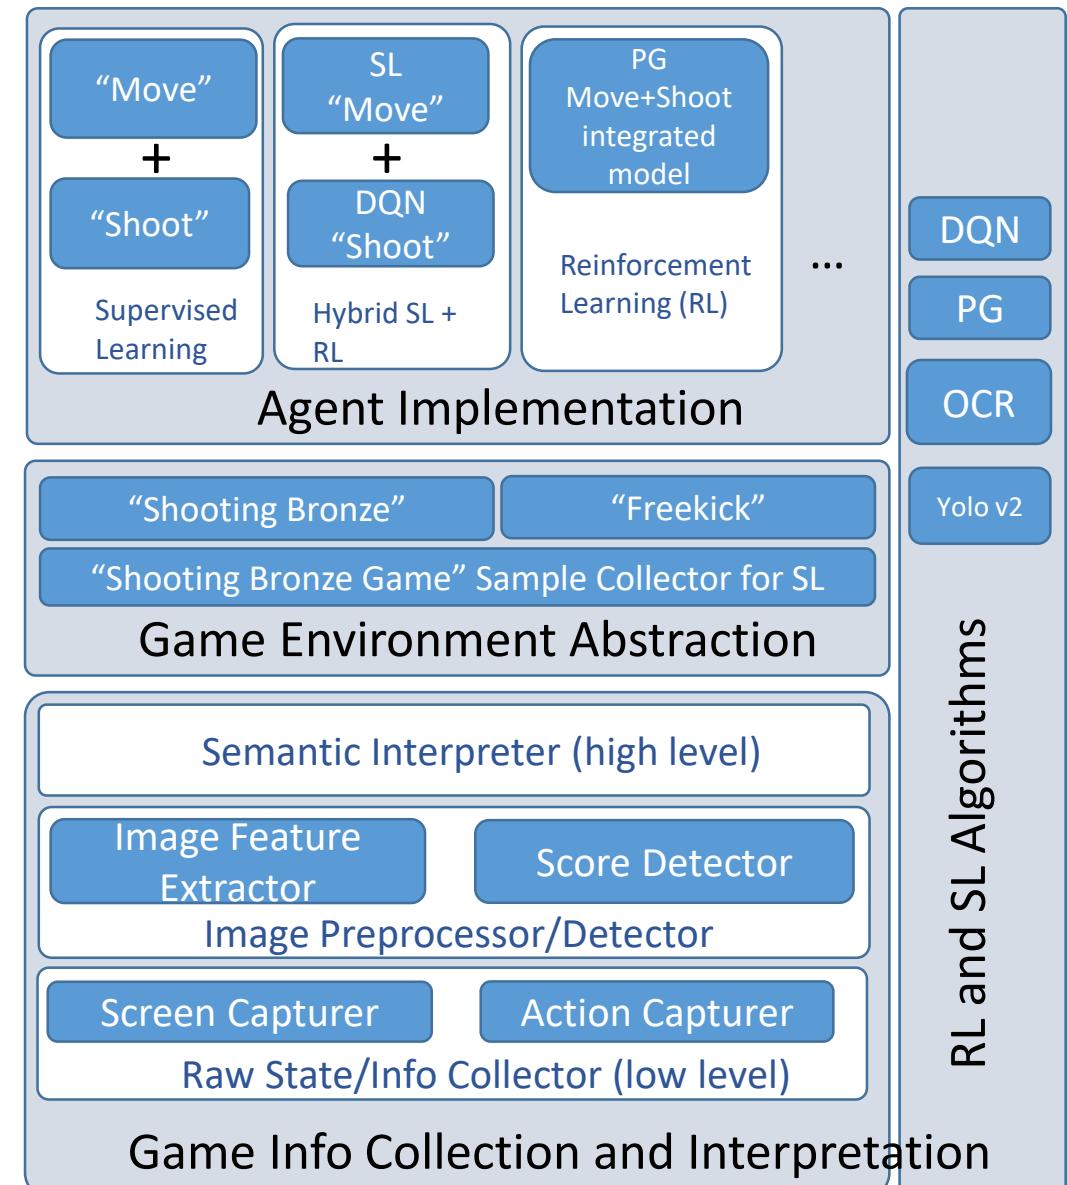
A platform for experimenting various agents and algorithms.

- It interacts with FIFA game by capturing screenshot as input, and issue keyboard input to control the game
- Major components
 - **Game Info Collection and Interpretation:** Hierarchical game info processing, extraction and understanding
 - **Game Environment Abstraction:** 2 available – “Shooting Bronze”, “Free-kick”
 - **Agent Implementation:** 3 available – SL, RL and hybrid SL+RL.
 - SL: Supervised Learning
 - RL: Reinforcement Learning

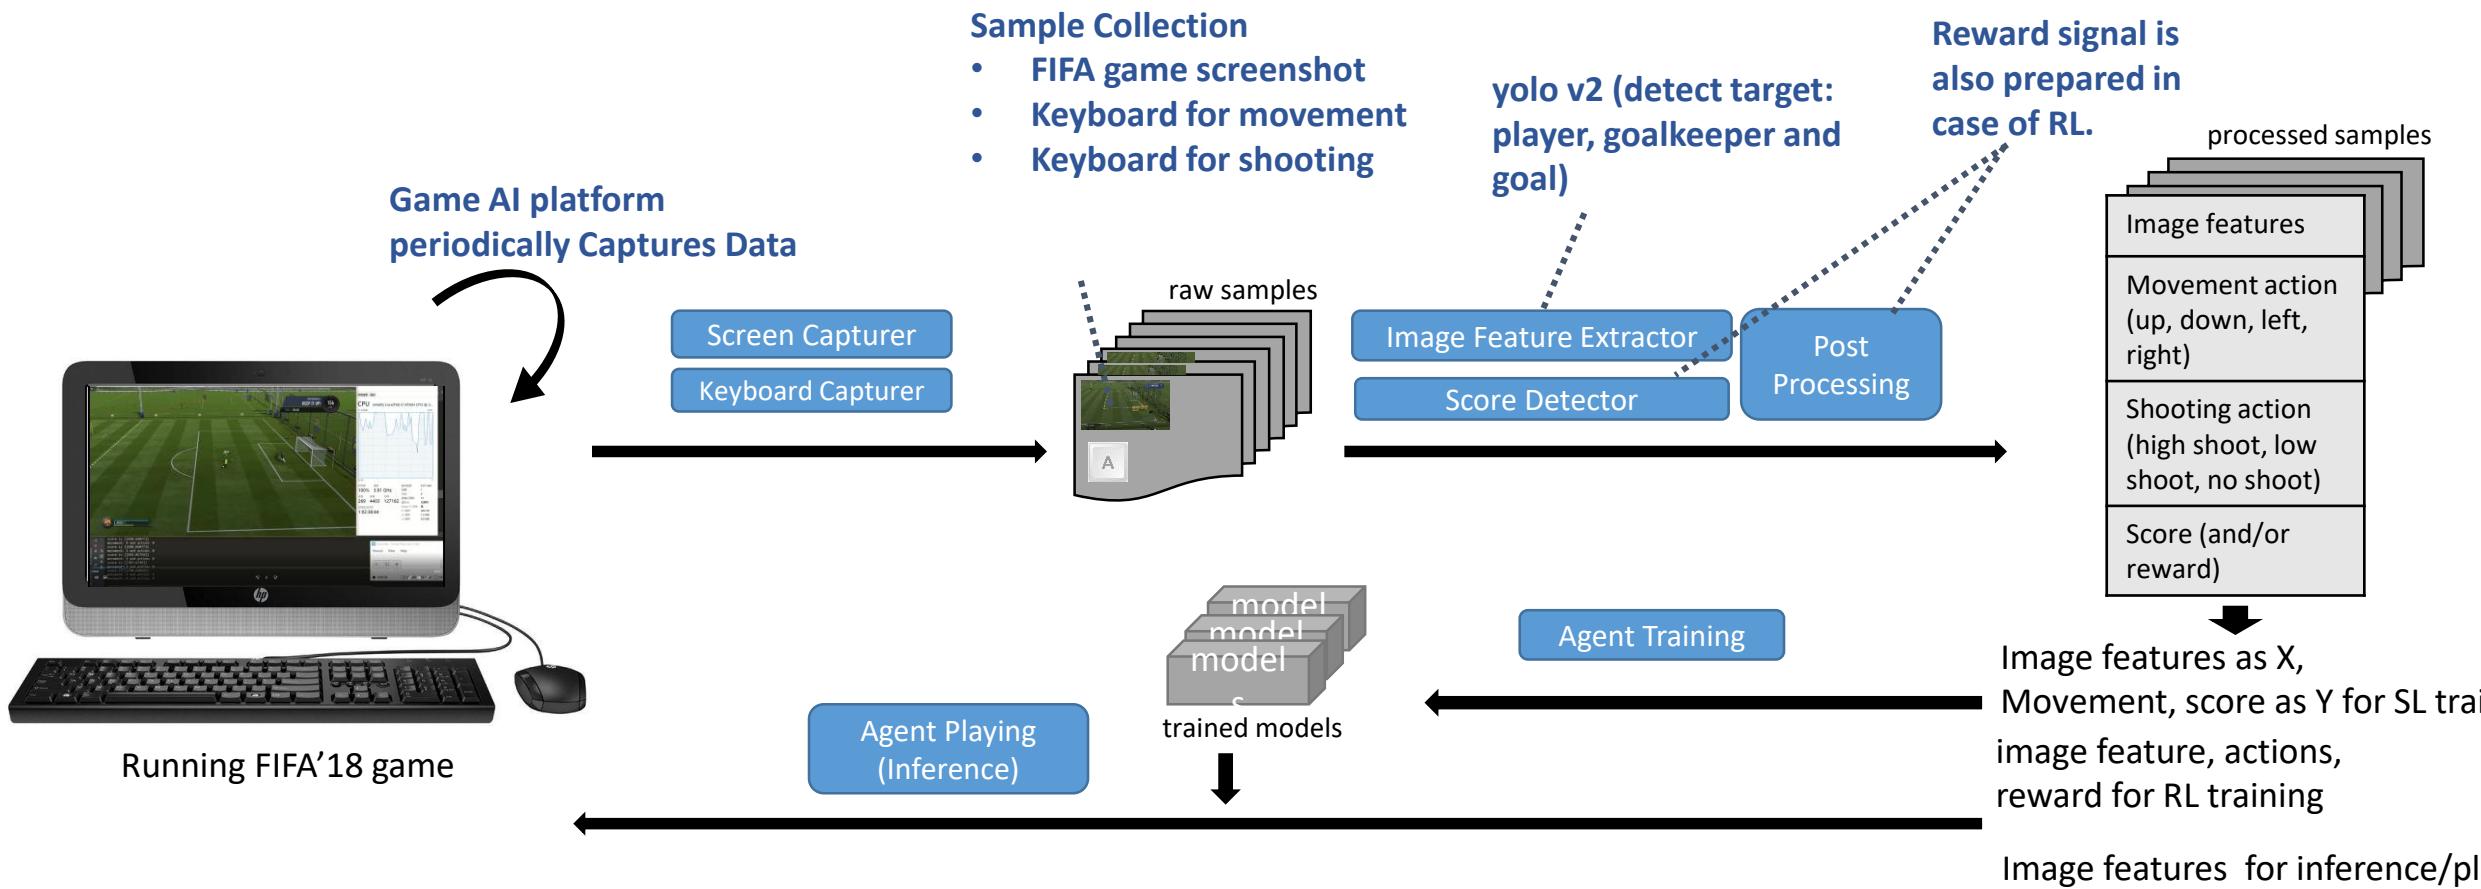


Notes 1: This is the score calculated by the game, considering scoring rate in ~45 secs, the direction of scoring and total number of shooting

Notes 2: This is calculated as $\frac{\text{number of successful shooting}}{\text{total number of shooting attempts}}$

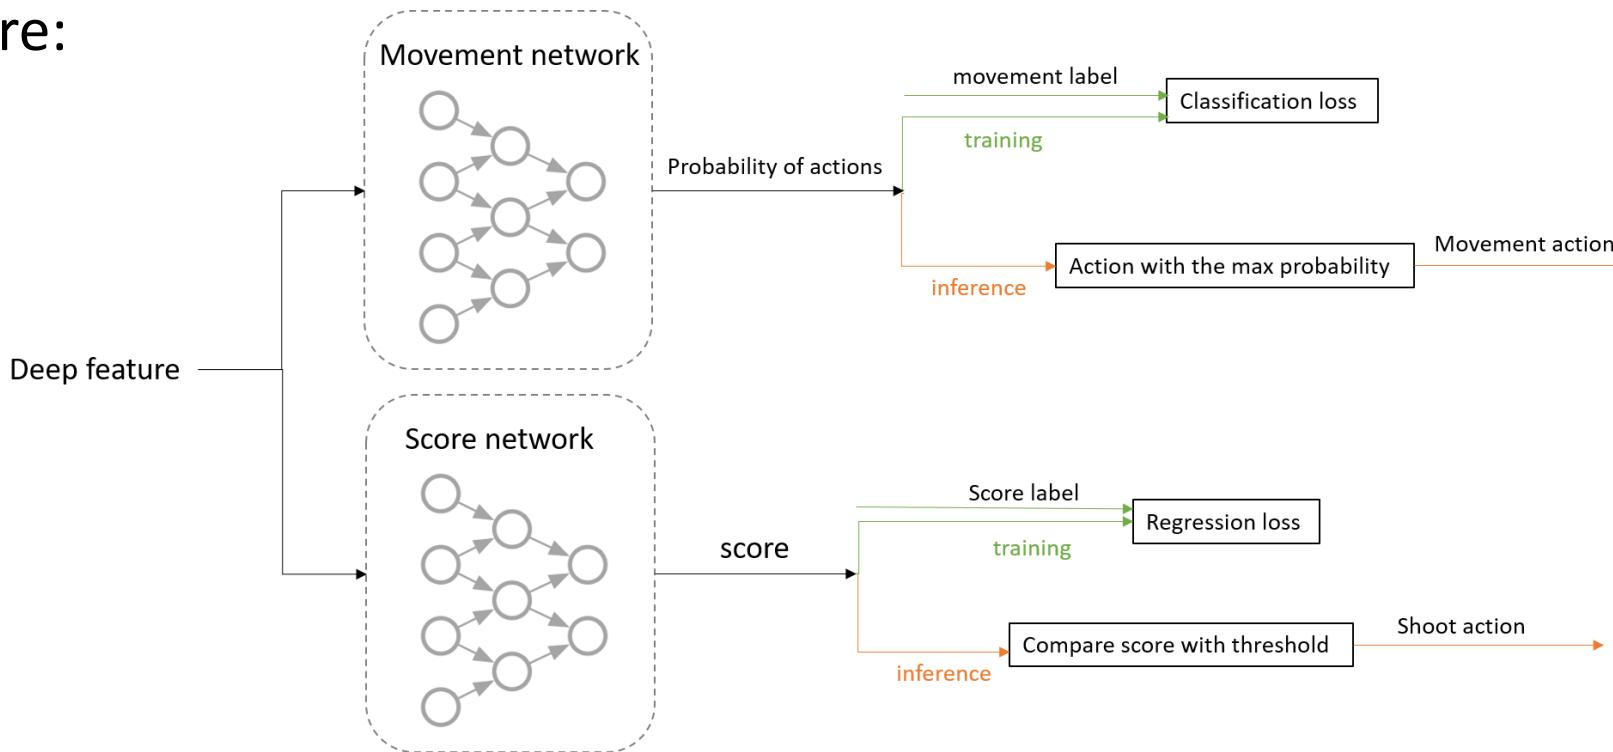
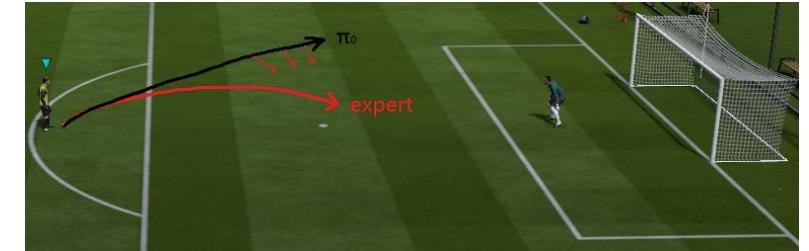


Internal Workflow



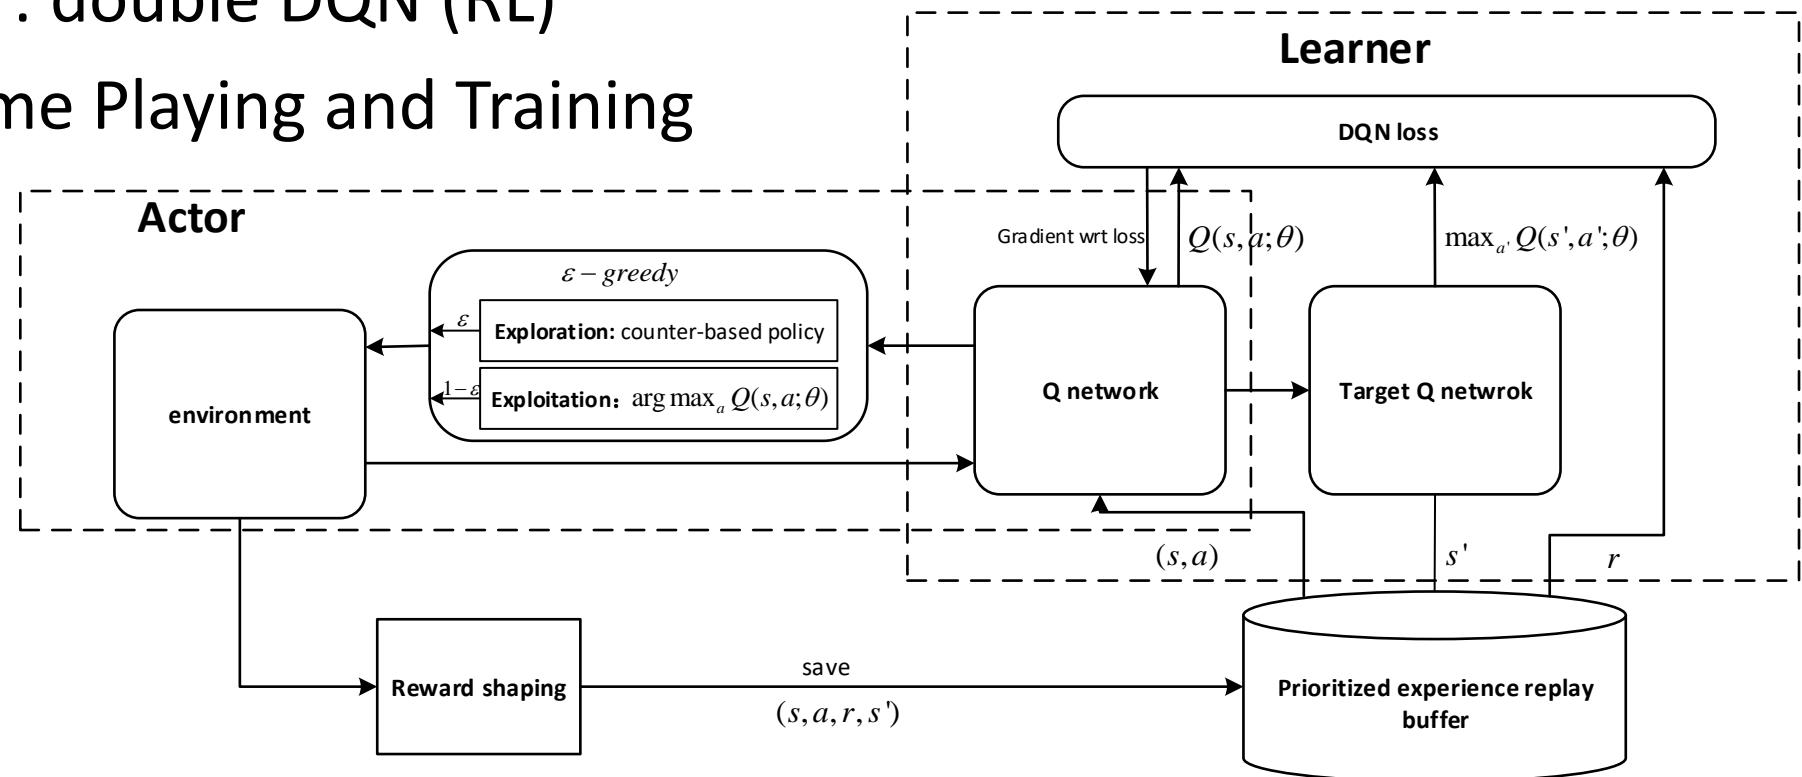
Agents implementation (SL Agent)

- Imitation learning
 - Behavior cloning: supervised learning by classification.
 - DAgger (Dataset Aggregation)
- Structure:



Agents implementation (SL+RL Agent)

- Movement predictor: pre-trained model with imitation learning (SL)
- Shoot predictor: double DQN (RL)
- Distributed Game Playing and Training

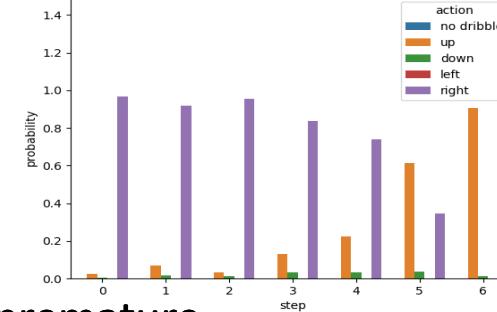


Demo

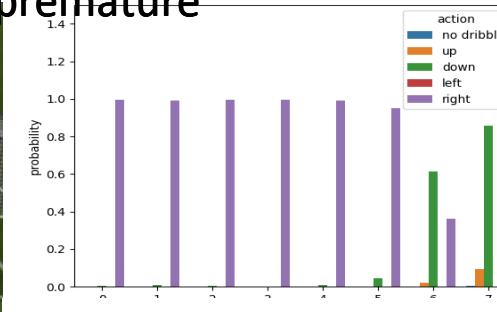
- [Human Demo](#) (the demonstrator used to train SL agent)
- [SL agent Demo](#) (plays better than the human it learnt from)
- [SL+RL agent Demo](#)

Typical trajectory analysis (SL+RL agent)

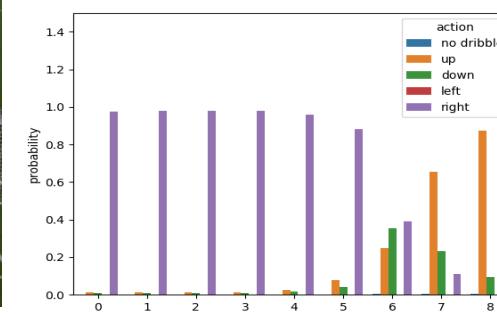
Typical trajectories



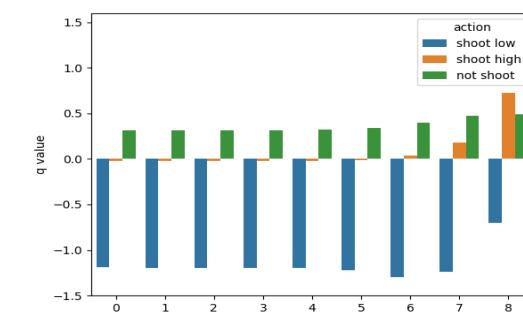
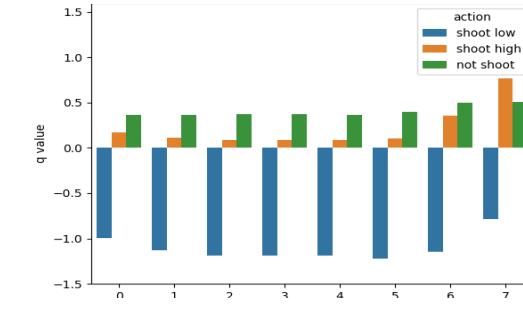
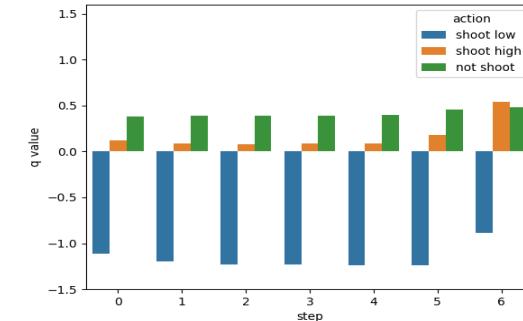
premature



1



Shoot Q-value (RL)



Result

		score	Score_ratio	Convergence speed
Human	beginner	5846.69	50%	-
	master	10112.78	92%	-
	demonstrator	7284.98	84.96%	-
Agent	SL	10345.18	92.54%	-
	RL	5606.31	40.25%	1069.5 epochs
	SL+RL	10514.43	95.59%	749.6 epochs

- SL+RL agent has the best performance and compare to the master level human players.
- SL agent outperforms its demonstrator.
- Performance of RL agent is similar to beginner human players

Key Takeaways

- RL + SL
- Spark is useful for supervised Learning and imitation learning training, especially when train data is very large (which is very possible if scenario is very complex, e.g. full-court).
premature
- Latency problem on small batch updates can be mitigated without hurting agent performance, as long as the (RL) algorithms are carefully chosen and tuned.

Future Work

- Better performant RL agents
- New hybrid agent implementations and more data for imitation learning.
- More difficult game modes (“Shooting Silver/Gold”, Full-court)
 - Need small object detection – (ball, path markers, goal keeper, etc.)
 - More complex scenes means more semantic info to process in game states or complicate keyboard actions.
 - Need to handle sparse direct reward signals, i.e. scoreboard – may need indirect performance indicators such as holding time



SPARK+AI
SUMMIT 2019

Thank you!

DON'T FORGET TO RATE
AND REVIEW THE SESSIONS

SEARCH SPARK + AI SUMMIT

SPARK+AI

Acknowledgements

(Alphabetical)

Huang, Shengsheng(Shane)

Yu, Shan

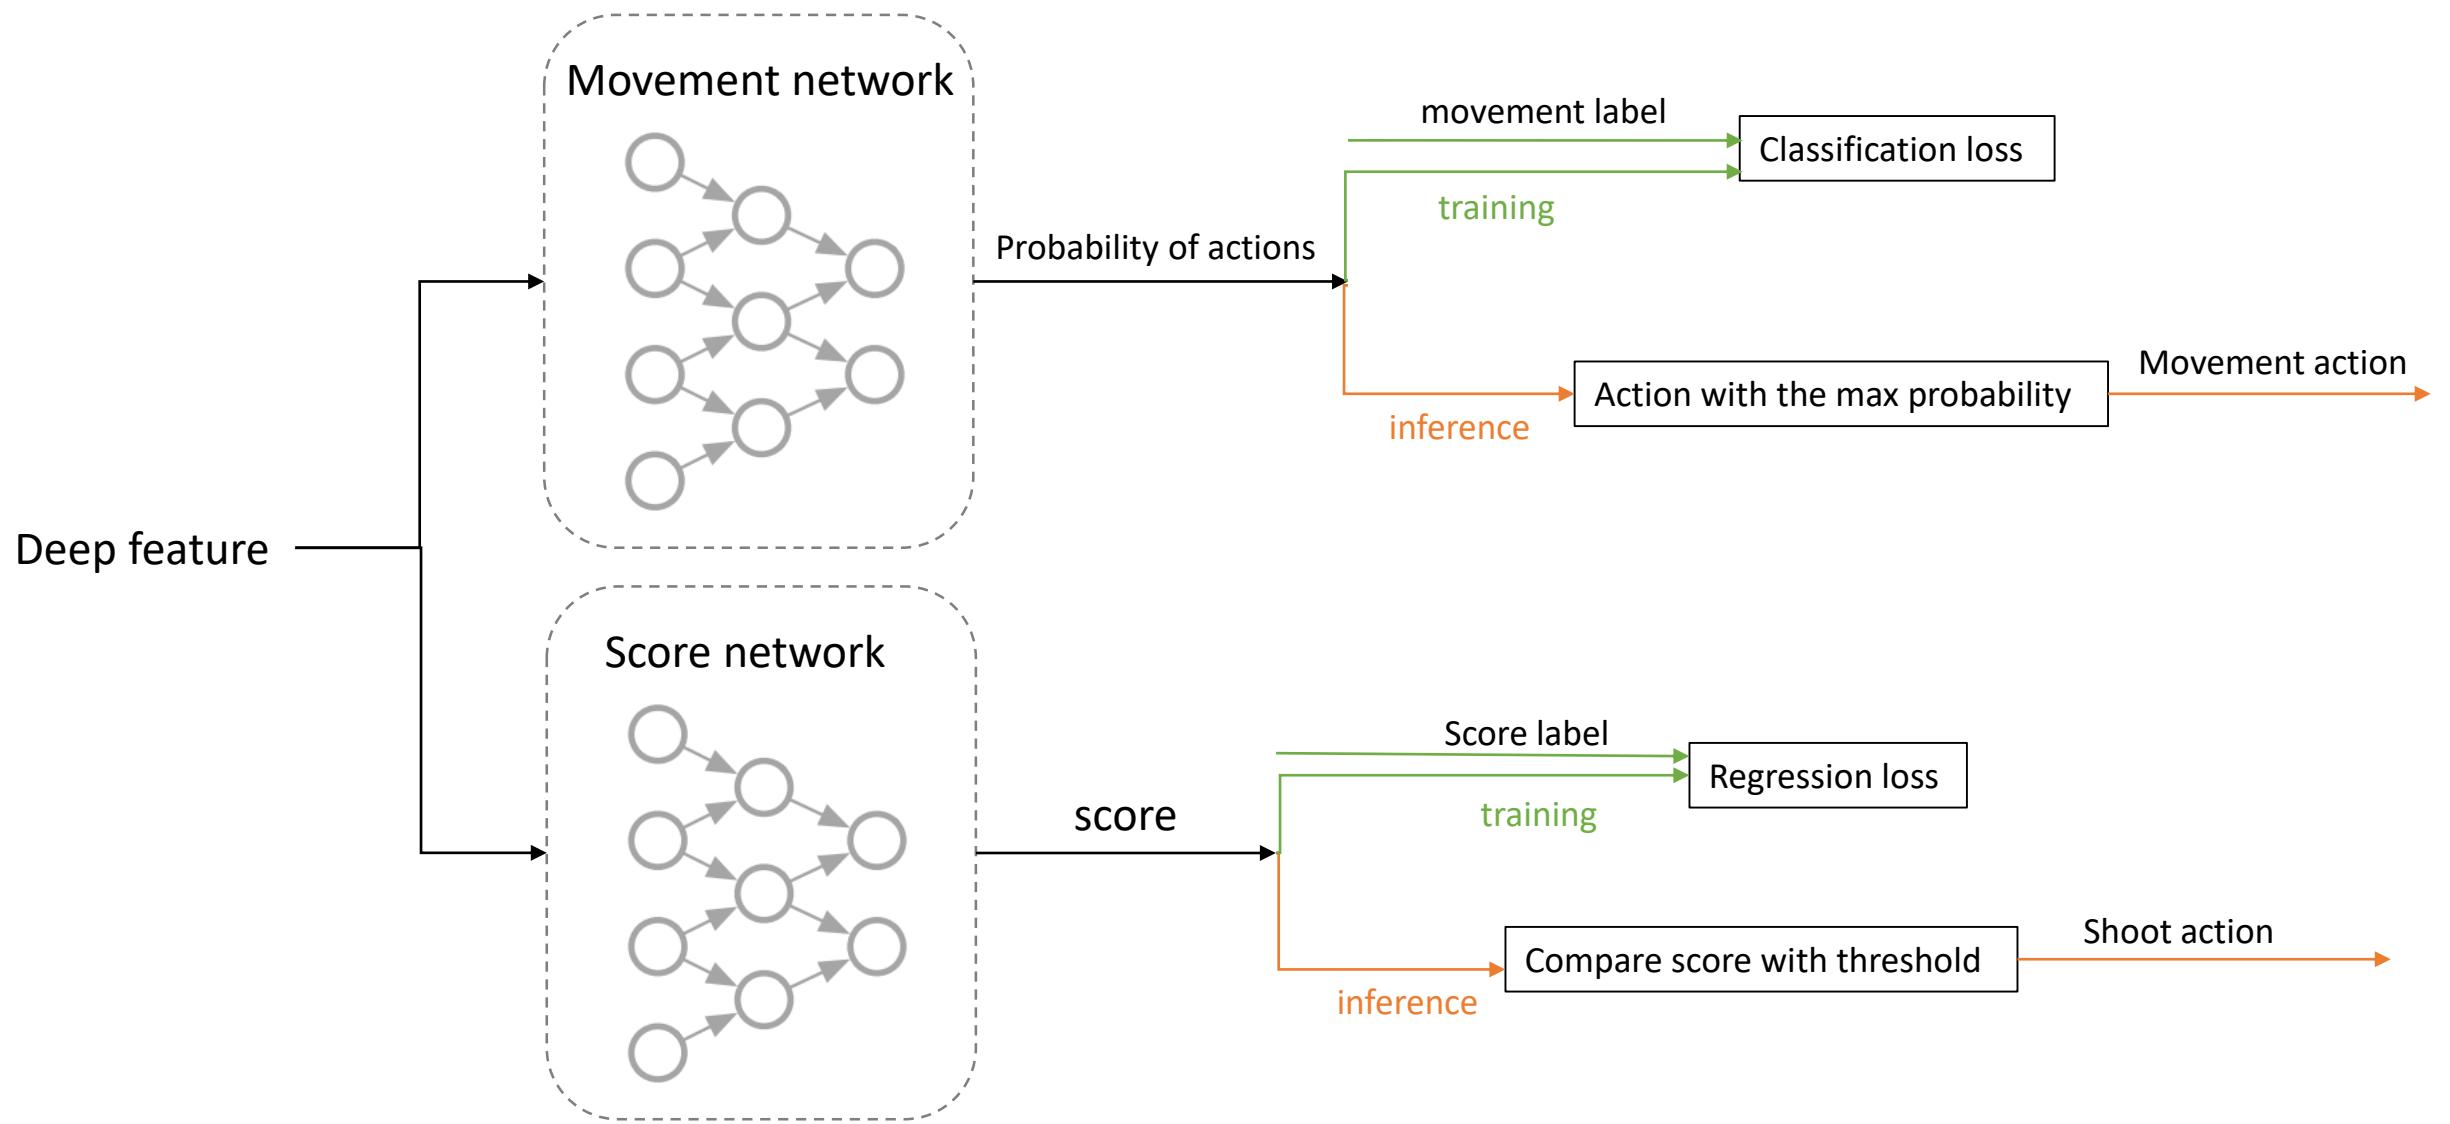
Appendix

miscellaneous

- Some tricks employed to make SL+RL more steady
 - In exploration stage of epsilon-Greedy, use a trick to prevent shooting happen too early.
 - Update DQN network every n-samples instead of 1-sample (Train at the end of each round of shooting instead of each step)
 - Reward shaping

$$R(s_t, a_t, a_{t+1}) = \begin{cases} -1 & \text{s.t. } s_t = \text{initial state and } a_t = \text{shoot} \\ -0.1 & \text{s.t. } s_t \neq \text{terminal and } a_t \neq \text{shoot} \\ 1 & \text{s.t. } s_{t+1} = \text{goal} \\ -1 & \text{s.t. } s_{t+1} = \text{not goal} \end{cases}$$

SL Algorithm Chart



System performance

Module	GPU	Tensorflow	ZOO (tfpark)
Feature_extractor			
movement_predictor	-		
Score_predictor (SL)	-		
dqn_actor (RL)	-		

- For inference only: analytics-zoo(w/ tfpark)'s overhead is ignorable
- For training:
 - For SL, training overhead does not matter, and distributed training had advantage for even larger dataset
 - For RL,
 - On-policy algorithms may not be a good fit using Spark due to the delay (frequent update on small batches make the overhead obvious).
 - We tried do update every n-samples for DQN. Delay can be tolerated and the scores is still not bad.

Tfpark Intro

- Do we need this?